

AI-Driven E-Commerce (Master's Level)  
Book Structure & 14-Week Plan (Placeholder)

Faculty of Computers and Information Sciences  
Mansoura University, Egypt

March 23, 2026



# Contents

|          |                                                                                                                   |           |
|----------|-------------------------------------------------------------------------------------------------------------------|-----------|
| <b>1</b> | <b>Weekly Plan Overview</b>                                                                                       | <b>29</b> |
| <b>2</b> | <b>An Introduction to E-Commerce with a Focus on Machine Learning, Deep Learning, and Artificial Intelligence</b> | <b>31</b> |
| 2.1      | Introduction to E-Commerce . . . . .                                                                              | 31        |
| 2.1.1    | E-Commerce Business Models and Value Chain . . . . .                                                              | 32        |
| 2.2      | Foundations of AI, Machine Learning, and Deep Learning . . . . .                                                  | 32        |
| 2.2.1    | Definitions and Relationships . . . . .                                                                           | 32        |
| 2.2.2    | Types of Learning and E-Commerce Examples . . . . .                                                               | 33        |
| 2.3      | Machine Learning Across the E-Commerce Lifecycle . . . . .                                                        | 33        |
| 2.3.1    | Personalized Recommendations . . . . .                                                                            | 33        |
| 2.3.2    | Search and Ranking . . . . .                                                                                      | 33        |
| 2.3.3    | Segmentation, churn, and CLV . . . . .                                                                            | 33        |
| 2.3.4    | Dynamic Pricing and Promotion Optimization . . . . .                                                              | 33        |
| 2.3.5    | Demand Forecasting and Inventory Management . . . . .                                                             | 34        |
| 2.3.6    | Fraud Detection and Transaction Security . . . . .                                                                | 34        |
| 2.3.7    | Customer Service and Conversational Commerce . . . . .                                                            | 34        |
| 2.4      | Deep Learning in E-Commerce . . . . .                                                                             | 34        |
| 2.4.1    | Neural Recommendation Systems . . . . .                                                                           | 34        |
| 2.4.2    | NLP for Product Text and Reviews . . . . .                                                                        | 34        |
| 2.4.3    | Computer Vision and Visual Search . . . . .                                                                       | 34        |
| 2.5      | Data, Architecture, and Deployment Considerations . . . . .                                                       | 34        |
| 2.6      | Challenges and Responsible AI . . . . .                                                                           | 35        |

|          |                                                                                   |           |
|----------|-----------------------------------------------------------------------------------|-----------|
| 2.7      | Multiple-choice questions (MCQs)                                                  | 35        |
| <b>3</b> | <b>Platforms, Marketplaces, and Business Models (Week 2)</b>                      | <b>37</b> |
| 3.1      | Learning objectives                                                               | 37        |
| 3.2      | From storefront to platform                                                       | 37        |
| 3.3      | Value creation, value capture, and growth                                         | 37        |
| 3.3.1    | Value creation and matching                                                       | 37        |
| 3.3.2    | Value capture models                                                              | 38        |
| 3.3.3    | Network effects                                                                   | 38        |
| 3.4      | Governance and trust as system requirements                                       | 38        |
| 3.5      | Business model patterns (and architectural implications)                          | 38        |
| 3.5.1    | B2C retail                                                                        | 38        |
| 3.5.2    | B2B commerce                                                                      | 39        |
| 3.5.3    | Marketplace                                                                       | 39        |
| 3.5.4    | Digital services (subscriptions)                                                  | 39        |
| 3.6      | Where AI fits (system-level view)                                                 | 39        |
| 3.7      | Hands-on (placeholder)                                                            | 39        |
| 3.8      | Case studies (placeholder)                                                        | 39        |
| 3.9      | Multiple-choice questions (MCQs)                                                  | 40        |
| <b>4</b> | <b>E-Commerce Data Foundations: Events, Experimentation, and Metrics (Week 3)</b> | <b>41</b> |
| 4.1      | Learning objectives                                                               | 41        |
| 4.2      | Why data foundations matter                                                       | 41        |
| 4.3      | From business process to event model                                              | 42        |
| 4.3.1    | Events versus state                                                               | 42        |
| 4.3.2    | Event taxonomies                                                                  | 43        |
| 4.4      | Event schema design                                                               | 43        |
| 4.4.1    | Required design principles                                                        | 43        |
| 4.4.2    | Example event record                                                              | 44        |
| 4.5      | Identity resolution and entity modeling                                           | 44        |

|          |                                                             |           |
|----------|-------------------------------------------------------------|-----------|
| 4.5.1    | Identifiers are not interchangeable . . . . .               | 45        |
| 4.5.2    | Identity stitching . . . . .                                | 45        |
| 4.5.3    | Privacy-aware identity design . . . . .                     | 45        |
| 4.6      | Data quality as a product requirement . . . . .             | 46        |
| 4.6.1    | Data contracts . . . . .                                    | 46        |
| 4.6.2    | Validation and observability . . . . .                      | 46        |
| 4.7      | Attribution and its limits . . . . .                        | 47        |
| 4.7.1    | Common attribution models . . . . .                         | 47        |
| 4.7.2    | Selection bias and measurement bias . . . . .               | 47        |
| 4.8      | Metrics layers and decision quality . . . . .               | 48        |
| 4.8.1    | North Star and supporting metrics . . . . .                 | 48        |
| 4.8.2    | Metric dimensions . . . . .                                 | 48        |
| 4.8.3    | Retention and cohort logic . . . . .                        | 49        |
| 4.9      | Experimentation in e-commerce . . . . .                     | 49        |
| 4.9.1    | A/B testing fundamentals . . . . .                          | 49        |
| 4.9.2    | Common experimental pitfalls . . . . .                      | 49        |
| 4.9.3    | Bandits and adaptive experimentation . . . . .              | 50        |
| 4.10     | From analytics to machine learning features . . . . .       | 50        |
| 4.11     | Hands-on lab: build a minimal metrics layer . . . . .       | 50        |
| 4.11.1   | Lab scenario . . . . .                                      | 51        |
| 4.11.2   | Lab tasks . . . . .                                         | 51        |
| 4.11.3   | Expected learning outcome . . . . .                         | 51        |
| 4.12     | Managerial and architectural implications . . . . .         | 51        |
| 4.13     | Multiple-choice questions (MCQs) . . . . .                  | 52        |
| 4.14     | Exercises (short) . . . . .                                 | 53        |
| 4.15     | Mini-case (Odoo-linked, vendor-neutral) . . . . .           | 53        |
| <b>5</b> | <b>Enterprise Architecture for E-Commerce (Week 4)</b>      | <b>55</b> |
| 5.1      | Learning objectives . . . . .                               | 55        |
| 5.2      | Why enterprise architecture matters in e-commerce . . . . . | 55        |

|        |                                                                    |    |
|--------|--------------------------------------------------------------------|----|
| 5.3    | Enterprise architecture as a set of views . . . . .                | 56 |
| 5.4    | Capability maps . . . . .                                          | 56 |
| 5.4.1  | Core e-commerce capabilities . . . . .                             | 57 |
| 5.4.2  | Why capability maps are useful . . . . .                           | 57 |
| 5.5    | Value streams . . . . .                                            | 57 |
| 5.5.1  | Typical e-commerce value streams . . . . .                         | 58 |
| 5.5.2  | Why value streams matter architecturally . . . . .                 | 58 |
| 5.6    | Domain boundaries and bounded contexts . . . . .                   | 58 |
| 5.6.1  | Typical commerce domains . . . . .                                 | 58 |
| 5.6.2  | Good boundaries versus bad boundaries . . . . .                    | 59 |
| 5.7    | Systems of record and systems of engagement . . . . .              | 59 |
| 5.7.1  | Typical patterns . . . . .                                         | 59 |
| 5.7.2  | Systems of engagement . . . . .                                    | 60 |
| 5.8    | Reference architecture for AI-driven e-commerce . . . . .          | 60 |
| 5.8.1  | Where AI fits . . . . .                                            | 60 |
| 5.9    | Current state, target state, and transition architecture . . . . . | 61 |
| 5.9.1  | Current-state analysis . . . . .                                   | 61 |
| 5.9.2  | Target architecture . . . . .                                      | 61 |
| 5.9.3  | Transition architecture . . . . .                                  | 61 |
| 5.10   | Architectural concerns specific to AI . . . . .                    | 62 |
| 5.11   | Operating model and team topology . . . . .                        | 62 |
| 5.11.1 | Typical team patterns . . . . .                                    | 62 |
| 5.12   | Artifacts for students and architects . . . . .                    | 62 |
| 5.13   | Hands-on artifacts . . . . .                                       | 63 |
| 5.14   | Managerial implications . . . . .                                  | 63 |
| 5.15   | Multiple-choice questions (MCQs) . . . . .                         | 64 |
| 5.16   | Exercises (short) . . . . .                                        | 65 |
| 5.17   | Mini-case (Odoo-linked, vendor-neutral) . . . . .                  | 65 |

|        |                                                         |    |
|--------|---------------------------------------------------------|----|
| 6.1    | Learning objectives . . . . .                           | 67 |
| 6.2    | Why integration patterns matter . . . . .               | 67 |
| 6.3    | Integration styles . . . . .                            | 68 |
| 6.3.1  | Synchronous request-response . . . . .                  | 68 |
| 6.3.2  | Asynchronous events . . . . .                           | 68 |
| 6.3.3  | Batch and file-based integration . . . . .              | 68 |
| 6.4    | Choosing the right integration style . . . . .          | 69 |
| 6.4.1  | Questions to ask . . . . .                              | 69 |
| 6.4.2  | Typical choices in commerce . . . . .                   | 69 |
| 6.5    | API-first integration . . . . .                         | 70 |
| 6.5.1  | Good API design principles . . . . .                    | 70 |
| 6.5.2  | Commands versus queries . . . . .                       | 70 |
| 6.6    | Idempotency . . . . .                                   | 70 |
| 6.6.1  | Why idempotency is required . . . . .                   | 70 |
| 6.6.2  | Idempotency keys . . . . .                              | 70 |
| 6.7    | Events and event-driven architecture . . . . .          | 71 |
| 6.7.1  | Domain events versus technical events . . . . .         | 71 |
| 6.7.2  | Benefits of event-driven integration . . . . .          | 71 |
| 6.7.3  | Costs of event-driven integration . . . . .             | 71 |
| 6.8    | Reliability patterns in distributed workflows . . . . . | 72 |
| 6.8.1  | Retries and backoff . . . . .                           | 72 |
| 6.8.2  | Dead-letter handling . . . . .                          | 72 |
| 6.8.3  | Timeouts and circuit breakers . . . . .                 | 72 |
| 6.8.4  | Reconciliation jobs . . . . .                           | 72 |
| 6.9    | The outbox pattern . . . . .                            | 72 |
| 6.9.1  | How the outbox pattern works . . . . .                  | 73 |
| 6.9.2  | Why it matters in commerce . . . . .                    | 73 |
| 6.10   | Message contracts and schema governance . . . . .       | 73 |
| 6.10.1 | What a contract should define . . . . .                 | 73 |
| 6.10.2 | Backward compatibility . . . . .                        | 74 |

|          |                                                              |           |
|----------|--------------------------------------------------------------|-----------|
| 6.11     | Observability for integrations . . . . .                     | 74        |
| 6.11.1   | Minimum signals . . . . .                                    | 74        |
| 6.11.2   | Correlation and traceability . . . . .                       | 74        |
| 6.12     | iPaaS, ESB, brokers, and workflow orchestration . . . . .    | 75        |
| 6.12.1   | iPaaS and ESB concepts . . . . .                             | 75        |
| 6.12.2   | Event brokers and streams . . . . .                          | 75        |
| 6.12.3   | Workflow orchestration . . . . .                             | 75        |
| 6.13     | Integration anti-patterns . . . . .                          | 75        |
| 6.14     | Hands-on lab: modeling an order lifecycle . . . . .          | 76        |
| 6.14.1   | Lab scenario . . . . .                                       | 76        |
| 6.14.2   | Lab tasks . . . . .                                          | 76        |
| 6.14.3   | Expected learning outcome . . . . .                          | 76        |
| 6.15     | Managerial and architectural implications . . . . .          | 76        |
| 6.16     | Multiple-choice questions (MCQs) . . . . .                   | 77        |
| 6.17     | Exercises (short) . . . . .                                  | 78        |
| 6.18     | Mini-case (Odoo-linked, vendor-neutral) . . . . .            | 78        |
| <b>7</b> | <b>ERP/CRM/SCM Integration in Commerce (Week 6)</b>          | <b>79</b> |
| 7.1      | Learning objectives . . . . .                                | 79        |
| 7.2      | Why ERP/CRM/SCM integration is central to commerce . . . . . | 79        |
| 7.3      | The enterprise systems perspective . . . . .                 | 80        |
| 7.3.1    | ERP . . . . .                                                | 80        |
| 7.3.2    | CRM . . . . .                                                | 80        |
| 7.3.3    | SCM and warehouse systems . . . . .                          | 80        |
| 7.4      | Commerce versus enterprise ownership . . . . .               | 80        |
| 7.4.1    | Common ownership patterns . . . . .                          | 81        |
| 7.4.2    | Why ownership must be explicit . . . . .                     | 81        |
| 7.5      | Systems of record for key entities . . . . .                 | 81        |
| 7.5.1    | Customer . . . . .                                           | 81        |
| 7.5.2    | Product . . . . .                                            | 81        |



|        |                                                           |    |
|--------|-----------------------------------------------------------|----|
| 7.5.3  | Price list . . . . .                                      | 82 |
| 7.5.4  | Order . . . . .                                           | 82 |
| 7.5.5  | Invoice and payment reconciliation . . . . .              | 82 |
| 7.5.6  | Shipment . . . . .                                        | 82 |
| 7.6    | Master data management (MDM) . . . . .                    | 82 |
| 7.6.1  | Why MDM matters . . . . .                                 | 82 |
| 7.6.2  | Practical MDM principles . . . . .                        | 83 |
| 7.7    | Order-to-cash in integrated commerce . . . . .            | 83 |
| 7.7.1  | Typical flow . . . . .                                    | 83 |
| 7.7.2  | Architectural decisions inside the flow . . . . .         | 84 |
| 7.7.3  | Failure modes . . . . .                                   | 84 |
| 7.8    | Inventory, fulfillment, and availability . . . . .        | 84 |
| 7.8.1  | Inventory concepts that should not be conflated . . . . . | 84 |
| 7.8.2  | Synchronization strategies . . . . .                      | 85 |
| 7.9    | Returns and refunds . . . . .                             | 85 |
| 7.9.1  | Typical return flow . . . . .                             | 85 |
| 7.9.2  | Architectural risks . . . . .                             | 85 |
| 7.10   | B2B pricing and account structures . . . . .              | 86 |
| 7.10.1 | Price lists and negotiated terms . . . . .                | 86 |
| 7.10.2 | Account hierarchies . . . . .                             | 86 |
| 7.11   | Using Odoo as a reference ERP . . . . .                   | 87 |
| 7.11.1 | Illustrative Odoo-aligned mapping . . . . .               | 87 |
| 7.11.2 | What should remain vendor-neutral . . . . .               | 87 |
| 7.12   | Integration risks and mitigation strategies . . . . .     | 87 |
| 7.13   | Artifacts for this chapter . . . . .                      | 88 |
| 7.14   | Managerial implications . . . . .                         | 88 |
| 7.15   | Multiple-choice questions (MCQs) . . . . .                | 88 |
| 7.16   | Exercises (short) . . . . .                               | 89 |
| 7.17   | Mini-case (Odoo-linked, vendor-neutral) . . . . .         | 89 |

|          |                                                                      |           |
|----------|----------------------------------------------------------------------|-----------|
| <b>8</b> | <b>Recommenders I: Classical Methods + Learning-to-Rank (Week 7)</b> | <b>91</b> |
| 8.1      | Learning objectives . . . . .                                        | 91        |
| 8.2      | Why recommenders matter in e-commerce . . . . .                      | 91        |
| 8.3      | Recommendation surfaces and tasks . . . . .                          | 92        |
| 8.3.1    | Common recommendation tasks . . . . .                                | 92        |
| 8.3.2    | Retrieval versus ranking . . . . .                                   | 92        |
| 8.4      | Feedback in commerce recommendation . . . . .                        | 92        |
| 8.4.1    | Explicit feedback . . . . .                                          | 92        |
| 8.4.2    | Implicit feedback . . . . .                                          | 93        |
| 8.4.3    | Signal strength . . . . .                                            | 93        |
| 8.5      | Memory-based collaborative filtering . . . . .                       | 93        |
| 8.5.1    | User-based collaborative filtering . . . . .                         | 94        |
| 8.5.2    | Item-based collaborative filtering . . . . .                         | 94        |
| 8.5.3    | Similarity measures . . . . .                                        | 94        |
| 8.6      | Matrix factorization . . . . .                                       | 94        |
| 8.6.1    | Basic idea . . . . .                                                 | 94        |
| 8.6.2    | Why factorization works well . . . . .                               | 95        |
| 8.6.3    | Limitations . . . . .                                                | 95        |
| 8.7      | Implicit-feedback recommendation . . . . .                           | 95        |
| 8.7.1    | Positive-only observations . . . . .                                 | 95        |
| 8.7.2    | Confidence weighting . . . . .                                       | 96        |
| 8.8      | Cold start and sparsity . . . . .                                    | 96        |
| 8.8.1    | User cold start . . . . .                                            | 96        |
| 8.8.2    | Item cold start . . . . .                                            | 96        |
| 8.8.3    | Marketplace and fairness considerations . . . . .                    | 96        |
| 8.9      | Learning-to-rank . . . . .                                           | 97        |
| 8.9.1    | Pointwise, pairwise, and listwise views . . . . .                    | 97        |
| 8.9.2    | Why LTR is important in commerce . . . . .                           | 97        |
| 8.9.3    | Feature design for LTR . . . . .                                     | 97        |
| 8.10     | Offline evaluation . . . . .                                         | 98        |

|          |                                                                           |            |
|----------|---------------------------------------------------------------------------|------------|
| 8.10.1   | Common ranking metrics . . . . .                                          | 98         |
| 8.10.2   | What these metrics measure . . . . .                                      | 98         |
| 8.10.3   | Limits of offline evaluation . . . . .                                    | 98         |
| 8.11     | Online evaluation and business impact . . . . .                           | 99         |
| 8.12     | Biases and practical risks . . . . .                                      | 99         |
| 8.12.1   | Popularity bias . . . . .                                                 | 99         |
| 8.12.2   | Selection bias . . . . .                                                  | 99         |
| 8.12.3   | Position bias . . . . .                                                   | 99         |
| 8.12.4   | Business-constraint tension . . . . .                                     | 100        |
| 8.13     | A practical recommender architecture . . . . .                            | 100        |
| 8.14     | Hands-on lab: a baseline recommender . . . . .                            | 100        |
| 8.14.1   | Lab scenario . . . . .                                                    | 100        |
| 8.14.2   | Lab tasks . . . . .                                                       | 101        |
| 8.14.3   | Expected learning outcome . . . . .                                       | 101        |
| 8.15     | Managerial implications . . . . .                                         | 101        |
| 8.16     | Multiple-choice questions (MCQs) . . . . .                                | 101        |
| 8.17     | Exercises (short) . . . . .                                               | 103        |
| 8.18     | Mini-case (Odoo-linked, vendor-neutral) . . . . .                         | 103        |
| <b>9</b> | <b>Recommenders II: Deep Learning, Embeddings, and Retrieval (Week 8)</b> | <b>105</b> |
| 9.1      | Learning objectives . . . . .                                             | 105        |
| 9.2      | From classical recommenders to deep retrieval . . . . .                   | 105        |
| 9.3      | Modern recommender architecture: retrieval and ranking . . . . .          | 106        |
| 9.3.1    | Candidate generation . . . . .                                            | 106        |
| 9.3.2    | Re-ranking . . . . .                                                      | 106        |
| 9.3.3    | Why multi-stage design matters . . . . .                                  | 106        |
| 9.4      | Embeddings as learned representations . . . . .                           | 106        |
| 9.4.1    | Why embeddings are powerful . . . . .                                     | 106        |
| 9.4.2    | What can be embedded . . . . .                                            | 107        |
| 9.4.3    | Interpretation . . . . .                                                  | 107        |

|        |                                                            |     |
|--------|------------------------------------------------------------|-----|
| 9.5    | Two-tower retrieval models . . . . .                       | 107 |
| 9.5.1  | Basic structure . . . . .                                  | 107 |
| 9.5.2  | Why two towers are useful . . . . .                        | 108 |
| 9.5.3  | Typical inputs . . . . .                                   | 108 |
| 9.6    | Approximate nearest-neighbor retrieval . . . . .           | 108 |
| 9.6.1  | Why exact search is often too expensive . . . . .          | 108 |
| 9.6.2  | Architectural implication . . . . .                        | 109 |
| 9.7    | Sequence models and session-aware recommendation . . . . . | 109 |
| 9.7.1  | Session behavior . . . . .                                 | 109 |
| 9.7.2  | Sequence model families . . . . .                          | 109 |
| 9.7.3  | Why sequence models help . . . . .                         | 110 |
| 9.8    | Multi-modal recommendation . . . . .                       | 110 |
| 9.8.1  | Text and semantic understanding . . . . .                  | 110 |
| 9.8.2  | Image understanding . . . . .                              | 110 |
| 9.8.3  | Fusion . . . . .                                           | 111 |
| 9.9    | Deep re-ranking . . . . .                                  | 111 |
| 9.9.1  | Why re-ranking differs from retrieval . . . . .            | 111 |
| 9.9.2  | Business role of re-ranking . . . . .                      | 111 |
| 9.10   | Training data and objectives . . . . .                     | 111 |
| 9.10.1 | Positive and negative examples . . . . .                   | 111 |
| 9.10.2 | Objective choices . . . . .                                | 112 |
| 9.10.3 | Label quality . . . . .                                    | 112 |
| 9.11   | Cold start in deep recommendation . . . . .                | 112 |
| 9.11.1 | User cold start . . . . .                                  | 112 |
| 9.11.2 | Item cold start . . . . .                                  | 112 |
| 9.11.3 | Catalog churn . . . . .                                    | 113 |
| 9.12   | Diversity, exploration, and bias . . . . .                 | 113 |
| 9.12.1 | Diversity . . . . .                                        | 113 |
| 9.12.2 | Exploration . . . . .                                      | 113 |
| 9.12.3 | Bias . . . . .                                             | 113 |

|           |                                                               |            |
|-----------|---------------------------------------------------------------|------------|
| 9.13      | Serving and latency constraints . . . . .                     | 113        |
| 9.13.1    | Key serving questions . . . . .                               | 114        |
| 9.13.2    | Caching and freshness . . . . .                               | 114        |
| 9.13.3    | Fallback behavior . . . . .                                   | 114        |
| 9.14      | Evaluation beyond offline metrics . . . . .                   | 114        |
| 9.14.1    | Why deep models can mislead offline . . . . .                 | 114        |
| 9.14.2    | Online considerations . . . . .                               | 115        |
| 9.15      | A practical deep recommender stack . . . . .                  | 115        |
| 9.16      | Hands-on lab: retrieval plus re-ranking . . . . .             | 115        |
| 9.16.1    | Lab scenario . . . . .                                        | 116        |
| 9.16.2    | Lab tasks . . . . .                                           | 116        |
| 9.16.3    | Expected learning outcome . . . . .                           | 116        |
| 9.17      | Managerial implications . . . . .                             | 116        |
| 9.18      | Multiple-choice questions (MCQs) . . . . .                    | 116        |
| 9.19      | Exercises (short) . . . . .                                   | 118        |
| 9.20      | Mini-case (Odoo-linked, vendor-neutral) . . . . .             | 118        |
| <b>10</b> | <b>Pricing and Promotions with ML + Optimization (Week 9)</b> | <b>119</b> |
| 10.1      | Learning objectives . . . . .                                 | 119        |
| 10.2      | Why pricing is a core AI problem in commerce . . . . .        | 119        |
| 10.3      | Pricing objectives . . . . .                                  | 120        |
| 10.3.1    | Common objectives . . . . .                                   | 120        |
| 10.3.2    | Why objective clarity matters . . . . .                       | 120        |
| 10.4      | Price, promotion, and discount structure . . . . .            | 120        |
| 10.4.1    | Promotions as structured interventions . . . . .              | 121        |
| 10.5      | Demand forecasting . . . . .                                  | 121        |
| 10.5.1    | What is being forecast . . . . .                              | 121        |
| 10.5.2    | Inputs to demand models . . . . .                             | 122        |
| 10.5.3    | Why forecasting is not enough . . . . .                       | 122        |
| 10.6      | Price elasticity . . . . .                                    | 122        |

|         |                                                                |     |
|---------|----------------------------------------------------------------|-----|
| 10.6.1  | Intuition . . . . .                                            | 122 |
| 10.6.2  | Why elasticity is difficult to estimate . . . . .              | 122 |
| 10.6.3  | Practical considerations . . . . .                             | 123 |
| 10.7    | Promotions and uplift . . . . .                                | 123 |
| 10.7.1  | Observed uplift versus causal uplift . . . . .                 | 123 |
| 10.7.2  | Why uplift matters . . . . .                                   | 123 |
| 10.8    | Optimization under business constraints . . . . .              | 124 |
| 10.8.1  | Typical constraints . . . . .                                  | 124 |
| 10.8.2  | Why unconstrained optimization is unrealistic . . . . .        | 124 |
| 10.9    | Dynamic pricing . . . . .                                      | 124 |
| 10.9.1  | When dynamic pricing can help . . . . .                        | 125 |
| 10.9.2  | When dynamic pricing can hurt . . . . .                        | 125 |
| 10.10   | Competition-aware pricing . . . . .                            | 125 |
| 10.10.1 | Competitive signals . . . . .                                  | 125 |
| 10.10.2 | Strategic caution . . . . .                                    | 126 |
| 10.11   | B2C versus B2B pricing . . . . .                               | 126 |
| 10.11.1 | B2C . . . . .                                                  | 126 |
| 10.11.2 | B2B . . . . .                                                  | 126 |
| 10.12   | Feature design for pricing models . . . . .                    | 127 |
| 10.12.1 | Feature governance . . . . .                                   | 127 |
| 10.13   | Experimentation in pricing . . . . .                           | 127 |
| 10.13.1 | Why pricing experiments are sensitive . . . . .                | 127 |
| 10.13.2 | What to measure . . . . .                                      | 127 |
| 10.13.3 | Guardrails . . . . .                                           | 128 |
| 10.14   | A practical pricing architecture . . . . .                     | 128 |
| 10.14.1 | Decision flow . . . . .                                        | 129 |
| 10.14.2 | Why architecture matters . . . . .                             | 129 |
| 10.15   | Common failure modes . . . . .                                 | 129 |
| 10.16   | Hands-on lab: forecast plus constrained optimization . . . . . | 129 |
| 10.16.1 | Lab scenario . . . . .                                         | 130 |

|                                                         |            |
|---------------------------------------------------------|------------|
| 10.16.2 Lab tasks . . . . .                             | 130        |
| 10.16.3 Expected learning outcome . . . . .             | 130        |
| 10.17 Managerial implications . . . . .                 | 130        |
| 10.18 Multiple-choice questions (MCQs) . . . . .        | 130        |
| 10.19 Exercises (short) . . . . .                       | 132        |
| 10.20 Mini-case (Odoo-linked, vendor-neutral) . . . . . | 132        |
| <b>11 Fraud, Risk, and Trust &amp; Safety (Week 10)</b> | <b>133</b> |
| 11.1 Learning objectives . . . . .                      | 133        |
| 11.2 Why fraud and trust & safety matter . . . . .      | 133        |
| 11.3 Fraud and abuse categories . . . . .               | 134        |
| 11.3.1 Payment fraud . . . . .                          | 134        |
| 11.3.2 Account takeover . . . . .                       | 134        |
| 11.3.3 Promotion and coupon abuse . . . . .             | 134        |
| 11.3.4 Refund and return abuse . . . . .                | 134        |
| 11.3.5 Marketplace and seller abuse . . . . .           | 135        |
| 11.4 Fraud as a machine learning problem . . . . .      | 135        |
| 11.4.1 Class imbalance . . . . .                        | 135        |
| 11.4.2 Adversarial behavior . . . . .                   | 135        |
| 11.4.3 Label delay and label noise . . . . .            | 135        |
| 11.4.4 Asymmetric costs . . . . .                       | 135        |
| 11.5 Signals and features for fraud detection . . . . . | 135        |
| 11.5.1 Transaction signals . . . . .                    | 136        |
| 11.5.2 Account signals . . . . .                        | 136        |
| 11.5.3 Device and network signals . . . . .             | 136        |
| 11.5.4 Behavioral and sequence signals . . . . .        | 136        |
| 11.6 Rule-based systems . . . . .                       | 137        |
| 11.6.1 Examples of useful rules . . . . .               | 137        |
| 11.6.2 Strengths and weaknesses . . . . .               | 137        |
| 11.7 Supervised learning . . . . .                      | 137        |

|                                                             |     |
|-------------------------------------------------------------|-----|
| 11.7.1 Typical model families . . . . .                     | 138 |
| 11.7.2 Why supervised learning helps . . . . .              | 138 |
| 11.7.3 Operational caution . . . . .                        | 138 |
| 11.8 Anomaly detection . . . . .                            | 138 |
| 11.8.1 What anomaly methods look for . . . . .              | 138 |
| 11.8.2 Strengths and weaknesses . . . . .                   | 139 |
| 11.9 Graph-based fraud analysis . . . . .                   | 139 |
| 11.9.1 Why graph structure matters . . . . .                | 139 |
| 11.9.2 Graph features and graph ML . . . . .                | 139 |
| 11.10 Hybrid rules + ML systems . . . . .                   | 140 |
| 11.10.1 Why hybrids work . . . . .                          | 140 |
| 11.10.2 Example decision flow . . . . .                     | 140 |
| 11.11 Thresholds and decision economics . . . . .           | 140 |
| 11.11.1 Approve, review, reject . . . . .                   | 140 |
| 11.11.2 Cost-sensitive tuning . . . . .                     | 141 |
| 11.12 Human-in-the-loop review . . . . .                    | 141 |
| 11.12.1 What reviewers need . . . . .                       | 141 |
| 11.12.2 Feedback loop . . . . .                             | 141 |
| 11.13 Fraud evaluation metrics . . . . .                    | 142 |
| 11.13.1 Useful metrics . . . . .                            | 142 |
| 11.13.2 Business interpretation . . . . .                   | 142 |
| 11.14 Drift, adaptation, and monitoring . . . . .           | 142 |
| 11.14.1 What to monitor . . . . .                           | 142 |
| 11.14.2 Why monitoring is essential . . . . .               | 143 |
| 11.15 Trust & safety beyond payment fraud . . . . .         | 143 |
| 11.15.1 Examples beyond transactions . . . . .              | 143 |
| 11.15.2 Why this matters architecturally . . . . .          | 143 |
| 11.16 Hands-on lab: risk scoring and thresholding . . . . . | 144 |
| 11.16.1 Lab scenario . . . . .                              | 144 |
| 11.16.2 Lab tasks . . . . .                                 | 144 |



|                                                               |            |
|---------------------------------------------------------------|------------|
| 11.16.3 Expected learning outcome . . . . .                   | 144        |
| 11.17 Managerial implications . . . . .                       | 144        |
| 11.18 Multiple-choice questions (MCQs) . . . . .              | 145        |
| 11.19 Exercises (short) . . . . .                             | 146        |
| 11.20 Mini-case (Odoo-linked, vendor-neutral) . . . . .       | 146        |
| <b>12 Customer Service Automation with LLMs (Week 11)</b>     | <b>147</b> |
| 12.1 Learning objectives . . . . .                            | 147        |
| 12.2 Why customer service is a natural LLM use case . . . . . | 147        |
| 12.3 Customer service tasks in e-commerce . . . . .           | 148        |
| 12.3.1 Informational tasks . . . . .                          | 148        |
| 12.3.2 Transactional tasks . . . . .                          | 148        |
| 12.3.3 Judgment-heavy tasks . . . . .                         | 149        |
| 12.4 Rules, classical NLP, and LLMs . . . . .                 | 149        |
| 12.4.1 When rules are best . . . . .                          | 149        |
| 12.4.2 When classical NLP may still be enough . . . . .       | 149        |
| 12.4.3 When LLMs are most valuable . . . . .                  | 150        |
| 12.5 Retrieval-augmented generation (RAG) . . . . .           | 150        |
| 12.5.1 Why RAG matters . . . . .                              | 150        |
| 12.5.2 Knowledge sources . . . . .                            | 150        |
| 12.5.3 Retrieval quality . . . . .                            | 151        |
| 12.6 Tool use and operational actions . . . . .               | 151        |
| 12.6.1 Examples of service tools . . . . .                    | 151        |
| 12.6.2 Why tools are necessary . . . . .                      | 151        |
| 12.6.3 Action control . . . . .                               | 151        |
| 12.7 Workflow orchestration . . . . .                         | 152        |
| 12.7.1 Typical service flow . . . . .                         | 152        |
| 12.7.2 Why orchestration matters . . . . .                    | 152        |
| 12.8 Conversation memory and context . . . . .                | 152        |
| 12.8.1 Useful forms of context . . . . .                      | 153        |

|                                                                    |     |
|--------------------------------------------------------------------|-----|
| 12.8.2 Context hygiene . . . . .                                   | 153 |
| 12.9 Grounding, citations, and answer control . . . . .            | 153 |
| 12.9.1 Why citations help . . . . .                                | 153 |
| 12.9.2 Controlled answer styles . . . . .                          | 154 |
| 12.10 Escalation and human handoff . . . . .                       | 154 |
| 12.10.1 Escalation triggers . . . . .                              | 154 |
| 12.10.2 Good handoff design . . . . .                              | 154 |
| 12.11 Evaluation of LLM service systems . . . . .                  | 155 |
| 12.11.1 Evaluation dimensions . . . . .                            | 155 |
| 12.11.2 Offline versus online evaluation . . . . .                 | 155 |
| 12.11.3 Human review in evaluation . . . . .                       | 156 |
| 12.12 Guardrails and safety . . . . .                              | 156 |
| 12.12.1 Common risk areas . . . . .                                | 156 |
| 12.12.2 Guardrail mechanisms . . . . .                             | 156 |
| 12.13 Prompt injection and retrieval risks . . . . .               | 157 |
| 12.13.1 Examples of injection risk . . . . .                       | 157 |
| 12.13.2 Mitigations . . . . .                                      | 157 |
| 12.14 A practical architecture for LLM support . . . . .           | 157 |
| 12.14.1 Read and write separation . . . . .                        | 158 |
| 12.14.2 Role of CRM and ERP integration . . . . .                  | 158 |
| 12.15 Hands-on lab: a grounded FAQ and service assistant . . . . . | 158 |
| 12.15.1 Lab scenario . . . . .                                     | 158 |
| 12.15.2 Lab tasks . . . . .                                        | 158 |
| 12.15.3 Expected learning outcome . . . . .                        | 159 |
| 12.16 Managerial implications . . . . .                            | 159 |
| 12.17 Multiple-choice questions (MCQs) . . . . .                   | 159 |
| 12.18 Exercises (short) . . . . .                                  | 160 |
| 12.19 Mini-case (Odoo-linked, vendor-neutral) . . . . .            | 161 |

|         |                                                         |     |
|---------|---------------------------------------------------------|-----|
| 13.1    | Learning objectives . . . . .                           | 163 |
| 13.2    | Why MLOps and DataOps matter in e-commerce . . . . .    | 163 |
| 13.3    | From experimentation to production systems . . . . .    | 164 |
| 13.3.1  | The production gap . . . . .                            | 164 |
| 13.3.2  | Why e-commerce is especially operational . . . . .      | 164 |
| 13.4    | What MLOps covers . . . . .                             | 164 |
| 13.4.1  | Typical MLOps responsibilities . . . . .                | 165 |
| 13.4.2  | What DataOps covers . . . . .                           | 165 |
| 13.5    | Reproducibility . . . . .                               | 165 |
| 13.5.1  | What should be reproducible . . . . .                   | 165 |
| 13.5.2  | Why this matters in commerce . . . . .                  | 166 |
| 13.6    | Pipelines and orchestration . . . . .                   | 166 |
| 13.6.1  | Typical pipeline stages . . . . .                       | 166 |
| 13.6.2  | Batch and streaming coexistence . . . . .               | 167 |
| 13.7    | Feature consistency and training-serving skew . . . . . | 167 |
| 13.7.1  | How skew happens . . . . .                              | 167 |
| 13.7.2  | Why skew is dangerous . . . . .                         | 167 |
| 13.7.3  | Feature stores as a concept . . . . .                   | 167 |
| 13.8    | Model registry and release management . . . . .         | 168 |
| 13.8.1  | Why a registry matters . . . . .                        | 168 |
| 13.8.2  | Release strategies . . . . .                            | 168 |
| 13.9    | Monitoring production ML systems . . . . .              | 168 |
| 13.9.1  | Technical monitoring . . . . .                          | 168 |
| 13.9.2  | Model monitoring . . . . .                              | 169 |
| 13.9.3  | Business monitoring . . . . .                           | 169 |
| 13.10   | Data drift and concept drift . . . . .                  | 169 |
| 13.10.1 | Data drift . . . . .                                    | 170 |
| 13.10.2 | Concept drift . . . . .                                 | 170 |
| 13.10.3 | Response to drift . . . . .                             | 170 |
| 13.11   | CI/CD for ML . . . . .                                  | 170 |

|                                                                        |            |
|------------------------------------------------------------------------|------------|
| 13.11.1 What CI for ML should validate . . . . .                       | 171        |
| 13.11.2 What CD for ML should consider . . . . .                       | 171        |
| 13.12 Experiment tracking and model comparison . . . . .               | 171        |
| 13.12.1 What to track . . . . .                                        | 171        |
| 13.12.2 Why this matters . . . . .                                     | 172        |
| 13.13 Governance and approvals . . . . .                               | 172        |
| 13.13.1 Examples of governed models . . . . .                          | 172        |
| 13.13.2 Typical governance artifacts . . . . .                         | 172        |
| 13.14 Incident response for ML systems . . . . .                       | 173        |
| 13.14.1 Examples of ML incidents . . . . .                             | 173        |
| 13.14.2 Preparedness . . . . .                                         | 173        |
| 13.15 Artifacts for this chapter . . . . .                             | 173        |
| 13.16 Hands-on artifacts . . . . .                                     | 174        |
| 13.17 Managerial implications . . . . .                                | 174        |
| 13.18 Multiple-choice questions (MCQs) . . . . .                       | 174        |
| <b>14 Security, Privacy, Compliance, and Responsible AI (Week 13)</b>  | <b>177</b> |
| 14.1 Learning objectives . . . . .                                     | 177        |
| 14.2 Why this chapter matters . . . . .                                | 177        |
| 14.3 Security as a business capability . . . . .                       | 178        |
| 14.4 Threat modeling for e-commerce platforms . . . . .                | 178        |
| 14.4.1 Threat modeling example: intelligent customer support . . . . . | 179        |
| 14.5 Identity, authentication, and access control . . . . .            | 179        |
| 14.6 Payment security and transactional integrity . . . . .            | 180        |
| 14.7 API and integration security . . . . .                            | 180        |
| 14.8 Data protection and privacy-by-design . . . . .                   | 181        |
| 14.8.1 Data minimization in AI systems . . . . .                       | 181        |
| 14.9 Compliance in the commerce environment . . . . .                  | 182        |
| 14.10 Responsible AI in e-commerce . . . . .                           | 182        |
| 14.10.1 High-risk and low-risk AI use cases . . . . .                  | 183        |

|                                                                           |            |
|---------------------------------------------------------------------------|------------|
| 14.11 Bias, proxies, and unintended discrimination . . . . .              | 183        |
| 14.12 Explainability, contestability, and customer trust . . . . .        | 183        |
| 14.13 LLM-specific risks and controls . . . . .                           | 184        |
| 14.14 Logging, monitoring, and auditability . . . . .                     | 185        |
| 14.15 Incident response and crisis handling . . . . .                     | 185        |
| 14.16 Operating model: who owns what . . . . .                            | 186        |
| 14.17 A practical risk assessment template . . . . .                      | 186        |
| 14.18 Managerial implications . . . . .                                   | 187        |
| 14.19 Hands-on lab: assessing an AI returns assistant . . . . .           | 187        |
| 14.20 Multiple-choice questions (MCQs) . . . . .                          | 187        |
| <b>15 Capstone: Enterprise-Grade AI E-Commerce Architecture (Week 14)</b> | <b>191</b> |
| 15.1 Scope (placeholder) . . . . .                                        | 191        |
| 15.2 Capstone deliverables (placeholder) . . . . .                        | 191        |
| <b>A Datasets and Synthetic Data for Teaching (placeholder)</b>           | <b>193</b> |
| <b>B Template: System Design Document (placeholder)</b>                   | <b>195</b> |
| <b>C Template: Data Contract (placeholder)</b>                            | <b>197</b> |
| <b>D Template: Model Card and Risk Assessment (placeholder)</b>           | <b>199</b> |



# How to Use This Draft

This document is a *structure-only* placeholder for an AI-driven, practical Master's-level e-commerce book. Each chapter will be written later; for now, sections contain brief bullets stating scope, key skills, and deliverables.





# Intended Audience and Prerequisites

- Audience: Master's students in Information Systems / Computer & Information Sciences.
- Prereqs (recommended): databases, basic programming, web fundamentals, and introductory ML.
- Tooling (suggested): Python, notebooks, APIs, cloud services, and an ERP/EA modeling tool (as available).



# Hands-on Track (Agreed)

- Approach: both tracks (lighter each).
- Track A (Enterprise/ERP): Odoo-focused integration labs (Sales/Inventory/Accounting flows).
- Track B (AI/ML): ML-heavy notebooks using e-commerce datasets (recsys, pricing, fraud, support).



# Course Outcomes (Draft)

- Design end-to-end e-commerce systems with enterprise architecture (EA) and integration concerns.
- Apply ML/DL/AI to personalization, search, pricing, fraud, operations, and customer service.
- Integrate e-commerce with ERP/CRM/SCM (reference: Odoo) via APIs, events, and modern integration patterns.
- Evaluate systems for security, privacy, governance, ethics, and measurable business value.



# Chapter 1

## Weekly Plan Overview

### Course emphasis (agreed)

- Overall emphasis: balanced (EA/ERP + ML/DL/AI).
- ERP reference stack: Odoo (concepts remain vendor-neutral where possible).

| Week | Topic (maps to the corresponding chapter)                                                  |
|------|--------------------------------------------------------------------------------------------|
| 1    | Foundations of modern e-commerce systems and AI-driven product thinking                    |
| 2    | Digital platforms, marketplaces, and business models (B2C/B2B/D2C, multi-sided platforms)  |
| 3    | E-commerce data foundations: events, tracking, data quality, and experimentation           |
| 4    | Enterprise architecture for e-commerce: capabilities, domains, and reference architectures |
| 5    | Enterprise integration: APIs, iPaaS, ESB, events/streaming, and integration patterns       |
| 6    | ERP/CRM/SCM integration: order-to-cash, procure-to-pay, inventory, and master data         |
| 7    | Recommenders I: classical + learning-to-rank for personalization and search                |
| 8    | Recommenders II: deep learning, embeddings, sequence models, and retrieval                 |
| 9    | Pricing and promotion analytics: forecasting, elasticity, and optimization (ML + OR)       |
| 10   | Fraud, risk, and trust & safety: anomaly detection, graph ML, and AML patterns             |
| 11   | Customer service automation: LLMs, RAG, chatbots, and agentic workflows                    |
| 12   | MLOps + DataOps for e-commerce: deployment, monitoring, drift, and governance              |
| 13   | Security, privacy, compliance, and responsible AI in commerce                              |
| 14   | Capstone: enterprise-grade AI e-commerce solution architecture + evaluation                |





## Chapter 2

# An Introduction to E-Commerce with a Focus on Machine Learning, Deep Learning, and Artificial Intelligence

### Abstract

Electronic commerce (e-commerce) has evolved from static online catalogs into highly dynamic, data-driven ecosystems. This transformation has been powered largely by advances in artificial intelligence (AI), particularly machine learning (ML) and deep learning (DL). This chapter introduces the foundations of e-commerce and explains how AI techniques are embedded across the e-commerce value chain—from customer acquisition and product discovery to pricing, logistics, and fraud prevention. The discussion is aimed at master’s students in computer and information sciences, and therefore emphasizes conceptual clarity, system-level thinking, and the mapping between theoretical models and real-world e-commerce applications.

**Keywords:** e-commerce, artificial intelligence, machine learning, deep learning, recommendation systems, dynamic pricing, fraud detection, personalization, logistics optimization

### 2.1 Introduction to E-Commerce

E-commerce refers to the buying and selling of goods and services via electronic networks, primarily the internet [? ]. It encompasses a broad set of transactional models: business-to-consumer (B2C) online retail, business-to-business (B2B) procurement platforms, consumer-to-consumer (C2C) marketplaces, and consumer-to-business (C2B) models such as influencer platforms and freelance marketplaces [? ].

From a computing perspective, an e-commerce platform is not just a web front-end. It is a socio-technical system composed of:

- User interfaces (web, mobile, conversational)

- Application logic (catalog, cart, checkout, account management)
- Payment and risk engines
- Logistics and fulfillment systems
- Data infrastructure and analytics pipelines
- AI/ML components that drive personalization, prediction, and automation

A key characteristic of e-commerce is its data richness. Every user interaction—page views, searches, clicks, scrolls, wishlists, add-to-carts, purchases, returns, and even time-to-decision—can be captured as fine-grained behavioral logs. Combined with product metadata, prices, promotions, and external contextual data, this creates an ideal environment for AI and ML [? ? ].

### 2.1.1 E-Commerce Business Models and Value Chain

Common e-commerce models include:

- **B2C retail:** Online stores selling directly to consumers.
- **Marketplaces:** Platforms mediating between many sellers and buyers.
- **B2B platforms:** Portals for wholesale ordering, procurement, and supply chain integration.
- **Digital services:** Software subscriptions, digital content, and online education.

Across these models, the **e-commerce value chain** typically includes:

1. Customer acquisition and traffic generation (marketing, ads, SEO)
2. Product discovery and decision support (search, recommendations, reviews)
3. Conversion and transaction processing (checkout, payments, risk checks)
4. Order fulfillment and logistics (inventory, warehousing, routing)
5. Post-purchase engagement (support, returns, loyalty, re-activation)

AI, ML, and DL are now present at each stage of this chain [? ].

## 2.2 Foundations of AI, Machine Learning, and Deep Learning

### 2.2.1 Definitions and Relationships

- **Artificial Intelligence (AI):** Systems that exhibit behaviors considered “intelligent” (perception, reasoning, learning, decision-making).

- **Machine Learning (ML):** Algorithms that learn patterns from data and improve with experience.
- **Deep Learning (DL):** ML methods based on multi-layer neural networks that learn representations from large-scale data.

Thus, deep learning  $\subset$  machine learning  $\subset$  artificial intelligence [? ? ? ].

### 2.2.2 Types of Learning and E-Commerce Examples

1. **Supervised learning:** learn  $f : X \rightarrow Y$  from labeled examples  $(x_i, y_i)$ . Examples: conversion prediction, fraud classification, demand forecasting.
2. **Unsupervised learning:** discover structure in unlabeled data. Examples: customer segmentation, anomaly detection, and learning embeddings from interactions.
3. **Reinforcement learning:** learn a policy to maximize long-term reward. Examples: recommendation policies and dynamic pricing under constraints.
4. **Representation learning:** learn features  $\phi(x)$  automatically (often via DL). Examples: user/item embeddings; Transformer-based session representations.

## 2.3 Machine Learning Across the E-Commerce Lifecycle

### 2.3.1 Personalized Recommendations

Recommendation systems use past behavior and product attributes to suggest items a user is likely to buy [? ? ]. Key evaluation ideas include ranking quality (precision@ $k$ , recall@ $k$ , NDCG) and business impact (incremental revenue) [? ].

### 2.3.2 Search and Ranking

Learning-to-rank uses user feedback (including clicks) to optimize product ranking in response to queries [? ]. Modern systems increasingly incorporate semantic retrieval and NLP components [? ].

### 2.3.3 Segmentation, churn, and CLV

Segmentation (e.g., via clustering) and predictive modeling (e.g., churn/CLV) connect behavioral data to marketing and retention actions [? ].

### 2.3.4 Dynamic Pricing and Promotion Optimization

Pricing combines prediction (demand response) with optimization under constraints; evaluation typically requires controlled online experimentation [? ? ].

### 2.3.5 Demand Forecasting and Inventory Management

Forecasting methods range from statistical time-series models to ML sequence models; forecasts drive replenishment and inventory decisions [? ? ].

### 2.3.6 Fraud Detection and Transaction Security

Fraud detection is an imbalanced classification problem with non-stationarity (attackers adapt), and must balance false positives against fraud loss [? ? ].

### 2.3.7 Customer Service and Conversational Commerce

Customer support automation mixes information retrieval, intent/entity extraction, and increasingly LLM-based interaction [? ].

## 2.4 Deep Learning in E-Commerce

DL is especially useful for unstructured and high-dimensional data (text, images, sequences) [? ].

### 2.4.1 Neural Recommendation Systems

Deep recommender architectures model non-linear user–item interactions and session dynamics [? ? ].

### 2.4.2 NLP for Product Text and Reviews

NLP supports query understanding, semantic search, review analysis, and conversational interfaces [? ].

### 2.4.3 Computer Vision and Visual Search

Vision models support category classification, attribute extraction, and similarity search via embeddings [? ].

## 2.5 Data, Architecture, and Deployment Considerations

E-commerce ML requires robust pipelines (data quality, governance), reliable deployment (latency, availability), and continuous evaluation (monitoring and experiments) [? ].

## 2.6 Challenges and Responsible AI

Key challenges include bias, privacy/security, explainability for high-impact decisions, and the organizational processes needed for safe deployment [? ? ].

## 2.7 Multiple-choice questions (MCQs)

1. Which relationship is correct?
  - (a)  $AI \subset ML \subset DL$
  - (b)  $DL \subset ML \subset AI$
  - (c)  $ML \subset DL \subset AI$
  - (d) AI, ML, and DL are disjoint fields
2. In e-commerce, which is **most** naturally framed as a *ranking* problem?
  - (a) Choosing the best warehouse location for a new facility
  - (b) Ordering products in a search results page for a query
  - (c) Estimating the total demand for next quarter
  - (d) Balancing accounting entries for an invoice
3. Which metric is **most** aligned with offline evaluation of a top- $k$  recommender?
  - (a)  $NDCG@k$
  - (b) Mean squared error (MSE) on prices
  - (c) Page load time (ms)
  - (d) Uptime percentage
4. Fraud detection is often difficult because:
  - (a) Fraud labels are always perfectly accurate
  - (b) The problem is typically extremely imbalanced and attackers adapt over time
  - (c) Fraud has no economic cost, only technical cost
  - (d) The best approach is always rule-based
5. Which statement best describes why A/B testing is important in e-commerce ML?
  - (a) It guarantees the model is unbiased
  - (b) It measures causal impact on business KPIs under real user behavior
  - (c) It replaces the need for offline evaluation entirely
  - (d) It eliminates the need for monitoring after deployment

**Answer key**

1. (b)
2. (b)
3. (a)
4. (b)
5. (b)

## Chapter 3

# Platforms, Marketplaces, and Business Models (Week 2)

### 3.1 Learning objectives

- Distinguish e-commerce *channels* (storefronts) from *platforms* (ecosystems) and *marketplaces* (multi-seller).
- Analyze platform economics: network effects, multi-homing, and pricing of participation.
- Translate business model choices into architectural requirements (identity, trust, data, integration, ERP).

### 3.2 From storefront to platform

- **Storefront:** a single merchant selling to buyers (B2C/B2B) via a web/app channel.
- **Platform:** a system that enables interactions among multiple participant groups (e.g., buyers, sellers, advertisers, logistics providers).
- **Marketplace:** a platform where multiple sellers list items and transactions are mediated by platform policies and infrastructure.

### 3.3 Value creation, value capture, and growth

#### 3.3.1 Value creation and matching

Platforms create value by reducing search and transaction costs, increasing variety, and improving trust; AI frequently strengthens these effects by improving matching and decision support.

### 3.3.2 Value capture models

- Take rate (commission on completed transactions)
- Subscription (seller or buyer membership tiers)
- Listing fees and value-added services (e.g., promotions, analytics)
- Advertising (sponsored listings, retail media)
- Fulfillment fees (warehousing, last-mile delivery, returns handling)

### 3.3.3 Network effects

- **Direct network effects:** more users attract more users (e.g., social commerce communities).
- **Indirect network effects:** more buyers attract more sellers and vice versa (classic marketplace dynamic).
- **Cold start:** bootstrapping supply and demand when network effects are weak.
- **Multi-homing:** sellers/buyers participate in multiple platforms; affects differentiation strategy.

## 3.4 Governance and trust as system requirements

Marketplace governance is not “policy only”: it becomes a set of product and system requirements.

- Seller onboarding, verification, and compliance workflows (KYC/KYB where relevant).
- Catalog integrity: listing policies, moderation, and counterfeit detection.
- Reviews/ratings, dispute resolution, refunds/returns policy enforcement.
- Trust & safety operations: rule+ML detection, case management, and human-in-the-loop escalation.

## 3.5 Business model patterns (and architectural implications)

### 3.5.1 B2C retail

- Key systems: PIM/catalog, promotions, payments, fulfillment, customer support.
- ERP/CRM integration: order-to-cash, inventory, accounting, returns.



### 3.5.2 B2B commerce

- Quotes, negotiated pricing, approval workflows, invoicing, credit limits.
- Integration-heavy: procurement, supplier management, and EDI/API flows.

### 3.5.3 Marketplace

- Split catalog ownership (platform vs sellers), seller tools, payout/settlement.
- Additional “control plane”: seller policies, listing moderation, abuse prevention.

### 3.5.4 Digital services (subscriptions)

- Recurring billing, entitlements, and churn/retention analytics.
- Usage-based pricing requires event design and metering.

## 3.6 Where AI fits (system-level view)

- **Matching layer:** search, recommendations, ranking.
- **Trust layer:** fraud/abuse detection, content moderation.
- **Operations layer:** forecasting, replenishment, routing.
- **Experience layer:** personalization and conversational commerce.

## 3.7 Hands-on (placeholder)

- Create a one-page *platform blueprint*: participant groups, value exchange, pricing model, and required capabilities.
- Map capabilities to an EA view and identify which capabilities are ERP-owned vs commerce-owned.

## 3.8 Case studies (placeholder)

- Modern marketplace patterns (generic; no vendor lock-in).

### 3.9 Multiple-choice questions (MCQs)

1. Which statement best distinguishes a *storefront* from a *marketplace*?
  - (a) A storefront is always B2B; a marketplace is always B2C.
  - (b) A storefront has a single merchant of record; a marketplace mediates transactions among multiple sellers and buyers.
  - (c) A storefront cannot use AI; a marketplace must use AI.
  - (d) A storefront requires ERP integration; a marketplace does not.
2. *Indirect network effects* in marketplaces most directly mean:
  - (a) The platform has a single user group and growth is linear.
  - (b) The value to buyers increases as more sellers join, and the value to sellers increases as more buyers join.
  - (c) Users prefer multiple platforms (multi-homing) regardless of differentiation.
  - (d) The platform must subsidize logistics to succeed.
3. In B2B commerce, which requirement is **most typical** compared to B2C?
  - (a) Session-based recommendations
  - (b) Negotiated pricing and approval workflows
  - (c) Visual search
  - (d) Social reviews and influencer marketing
4. *Multi-homing* refers to:
  - (a) Hosting the platform in multiple cloud regions.
  - (b) Sellers or buyers participating in multiple platforms simultaneously.
  - (c) Using multiple payment gateways for redundancy.
  - (d) Maintaining multiple warehouses for the same SKU.
5. Which is the **best** example of a marketplace *governance* capability?
  - (a) Re-ranking items using embeddings
  - (b) Seller onboarding verification and dispute resolution policies
  - (c) Demand forecasting for replenishment
  - (d) Offline model training with hyperparameter tuning

#### Answer key

1. (b)
2. (b)
3. (b)
4. (b)
5. (b)

## Chapter 4

# E-Commerce Data Foundations: Events, Experimentation, and Metrics (Week 3)

### 4.1 Learning objectives

- Design an event-driven data model for storefront, marketplace, and mobile commerce scenarios.
- Distinguish identifiers, identities, sessions, and entities across operational and analytical systems.
- Explain how attribution, data quality, and metric definitions affect decision-making and ML performance.
- Plan controlled experiments and understand when adaptive approaches such as bandits are appropriate.

### 4.2 Why data foundations matter

AI systems in e-commerce are only as reliable as the data substrate beneath them. Recommendation models, pricing systems, fraud detectors, and executive dashboards all depend on the same underlying facts: who did what, when, in which context, and with what downstream outcome. If those facts are delayed, duplicated, ambiguously defined, or disconnected across systems, both analytics and machine learning become fragile.

This chapter treats data foundations as a *system design* problem rather than a reporting problem. The goal is not merely to “collect more data,” but to create a coherent measurement layer that can support:

- Product analytics and funnel analysis
- Operational reporting across commerce and ERP systems

- Reliable feature generation for ML models
- Controlled online experimentation
- Governance, auditing, and cross-functional alignment

### 4.3 From business process to event model

E-commerce platforms contain several interacting process layers:

- **Customer interaction layer:** impressions, searches, clicks, scrolls, add-to-cart, checkout steps, support interactions.
- **Commercial transaction layer:** order placement, payment authorization, invoicing, shipment, cancellation, return, refund.
- **Operational layer:** inventory adjustments, picking, packing, supplier replenishment, carrier status changes.
- **Experimentation layer:** exposure to treatments, eligibility checks, metric assignment, and holdouts.

Each layer emits events, but those events serve different purposes. A page view event captures human behavior. An order-created event captures a business transaction. A shipment-delivered event captures operational state. An experiment-exposed event captures causal assignment context. Strong data design begins by separating these concerns while still linking them with stable keys.

#### 4.3.1 Events versus state

Students often confuse *events* with *tables of current state*. The distinction is fundamental:

- **State** answers “what is true now?” Examples: current inventory, current product price, current loyalty tier.
- **Events** answer “what happened?” Examples: product viewed, cart updated, payment failed, order shipped.

State is necessary for operations, but events are essential for causality, sequence modeling, root-cause analysis, and reconstructing past behavior. For example, a fraud model may need the sequence of failed payment attempts before an order was placed, not just the final order record. Similarly, retention analysis depends on when a user returned, not merely whether the user is currently “active.”

### 4.3.2 Event taxonomies

A useful event taxonomy groups events into classes that are stable even when the user interface changes. For a commerce platform, a practical taxonomy may include:

- **Discovery events:** homepage viewed, category viewed, search submitted, filter applied, result clicked.
- **Consideration events:** product viewed, review expanded, comparison opened, wishlist updated.
- **Intent events:** add-to-cart, remove-from-cart, coupon applied, shipping option selected.
- **Transaction events:** checkout started, payment submitted, order completed, order canceled.
- **Post-purchase events:** shipment delivered, return requested, refund issued, support ticket opened.

This style of taxonomy makes analytics more robust than mirroring UI element names directly. If a button label changes from “Buy Now” to “Checkout,” the business meaning should still map to a stable event family rather than creating an entirely new measurement definition.

## 4.4 Event schema design

An event is only analytically useful if its schema is consistent. At minimum, most commerce events should include:

- **Event name:** stable business-oriented label such as `product_viewed` or `checkout_started`
- **Event time:** when the event actually occurred
- **Ingestion time:** when the event reached the data platform
- **Actor identifiers:** user ID, guest ID, device ID, session ID, seller ID where relevant
- **Entity identifiers:** product ID, SKU ID, cart ID, order ID, payment ID
- **Context fields:** channel, app version, locale, campaign, experiment assignment, page type
- **Properties:** quantity, price, discount, position in ranking, search query, payment method

### 4.4.1 Required design principles

- **Business semantics before UI semantics:** schemas should reflect domain meaning, not implementation accidents.
- **Stable field names:** renaming fields casually breaks dashboards, pipelines, and models.
- **Explicit units:** store price currency, quantity units, and time zones unambiguously.

- **Immutable raw events:** avoid rewriting the historical log except for tightly governed correction procedures.
- **Versioned schemas:** consumers must know when a producer changed meaning or structure.

#### 4.4.2 Example event record

The following conceptual record shows how a product view may be represented:

```

event_name      = product_viewed
event_time      = 2026-03-23T14:03:11Z
ingestion_time  = 2026-03-23T14:03:13Z
user_id         = null
guest_id        = g_1842
session_id      = s_88fa
product_id      = p_501
sku_id          = sku_501_bl_m
page_type       = product_detail
channel         = mobile_app
campaign_id     = spring_push_7
rank_position   = 4
experiment_id   = rec_widget_v3
variant         = treatment

```

The example illustrates an important practice: authenticated and anonymous identifiers may coexist. This supports both guest behavior analysis and later identity stitching when a guest eventually logs in or purchases.

### 4.5 Identity resolution and entity modeling

One of the hardest problems in e-commerce analytics is deciding what a “customer” means across channels and systems. The same person may appear as:

- an anonymous web browser cookie,
- a mobile-app installation,
- a logged-in storefront account,
- a CRM contact,
- an ERP customer or partner record,
- a marketplace buyer account,
- or multiple records created through operational errors.

### 4.5.1 Identifiers are not interchangeable

Different identifiers serve different roles:

- **User/account ID:** application-level identity for authentication and personalization.
- **Customer/partner ID:** commercial identity used in ERP, billing, and support.
- **Session ID:** short-lived grouping of interactions.
- **Device or browser ID:** technical identifier for anonymous continuity.
- **Order ID:** transaction-level identity.
- **Household or organization ID:** aggregation identity for B2B or family accounts.

If analysts casually join these identifiers as though they were equivalent, metrics become unstable. For example, “active customers” may mean active accounts in one report but active billing entities in another. That inconsistency can distort conversion rates, retention cohorts, and lifetime value estimates.

### 4.5.2 Identity stitching

Identity stitching is the process of linking records that likely belong to the same real-world actor. Common rules include:

- linking a guest cart to a logged-in account after authentication,
- attaching a mobile device history to a user after sign-in,
- associating storefront orders with an ERP customer once order export completes,
- maintaining a separate unresolved pool when confidence is low.

The core design principle is to preserve raw evidence and build linkage logic transparently. Many downstream problems arise when teams overwrite raw IDs with guessed “master” identities and can no longer audit how the match occurred.

### 4.5.3 Privacy-aware identity design

Identity resolution must also respect privacy and compliance constraints. A technically convenient identifier is not automatically a legally or ethically appropriate one. Good practice includes:

- minimizing collection of directly identifying information,
- separating analytical IDs from operational personal data where possible,
- documenting retention periods,
- and ensuring that consent and deletion workflows can propagate to analytical systems.

## 4.6 Data quality as a product requirement

Data quality problems in commerce are rarely abstract. They directly alter revenue reporting, experiment reads, and ML labels. Representative failure modes include:

- **Missing events:** checkout steps are undercounted, making conversion funnels look healthier than reality.
- **Duplicate events:** add-to-cart counts inflate intent metrics and derived model features.
- **Late events:** daily dashboards and training jobs disagree because the data arrive after aggregation windows close.
- **Semantic drift:** an event name remains unchanged while the product team silently changes its meaning.
- **Broken joins:** product IDs in clickstream logs no longer match product IDs in the catalog or ERP.

### 4.6.1 Data contracts

An effective way to reduce these failures is to define *data contracts* between producers and consumers. A data contract specifies:

- event name and schema,
- field definitions and allowed values,
- freshness expectations,
- ownership and escalation path,
- versioning policy,
- and validation rules.

This shifts governance left. Instead of discovering breakage after dashboards fail, producers must acknowledge downstream dependencies before changing event payloads.

### 4.6.2 Validation and observability

Data platforms should monitor more than pipeline uptime. Useful checks include:

- row-count anomalies by event family,
- null-rate spikes in required fields,
- duplicate-rate thresholds by event key,



- cross-system reconciliation between orders, payments, and shipments,
- and schema drift alerts when producers add or remove fields.

In mature teams, these checks are treated as production monitoring. Broken analytics should be considered an operational incident, not a minor inconvenience.

## 4.7 Attribution and its limits

Attribution asks which touchpoints influenced an outcome such as purchase, repeat visit, or subscription renewal. In e-commerce, this sounds straightforward but is often confounded by user behavior, cross-device journeys, channel interactions, and pre-existing intent.

### 4.7.1 Common attribution models

- **Last-touch attribution:** credits the final recorded touchpoint before conversion.
- **First-touch attribution:** credits the earliest touchpoint in the conversion path.
- **Multi-touch heuristics:** divide credit across several interactions.
- **Incrementality-based approaches:** estimate causal lift rather than merely assigning credit.

Heuristic attribution is useful for operational reporting, but it should not be confused with causal inference. A campaign that appears near the end of many journeys may receive last-touch credit even if it did not truly create incremental demand.

### 4.7.2 Selection bias and measurement bias

Attribution systems are affected by:

- **Selection bias:** users exposed to an experience may differ systematically from users who are not.
- **Survivorship bias:** only recorded users or channels appear in the analysis.
- **Cross-device loss:** a user researches on one device and purchases on another without a reliable link.
- **Channel interference:** paid and organic channels interact, making isolated credit assignments misleading.

This is why attribution reporting is a descriptive tool, whereas experimentation is the primary causal tool.

## 4.8 Metrics layers and decision quality

Executives often ask for “the conversion rate” as though it were a single obvious number. In reality, metric definitions encode business rules. Without a shared metrics layer, each team may compute the same KPI differently.

### 4.8.1 North Star and supporting metrics

Many commerce organizations choose a high-level outcome metric such as gross merchandise value, contribution margin, or retained active buyers. That metric must then be decomposed into leading indicators:

- traffic volume,
- product-detail engagement,
- add-to-cart rate,
- checkout completion rate,
- repeat purchase rate,
- return rate,
- and support burden.

For educational purposes, it is useful to connect these metrics to event definitions explicitly. For instance:

$$\text{Add-to-cart rate} = \frac{\text{\#sessions with add\_to\_cart}}{\text{\#sessions with product\_view}}$$

This formula is simple, but each term still depends on sessionization rules, bot filtering, and event validity criteria.

### 4.8.2 Metric dimensions

Commerce metrics usually need to be sliced by:

- channel (web, app, store-assisted, marketplace),
- customer segment (new, returning, VIP, B2B account tier),
- geography and locale,
- device class,
- product category and seller,
- experiment cohort.

A robust metrics layer therefore requires both canonical definitions and dimensional consistency. It should be possible to explain why a number changes, not just to report the number itself.

### 4.8.3 Retention and cohort logic

Retention metrics are especially sensitive to definition. Possible questions include:

- Did the user return to browse?
- Did the user make another purchase?
- Did the organization place another order from any employee account?
- Did the customer remain active despite returns or failed payments?

Each version may be legitimate, but each reflects a different product and business question. Good analytical practice is to name retention metrics with explicit behavioral criteria rather than using ambiguous shorthand.

## 4.9 Experimentation in e-commerce

Offline evaluation is valuable, but it cannot fully predict live business impact. Ranking changes alter user behavior. Promotions affect demand patterns. Fraud rules shift attacker behavior. For these reasons, e-commerce systems rely heavily on online experiments.

### 4.9.1 A/B testing fundamentals

An A/B test compares a treatment against a control by random assignment. The point of randomization is not to guarantee improvement; it is to make the groups comparable in expectation so that observed differences can be interpreted causally.

Key design choices include:

- **Unit of randomization:** user, session, device, seller, or region.
- **Exposure logic:** where and when eligibility is checked.
- **Primary KPI:** the main outcome the experiment is intended to improve.
- **Guardrails:** metrics that protect against harmful side effects such as latency, return rate, or support contacts.
- **Analysis window:** how long outcomes are observed after exposure.

### 4.9.2 Common experimental pitfalls

- **Contamination:** the same user sees both variants through multiple devices or sessions.
- **Novelty effects:** short-term excitement masks weak long-term performance.
- **Interference:** one user's treatment influences another user, seller, or market.

- **Metric leakage:** the treatment changes how events are logged rather than how users behave.
- **Underpowered tests:** sample sizes are too small to detect meaningful differences.

These pitfalls illustrate why experimentation is partly a data-engineering problem. If exposure logs and outcome metrics do not align, the test result is not credible.

### 4.9.3 Bandits and adaptive experimentation

Multi-armed bandits are useful when the platform wants to balance learning with near-term reward. Compared with standard A/B tests, bandits typically allocate more traffic to arms that appear to perform better. This can be attractive for:

- promotions on fast-moving catalogs,
- ranking widgets with many interchangeable treatments,
- high-traffic surfaces where opportunity cost is large.

However, bandits are not a universal replacement for A/B testing. They complicate interpretation, can interact badly with delayed rewards, and are less appropriate when decision-makers need a clean causal estimate for a fixed policy choice. In practice, many organizations use A/B tests for clear comparisons and adaptive methods for optimization on mature surfaces.

## 4.10 From analytics to machine learning features

The same event stream that supports dashboards often becomes the basis for ML features. Examples include:

- recent category views for recommendation,
- failed login or payment attempts for fraud detection,
- coupon usage history for promotion targeting,
- time since last order for churn models,
- return behavior and order defects for seller-quality models.

This creates an important dependency: if event semantics drift, model features drift as well. A model can degrade not only because user behavior changes, but because the measurement pipeline changed underneath it. For this reason, feature definitions should be traceable to governed event definitions.

## 4.11 Hands-on lab: build a minimal metrics layer

The aim of the lab is to help students bridge raw event logs and business-facing KPIs.

### 4.11.1 Lab scenario

Assume a storefront emits five core events: `session_started`, `product_viewed`, `add_to_cart`, `checkout_started`, and `order_completed`. Students are given a small synthetic log containing timestamps, session IDs, customer IDs where available, product IDs, and order IDs.

### 4.11.2 Lab tasks

1. Create a canonical event table with required fields, event validation flags, and ingestion metadata.
2. Build daily metrics for sessions, product views, add-to-cart rate, checkout-start rate, purchase conversion rate, and day-7 repeat purchase.
3. Define at least three data-quality checks and show how failing checks affect one KPI.
4. Create a simple cohort table for first purchase date and compute repeat purchase behavior by cohort.
5. Simulate an experiment exposure field and compute the primary KPI by control versus treatment.

### 4.11.3 Expected learning outcome

By the end of the lab, students should be able to explain why a metric is not just an aggregation, but the result of event design choices, identity assumptions, and validation rules.

## 4.12 Managerial and architectural implications

Data foundations are not owned by analysts alone. They require coordination across product, engineering, data, and enterprise systems teams. Architecturally, the following questions matter:

- Which identifiers are authoritative, and where are they mastered?
- Which events belong in the clickstream versus operational event streams?
- How are storefront events linked to ERP objects such as orders, invoices, and shipments?
- What happens when a schema changes in the storefront while ERP integrations remain stable, or vice versa?
- Who approves metric definitions that drive executive dashboards and model targets?

These are governance questions, but they also influence software boundaries. Weak ownership at this layer creates confusion throughout the commerce stack.

### 4.13 Multiple-choice questions (MCQs)

1. In event-based analytics for e-commerce, which statement is most accurate?
  - (a) Events should be stored without timestamps to save space.
  - (b) Events are most useful when they have a consistent schema, a timestamp, and stable identifiers.
  - (c) Only purchase events matter; click and view events are noise.
  - (d) The best schema changes frequently to match UI updates.
2. Which is an example of a *data quality* problem that can directly harm ML models?
  - (a) A model uses a neural network instead of linear regression.
  - (b) Missing or duplicated events for key actions (e.g., add-to-cart) create biased labels/features.
  - (c) Using SQL instead of Python for ETL.
  - (d) Choosing a smaller batch size during training.
3. Why are offline metrics often insufficient for evaluating an e-commerce ranking model?
  - (a) Offline metrics are always higher than online metrics.
  - (b) Offline metrics cannot measure causal impact under real user behavior and feedback loops.
  - (c) Offline evaluation cannot be computed from logs.
  - (d) Offline metrics are not used in industry.
4. In an A/B test, the primary purpose of randomization is to:
  - (a) Increase model complexity.
  - (b) Ensure the treatment and control groups are comparable in expectation.
  - (c) Eliminate the need for monitoring.
  - (d) Guarantee the KPI improves.
5. A common *identity resolution* challenge in e-commerce is:
  - (a) Mapping product IDs to SKU IDs.
  - (b) Linking the same person across devices/sessions while respecting privacy constraints.
  - (c) Computing NDCG@ $k$  efficiently.
  - (d) Selecting a cloud region.

#### Answer key

1. (b)
2. (b)
3. (b)
4. (b)
5. (b)

## 4.14 Exercises (short)

1. Propose an event taxonomy for a storefront: list at least 10 events (search, view, add-to-cart, checkout steps, purchase, return) and for each event specify 3–5 key fields.
2. Define a North Star metric for an e-commerce product and break it into 3–5 leading indicators. Explain how you would compute each from event logs.
3. Design a simple A/B test for a recommendation widget. Specify: unit of randomization, primary KPI, guardrail metrics, and an example of a bias/confounder to watch for.

## 4.15 Mini-case (Odoo-linked, vendor-neutral)

**Scenario:** A hybrid B2C/B2B company uses Odoo (Sales, Inventory, Accounting, CRM). The storefront and mobile app generate clickstream events.

**Task:** Design a minimal data model that supports both analytics and ML:

- Define how you will link clickstream identities to Odoo customers/partners (when possible) and how you handle guests.
- Propose which entities should have stable IDs across systems (customer, product, order, shipment) and which are channel-specific.
- Specify 5 “data contracts” (producer  $\rightarrow$  consumer) that reduce breakage when the UI or ERP changes.





## Chapter 5

# Enterprise Architecture for E-Commerce (Week 4)

### 5.1 Learning objectives

- Explain how enterprise architecture (EA) helps align business goals, operating models, and technology decisions in e-commerce.
- Use capability maps, value streams, domain boundaries, and system-of-record thinking to structure a commerce platform.
- Distinguish current-state architecture from target architecture and transition planning.
- Analyze how AI capabilities fit into enterprise architecture without turning the architecture into a collection of isolated point solutions.

### 5.2 Why enterprise architecture matters in e-commerce

E-commerce systems evolve quickly. New channels, promotions, AI features, logistics partnerships, and regulatory constraints are introduced faster than many enterprise systems were originally designed to handle. Without architectural discipline, organizations often accumulate fragmented storefronts, duplicated customer records, inconsistent pricing logic, and analytics pipelines that are disconnected from operational truth.

Enterprise architecture is useful because it provides a way to reason across these layers at once. It connects:

- business strategy,
- organizational capabilities,
- application boundaries,
- data ownership,

- integration patterns,
- and technology constraints.

In an AI-driven commerce setting, the architectural question is not only “which model should we use?” It is also:

- Which systems generate the data that the model depends on?
- Which domain owns the business decision that the model influences?
- Which workflow must remain reliable when the model is unavailable or uncertain?
- How are predictions turned into operational actions in storefront, CRM, ERP, or support systems?

### 5.3 Enterprise architecture as a set of views

An architecture description is easier to manage when it is divided into views. For e-commerce, four views are especially practical:

- **Business architecture:** capabilities, value streams, organization roles, policies, and business outcomes.
- **Application architecture:** major systems, services, platform modules, ownership boundaries, and interactions.
- **Data architecture:** entities, identifiers, systems of record, analytical pipelines, and governance rules.
- **Technology architecture:** hosting environment, integration infrastructure, observability stack, identity/security, and platform services.

These views should not be treated as isolated documents. They are projections of one system. For example, a business capability such as pricing must connect to application services, authoritative data sources, and runtime controls.

### 5.4 Capability maps

A capability map describes what an organization must be able to do, independent of how the capability is implemented today. That distinction matters because teams, products, and vendor choices can change while core business abilities remain relatively stable.

### 5.4.1 Core e-commerce capabilities

An AI-enabled commerce organization often needs capabilities such as:

- customer acquisition and campaign orchestration,
- product information and catalog management,
- search, discovery, and merchandising,
- personalization and recommendation,
- pricing, promotions, and offer management,
- cart, checkout, and payment orchestration,
- order management and fulfillment coordination,
- returns, refunds, and post-purchase support,
- customer profile and loyalty management,
- analytics, experimentation, and decision support,
- security, privacy, fraud, and governance.

### 5.4.2 Why capability maps are useful

Capability maps help teams avoid jumping directly into solution design. For example, if an organization says “we need a recommendation engine,” the architectural question should be expanded:

- Is personalization a standalone capability, or is it part of a broader discovery capability?
- Which parts of the recommendation workflow belong to the storefront experience layer, and which belong to a shared data platform?
- Which existing operational capabilities must change to make recommendations effective, such as promotions, inventory awareness, or content quality?

Capability maps also expose imbalance. A company may invest heavily in customer acquisition and storefront design while underinvesting in returns, data governance, or seller operations. Architecture should reveal such weaknesses before they become scaling problems.

## 5.5 Value streams

Capabilities describe *what* the organization can do. Value streams describe *how value moves* through the business from trigger to outcome.

### 5.5.1 Typical e-commerce value streams

Important value streams include:

- **Browse-to-buy:** discovery, consideration, cart, checkout, payment, confirmation.
- **Order-to-cash:** order capture, validation, fulfillment, invoicing, settlement, financial posting.
- **Procure-to-stock:** replenishment planning, supplier ordering, receiving, storage, inventory updates.
- **Return-to-resolution:** return request, approval, pickup/drop-off, inspection, refund or exchange.
- **Lead-to-account:** acquisition, registration, qualification, first purchase, retention.

### 5.5.2 Why value streams matter architecturally

Value streams expose cross-domain handoffs. For example, “browse-to-buy” may touch search, catalog, pricing, promotions, checkout, payment, fraud, order management, and analytics. If the stream depends on unclear ownership or duplicated business rules, failures appear at handoff points rather than inside a single system.

Value-stream mapping is especially valuable in hybrid commerce environments where storefront systems and ERP systems divide responsibility. Students should learn to ask not only which system executes a step, but which system is accountable for the business meaning of that step.

## 5.6 Domain boundaries and bounded contexts

Large commerce platforms become difficult to scale when everything is treated as one application or one shared database. Architecture therefore needs domain boundaries with explicit ownership.

### 5.6.1 Typical commerce domains

Common domains include:

- **Catalog:** products, attributes, categories, media, content enrichment.
- **Pricing and promotions:** price lists, discount rules, coupon logic, campaign offers.
- **Inventory and availability:** stock position, reservations, safety stock, availability promises.
- **Cart and checkout:** session basket state, shipping options, tax estimation, checkout workflow.
- **Payments:** payment intents, authorization, capture, refund orchestration.
- **Orders:** order lifecycle, status management, amendments, cancellations.

- **Customer and identity:** profile, authentication, preferences, loyalty, segmentation.
- **Fulfillment and logistics:** warehouse tasks, shipment creation, delivery events, returns routing.
- **Analytics and experimentation:** event collection, metric definitions, experiment assignment.

### 5.6.2 Good boundaries versus bad boundaries

A good boundary reduces coupling and clarifies ownership of behavior and data. A poor boundary creates ambiguity such as:

- pricing rules duplicated in storefront, ERP, and marketing tools,
- customer attributes updated independently in CRM and commerce applications,
- order status interpreted differently by warehouse, finance, and support teams,
- product IDs that differ between catalog, analytics, and inventory systems.

Architectural boundaries should therefore be tested against real business scenarios. If a business change requires editing logic in five systems with no clear owner, the boundary is likely wrong.

## 5.7 Systems of record and systems of engagement

Not every system should be authoritative for every type of data. One of the most important EA concepts in commerce is deciding the *system of record* for each key entity.

### 5.7.1 Typical patterns

- The **commerce layer** often owns high-velocity interaction data such as carts, storefront sessions, and recommendation exposures.
- The **ERP layer** often owns invoices, financial postings, stock movements, supplier records, and contractual B2B customer data.
- A **PIM or catalog service** may own enriched product content while ERP owns procurement-oriented product master fields.
- A **data platform** may own analytical aggregates and features but should not silently become the operational source of truth.

### 5.7.2 Systems of engagement

Storefronts, mobile apps, support consoles, and seller portals are often systems of engagement rather than systems of record. They present, collect, and orchestrate information, but they may not be authoritative for final commercial truth. Confusing these roles leads to reconciliation failures.

For example, a storefront may display an order as “confirmed” while ERP has not yet accepted it due to credit rules or stock constraints. Architecture must define which state is customer-facing, which state is financially binding, and how transitions are synchronized.

## 5.8 Reference architecture for AI-driven e-commerce

There is no single universal architecture, but a practical reference pattern often includes the following layers:

- **Experience layer:** web, mobile, agent interfaces, seller/admin portals.
- **Commerce services layer:** catalog, search, pricing, cart, checkout, order orchestration, customer services.
- **Enterprise systems layer:** ERP, CRM, WMS, finance, procurement, supplier management.
- **Integration layer:** APIs, event streams, workflow orchestration, master-data synchronization.
- **Data and AI layer:** event collection, warehouse/lakehouse, feature pipelines, experimentation, model training, model serving.
- **Cross-cutting controls:** identity, security, privacy, audit, observability, governance.

### 5.8.1 Where AI fits

AI should be mapped to business capabilities rather than inserted as an isolated layer that “does intelligence.” Examples include:

- search ranking and recommendations within discovery,
- fraud scoring within checkout and payments,
- demand forecasting within inventory and replenishment,
- case summarization and response assistance within support,
- anomaly detection within operations and finance.

This mapping matters because it keeps accountability clear. The discovery domain owns business outcomes for search and recommendation even if the underlying models are developed by a centralized data team.

## 5.9 Current state, target state, and transition architecture

Architecture is not useful if it only documents the present. Organizations need a target state that guides investment and sequencing.

### 5.9.1 Current-state analysis

Current-state architecture should identify:

- duplicated systems and overlapping responsibilities,
- brittle integrations and manual reconciliation steps,
- data ownership conflicts,
- missing capabilities,
- and operational pain points such as latency, outages, or inconsistent KPIs.

### 5.9.2 Target architecture

Target architecture is a future-state design that expresses intended business and technical structure. It is not a vendor brochure and not a line-by-line implementation plan. It should answer:

- Which domains and systems will own which capabilities?
- Which integrations will be synchronous, asynchronous, or batch?
- Which entities will be mastered where?
- Which platforms are shared across business units?
- Which AI capabilities will be productized versus embedded locally?

### 5.9.3 Transition architecture

Most organizations cannot move directly from current state to target state. Transition architecture defines intermediate steps. Examples include:

- separating catalog and pricing ownership before replacing checkout,
- introducing a shared event stream before centralizing experimentation,
- establishing a customer identity service before rolling out cross-channel personalization.

This is important in enterprise settings because a theoretically elegant target architecture may fail if the migration path is not operationally realistic.

## 5.10 Architectural concerns specific to AI

AI adds architectural concerns beyond model selection. Representative concerns include:

- **Training-serving consistency:** features must mean the same thing offline and online.
- **Latency budgets:** recommendation or fraud decisions may need sub-second responses.
- **Fallback behavior:** the commerce workflow must continue safely if a model is unavailable.
- **Observability:** systems must track not only uptime but also prediction quality, drift, and business impact.
- **Governance:** model changes may alter business policy and therefore require approval, auditability, and rollback procedures.

In many organizations, the failure is not that models are weak. The failure is that model outputs are not architecturally integrated into business workflows with clear ownership and control points.

## 5.11 Operating model and team topology

Architecture is influenced by team structure. A domain with clear technical boundaries but no accountable product owner often performs badly in practice.

### 5.11.1 Typical team patterns

- **Domain-aligned product teams:** own specific business capabilities such as checkout or catalog.
- **Platform teams:** provide shared infrastructure such as identity, observability, experimentation, or data platforms.
- **Enterprise systems teams:** maintain ERP, CRM, accounting, and operational master-data platforms.
- **Central AI or data teams:** develop shared methods, governance, and model platforms while collaborating with domain teams on use cases.

An architecture is more likely to succeed when ownership in the diagrams matches ownership in the organization. If a system is “owned by everyone,” it is usually owned by no one.

## 5.12 Artifacts for students and architects

A minimal EA artifact set for an e-commerce program may include:



- a capability map,
- one or more value stream maps,
- a domain map with ownership boundaries,
- a system-context diagram,
- a system-of-record matrix for key entities,
- and a target-state architecture note with transition priorities.

These artifacts should be consistent with each other. If the domain map says pricing is a shared domain but the system-of-record matrix shows three authoritative pricing systems, the architecture still contains a contradiction.

### 5.13 Hands-on artifacts

The key deliverable for this chapter is a compact but defensible architecture package:

- an e-commerce capability map,
- a domain model showing clear ownership,
- and a short target-architecture note that explains systems of record and integration choices.

### 5.14 Managerial implications

Enterprise architecture is sometimes dismissed as documentation overhead. In e-commerce, that is a mistake. Architectural ambiguity produces real costs:

- slow product delivery because every change triggers cross-team negotiation,
- financial reconciliation issues from conflicting commercial truth,
- duplicated AI initiatives that solve the same problem inconsistently,
- and weak governance when no domain clearly owns a high-impact decision.

Good architecture does not eliminate complexity. It makes complexity explicit, owned, and therefore manageable.

### 5.15 Multiple-choice questions (MCQs)

1. In enterprise architecture, a *capability map* primarily describes:
  - (a) A list of specific microservices and their endpoints.
  - (b) *What* the business does (stable abilities), independent of *how* it is implemented.
  - (c) The physical network topology of the data center.
  - (d) A Gantt chart of the project plan.
2. A good domain boundary (e.g., for a “Catalog” domain) typically aims to:
  - (a) Maximize shared database tables across all teams.
  - (b) Minimize coupling and define clear ownership of data and behaviors.
  - (c) Ensure all operations are synchronous.
  - (d) Avoid having APIs.
3. Which artifact is most suitable for describing how business value is delivered end-to-end?
  - (a) Value stream map
  - (b) Entity-relationship diagram (ERD)
  - (c) Source code repository layout
  - (d) TLS configuration file
4. “Target architecture” is best described as:
  - (a) The current-state system diagram.
  - (b) A future-state design that guides decisions and sequencing of change.
  - (c) A vendor product brochure.
  - (d) A test plan for unit tests.
5. In an AI-driven e-commerce context, which is the best example of an *architectural concern* (not a model choice)?
  - (a) Whether to use XGBoost or logistic regression
  - (b) How to ensure training-serving consistency and monitor drift
  - (c) Whether to use  $k$ -means or DBSCAN
  - (d) Whether to use SGD or Adam

#### Answer key

1. (b)
2. (b)
3. (a)
4. (b)
5. (b)

## 5.16 Exercises (short)

1. Draft a capability map for an AI-enabled e-commerce organization. Include at least: customer acquisition, discovery, pricing/promotions, order management, fulfillment/returns, customer support, data/analytics, and governance/security.
2. Choose one value stream (e.g., “order-to-cash” or “browse-to-buy”). Identify 5–8 steps and list the owning domain/system for each step (commerce vs ERP vs logistics vs payment).
3. Propose domain boundaries for: catalog, pricing, orders, payments, and customer profiles. For each boundary, name the system-of-record for key data entities.

## 5.17 Mini-case (Odoo-linked, vendor-neutral)

**Scenario:** Your company is hybrid B2C/B2B. Odoo is used for core ERP flows (Sales, Inventory, Accounting, CRM). A separate commerce layer handles web/mobile experiences and AI features.

**Task:** Produce a short “target architecture” note:

- Draw a capability-to-system mapping (capabilities → commerce layer vs Odoo vs data platform).
- Define the *system of record* for customer, product, price list, order, invoice, and shipment.
- Identify 3 integration risks (data consistency, latency, duplicate sources of truth) and propose mitigations.



## Chapter 6

# Enterprise Integration Patterns for E-Commerce (Week 5)

### 6.1 Learning objectives

- Compare synchronous APIs, asynchronous events, and batch integrations in enterprise commerce environments.
- Explain why reliability patterns such as idempotency, retries, dead-letter handling, and the outbox pattern are essential in distributed commerce systems.
- Design integration contracts that reduce coupling between storefront, ERP, payment, logistics, and data platforms.
- Evaluate integration choices in terms of latency, resilience, consistency, and operational complexity.

### 6.2 Why integration patterns matter

E-commerce systems rarely operate as a single application. Even a modest commerce stack may involve a storefront, search service, promotions engine, payment provider, ERP, warehouse system, shipping partner, CRM, and analytics platform. These systems must exchange data continuously while still tolerating delays, retries, outages, and partial failures.

The architectural problem is not merely how to “connect systems.” The deeper question is how to coordinate business processes when no single runtime boundary contains the entire workflow. An order may be accepted in the storefront, authorized by a payment gateway, validated by fraud controls, persisted in commerce services, exported to ERP, and later updated by warehouse and carrier events. Each handoff introduces failure modes that integration design must anticipate.

## 6.3 Integration styles

Integration patterns can be grouped into a small number of common styles.

### 6.3.1 Synchronous request-response

In synchronous integration, one system calls another and waits for a response. Typical examples include:

- payment authorization during checkout,
- tax estimation,
- retrieving customer-specific pricing,
- checking address validation services,
- fetching live delivery options.

The main advantage is immediacy. The caller receives a result within the user flow and can make a direct decision. The main disadvantage is coupling. The caller now depends on the callee's availability, latency, contract stability, and error behavior.

### 6.3.2 Asynchronous events

In asynchronous integration, a producer publishes an event and does not wait for every downstream consumer to finish. Examples include:

- `OrderCreated`,
- `InvoicePosted`,
- `ShipmentDispatched`,
- `InventoryAdjusted`,
- `CustomerRegistered`.

This style supports loose coupling and fan-out. Multiple consumers can react to the same business event without the producer needing to know all of them in advance. However, event-driven systems trade immediacy for eventual consistency and demand stronger operational discipline.

### 6.3.3 Batch and file-based integration

Batch integration remains common in enterprise settings despite modern API and event tooling. Examples include nightly product exports, finance reconciliation files, bulk supplier updates, and large historical loads into analytical platforms.

Batch is often appropriate when:

- latency requirements are low,
- partner systems are legacy or externally controlled,
- very large data volumes must be moved efficiently,
- or the business process is administrative rather than interactive.

Students should not assume batch is “obsolete.” It is often the right pattern for the wrong place if teams are not explicit about where it belongs.

## 6.4 Choosing the right integration style

No single integration style is universally best. The right choice depends on business semantics.

### 6.4.1 Questions to ask

When selecting a pattern, architects should ask:

- Does the user or operator need an immediate answer?
- What is the tolerance for stale data or delayed propagation?
- What happens if the downstream system is unavailable?
- Can the workflow continue in a degraded mode?
- Is the interaction command-like, query-like, or event-like?
- How many downstream consumers must react to the information?

### 6.4.2 Typical choices in commerce

- **Checkout payment authorization:** usually synchronous because the customer needs immediate confirmation.
- **Order export to ERP:** often asynchronous because the storefront should not block on internal back-office processing.
- **Inventory availability updates:** often event-driven with cache refreshes and reconciliation support.
- **Financial reporting loads:** often batch.
- **Shipment tracking updates:** usually asynchronous because carriers and fulfillment processes are naturally event-driven.

## 6.5 API-first integration

APIs are essential in modern commerce because they define explicit service boundaries. An API-first mindset means the contract is treated as a product artifact rather than an afterthought.

### 6.5.1 Good API design principles

- Model APIs around business resources and actions, not internal database tables.
- Make required fields, validation rules, and error semantics explicit.
- Use stable identifiers and versioning policies.
- Design for failure and retries rather than assuming a perfect network.
- Document authorization requirements and rate limits.

### 6.5.2 Commands versus queries

Many integration problems become clearer when teams distinguish:

- **Queries:** read current information, such as “get available shipping options.”
- **Commands:** request a state change, such as “create order” or “capture payment.”

Commands are especially sensitive to retry behavior. If a request is repeated because of a timeout, the system must avoid creating duplicate side effects. This leads directly to idempotency design.

## 6.6 Idempotency

Idempotency means that repeating the same operation does not create additional unintended effects. This is one of the most important patterns in distributed order processing.

### 6.6.1 Why idempotency is required

Networks fail in ambiguous ways. A caller may send a command, the server may execute it, and the response may be lost. If the caller retries without an idempotency mechanism, duplicate orders, duplicate refunds, or duplicate invoice creation may occur.

### 6.6.2 Idempotency keys

For high-value commands such as order creation or payment capture, the client or upstream service often sends an idempotency key. The receiving service stores the key with the outcome of the first successful execution and returns the same result for subsequent retries with the same key.



Good practice includes:

- making the key stable for a business attempt,
- scoping it to the operation type,
- storing it long enough to cover realistic retry windows,
- and linking it to request payload validation to prevent misuse.

## 6.7 Events and event-driven architecture

An event represents something that already happened in the business domain. Examples include order creation, payment success, stock decrement, and shipment delivery.

### 6.7.1 Domain events versus technical events

Teams should distinguish:

- **Domain events:** meaningful business facts such as `OrderCanceled`.
- **Technical events:** operational notifications such as cache refresh or job completion.

Domain events are more reusable because multiple consumers can act on them without knowing the producer's implementation details. If an event stream is dominated by low-level technical signals, consumers become tightly coupled to internals rather than business meaning.

### 6.7.2 Benefits of event-driven integration

- Producers and consumers can evolve more independently.
- Multiple downstream systems can react to the same fact.
- Workloads can be buffered and retried.
- Slow consumers do not necessarily block user-facing systems.

### 6.7.3 Costs of event-driven integration

- Debugging becomes harder because workflows are distributed across time and services.
- Consumers must tolerate reordering, duplication, and replay.
- Monitoring and tracing become mandatory rather than optional.
- Business users must understand eventual consistency.

## 6.8 Reliability patterns in distributed workflows

Integration reliability is not achieved by a single technology choice. It is achieved by combining patterns that reduce ambiguity under failure.

### 6.8.1 Retries and backoff

Transient failures are normal. Consumers and callers should usually retry with bounded exponential backoff rather than failing immediately. However, retries without idempotency or circuit breaking can amplify incidents.

### 6.8.2 Dead-letter handling

If a message repeatedly fails processing, it should not block the entire flow indefinitely. Dead-letter queues or equivalent quarantine mechanisms allow operators to inspect problematic messages, repair data if needed, and replay safely.

### 6.8.3 Timeouts and circuit breakers

Synchronous integrations should use explicit timeouts. If a downstream system becomes slow or unavailable, circuit breakers can stop a cascade of blocked requests. In commerce, this is often paired with degraded behavior such as:

- temporarily disabling a non-critical recommendation widget,
- falling back to cached delivery promises,
- or routing support cases for manual review.

### 6.8.4 Reconciliation jobs

Not all inconsistencies can be prevented in real time. Periodic reconciliation is a practical safeguard, especially for orders, invoices, stock, and refunds. Reconciliation does not replace good design, but it provides a controlled recovery mechanism when distributed flows drift apart.

## 6.9 The outbox pattern

One of the most common distributed data bugs is the dual-write problem. For example, a service may write a new order to its database and then publish an event. If the database write succeeds but the event publish fails, downstream consumers never learn about the order. If the event is published but the database write fails, consumers act on a fact that is not actually committed.

### 6.9.1 How the outbox pattern works

The outbox pattern addresses this by:

- writing the business state change and an “outbox record” in the same local transaction,
- then having a separate relay process publish the outbox record to the event broker,
- marking the record as sent only after successful publish.

This pattern does not magically create exactly-once delivery across the whole system. What it does is make producer-side event publication reliable relative to the local transaction boundary.

### 6.9.2 Why it matters in commerce

Commerce workflows are especially sensitive to dual-write failures because duplicated or missing side effects can affect revenue, customer trust, and financial records. Orders, refunds, shipment creation, and invoice triggers are all candidates for outbox-style publishing.

## 6.10 Message contracts and schema governance

Events and APIs require contracts. Without governed contracts, integration platforms become brittle collections of undocumented assumptions.

### 6.10.1 What a contract should define

A useful contract specifies:

- message or API name,
- required and optional fields,
- field semantics and units,
- identifier rules,
- allowed status transitions,
- versioning and deprecation policy,
- ownership and escalation path.

### 6.10.2 Backward compatibility

Contract evolution should aim for backward compatibility whenever possible. Common practices include:

- adding optional fields rather than repurposing existing ones,
- preserving existing semantics across versions,
- introducing new event names or explicit versions for breaking changes,
- and validating consumer readiness before removing old fields.

In enterprise commerce, breaking contract changes are expensive because downstream systems often include ERP integrations, partner processes, and compliance reporting.

## 6.11 Observability for integrations

Integration systems need observability that spans services and transport layers. Logs alone are insufficient.

### 6.11.1 Minimum signals

An integration operating model should include:

- **Logs:** structured records with correlation IDs, idempotency keys, and error codes.
- **Metrics:** request rates, latency, failure rates, retry counts, dead-letter counts, consumer lag.
- **Traces:** end-to-end flow visibility across API calls, event publication, and downstream processing.
- **Business counters:** order export success, invoice posting delay, shipment-update freshness, reconciliation mismatch rates.

### 6.11.2 Correlation and traceability

Cross-system workflows should carry shared identifiers such as order ID, correlation ID, and experiment or campaign context where relevant. Without this, incident response becomes guesswork. In distributed commerce systems, the ability to answer “what happened to order X?” is an operational requirement, not a nice-to-have feature.

## 6.12 iPaaS, ESB, brokers, and workflow orchestration

Enterprise integration discussions often involve platform choices such as ESBs, iPaaS products, event brokers, and orchestration tools. Students should understand their roles conceptually rather than treating them as interchangeable buzzwords.

### 6.12.1 iPaaS and ESB concepts

- **ESB-style platforms:** historically emphasized centralized mediation, transformation, and routing.
- **iPaaS platforms:** emphasize managed connectors, workflow automation, and cloud integration convenience.

These tools can accelerate delivery, especially where enterprise applications and partner ecosystems are involved. But they can also become hidden concentrations of logic if every business rule is embedded in a central integration layer.

### 6.12.2 Event brokers and streams

Message brokers and streaming platforms provide transport for asynchronous communication. Their value lies not just in moving messages, but in enabling buffering, replay, scaling, and decoupling. However, the transport does not solve data modeling or ownership problems by itself.

### 6.12.3 Workflow orchestration

Some workflows are easier to manage with explicit orchestration, especially long-running processes such as returns handling, supplier exceptions, or manual fraud review. Orchestration can improve clarity, but over-centralization can also produce a fragile control tower if every process depends on one orchestrator.

## 6.13 Integration anti-patterns

Common anti-patterns in commerce integration include:

- direct database-to-database coupling across domains,
- duplicating business rules independently in multiple systems,
- publishing events without stable schemas or owners,
- assuming exactly-once delivery where only at-least-once semantics are realistic,
- using synchronous calls for every workflow step,

- and treating reconciliation as proof that real-time design can be ignored.

Students should be able to identify these anti-patterns because they appear frequently in systems that “worked at small scale” but fail under growth.

## 6.14 Hands-on lab: modeling an order lifecycle

The practical goal of this chapter is to connect abstract patterns to one realistic workflow.

### 6.14.1 Lab scenario

Assume a headless storefront, a payment service, an order service, Odoo, and a shipment service. Students are asked to model the lifecycle from checkout submission to shipment dispatch.

### 6.14.2 Lab tasks

1. Identify which interactions are synchronous and which are asynchronous.
2. Define an idempotency strategy for order creation and payment capture.
3. Design an `OrderCreated` event contract and one downstream consumer.
4. Show where the outbox pattern would be applied.
5. Define the operational metrics needed to detect broken flows.

### 6.14.3 Expected learning outcome

By the end of the lab, students should be able to explain integration design in terms of business guarantees rather than technology labels alone.

## 6.15 Managerial and architectural implications

Integration choices shape business agility. If every new channel or AI feature requires fragile point-to-point connections, innovation slows. If teams overuse asynchronous flows without good observability, the organization loses operational clarity.

Architecture therefore requires a balanced approach:

- synchronous calls where immediate business decisions are required,
- events where decoupling and fan-out matter,
- batch where latency is not critical,

- and governance across all three.

Well-designed integration is not invisible plumbing. It is a core business capability in enterprise commerce.

## 6.16 Multiple-choice questions (MCQs)

1. Which statement best characterizes synchronous vs asynchronous integration?
  - (a) Synchronous integration cannot fail; asynchronous always fails.
  - (b) Synchronous integration couples availability and latency; asynchronous trades immediacy for resilience and decoupling.
  - (c) Asynchronous integration is only for analytics; synchronous is only for ERP.
  - (d) They are equivalent as long as JSON is used.
2. Why is idempotency important in order processing APIs?
  - (a) It makes responses larger.
  - (b) It prevents duplicate effects when requests are retried.
  - (c) It eliminates the need for authentication.
  - (d) It guarantees exactly-once message delivery in all systems.
3. The *outbox pattern* is primarily used to:
  - (a) Store images for the catalog.
  - (b) Reliably publish events when database writes succeed, avoiding dual-write inconsistencies.
  - (c) Encrypt API traffic.
  - (d) Replace a message broker.
4. In event-driven systems, a *consumer* should typically be designed to be:
  - (a) State-less but non-retryable
  - (b) Idempotent and retryable
  - (c) Dependent on message ordering always being perfect
  - (d) Tightly coupled to producer database schemas
5. Which is a common reason to choose events over direct API calls for some flows?
  - (a) To avoid defining schemas/contracts
  - (b) To decouple systems and allow multiple downstream consumers without changing the producer
  - (c) To guarantee zero latency
  - (d) To avoid monitoring

**Answer key**

1. (b)
2. (b)
3. (b)
4. (b)
5. (b)

**6.17 Exercises (short)**

1. For each flow, choose sync API or async events and justify: (i) “place order”, (ii) “update inventory availability”, (iii) “invoice posted”, (iv) “shipment delivered”, (v) “customer profile updated”.
2. Design an idempotency strategy for the “create order” API. Specify the idempotency key and how long it is retained.
3. Write a short failure story: a dual-write bug between database and message broker. Explain how the outbox pattern prevents it.

**6.18 Mini-case (Odoo-linked, vendor-neutral)**

**Scenario:** A headless storefront integrates with Odoo for orders, inventory, invoices, and CRM updates.

**Task:** Propose an integration design that supports scale and reliability:

- Define the key domain events (at least 8) and which system publishes each.
- Identify where you need synchronous APIs (e.g., payment authorization) vs asynchronous events (e.g., shipment updates).
- Specify two data contracts and a versioning strategy to avoid breaking changes.
- List the minimum observability signals (logs/metrics/traces) required to operate the integration in production.



## Chapter 7

# ERP/CRM/SCM Integration in Commerce (Week 6)

### 7.1 Learning objectives

- Explain how commerce platforms integrate with ERP, CRM, and supply-chain systems across operational workflows.
- Identify systems of record for key entities such as customer, product, price list, order, invoice, and shipment.
- Analyze master data management, process ownership, and consistency risks in hybrid commerce architectures.
- Design integration flows for order-to-cash, returns, and B2B pricing using Odoo as a reference ERP environment.

### 7.2 Why ERP/CRM/SCM integration is central to commerce

Modern e-commerce experiences are often built in specialized commerce platforms, but commercial truth does not live only in the storefront. Financial commitments, stock movements, supplier processes, customer account structures, invoicing, and back-office controls typically depend on enterprise systems such as ERP, CRM, and warehouse or supply-chain platforms.

This means a commerce platform cannot be architected only as a customer-facing product. It must also participate in enterprise operating flows. If that integration is weak, the result is familiar:

- orders that appear confirmed online but fail in back-office processing,
- inventory figures that differ between storefront and warehouse,
- duplicate customer records across channels,
- B2B price lists that drift between sales and commerce systems,

- returns and refunds that create reconciliation problems in finance.

The core architectural task is not simply to “connect Odoo” or any ERP. It is to define ownership, process boundaries, and trustworthy data movement across the enterprise stack.

## 7.3 The enterprise systems perspective

In a hybrid commerce architecture, enterprise systems typically serve different roles.

### 7.3.1 ERP

ERP systems coordinate structured operational and financial processes. Using Odoo as a reference point, relevant modules often include:

- **Sales:** quotations, sales orders, customer terms, sales workflows,
- **Inventory:** stock moves, reservations, warehouse operations, replenishment,
- **Accounting:** invoices, credit notes, payments, reconciliation,
- **CRM:** leads, opportunities, account history, commercial context.

### 7.3.2 CRM

CRM systems focus on relationship history, account management, sales interactions, lead qualification, and customer-service context. In some organizations CRM and ERP overlap. What matters architecturally is not the product label but the business responsibility assigned to the system.

### 7.3.3 SCM and warehouse systems

Supply-chain and warehouse platforms manage sourcing, supplier coordination, inbound logistics, warehouse execution, and delivery orchestration. In smaller environments these functions may be embedded in ERP. In larger environments they may be split into specialized systems such as WMS, TMS, or procurement platforms.

## 7.4 Commerce versus enterprise ownership

A recurring source of failure in commerce architecture is treating ownership as a technical detail rather than a business decision. Teams must decide which domain owns each important entity and process state.

### 7.4.1 Common ownership patterns

- The **commerce layer** often owns web sessions, carts, checkout state, personalization context, and channel-specific merchandising behavior.
- The **ERP layer** often owns invoices, accounting entries, stock movements, contractual customer records, and internal fulfillment commitments.
- The **CRM layer** may own lead status, sales activity, account segmentation, and service interactions.
- The **data platform** may own analytical aggregates and experimentation outputs but should not silently become the master for operational entities.

### 7.4.2 Why ownership must be explicit

If pricing is editable in both storefront and ERP, or if both commerce and CRM create customer records independently, inconsistencies are not accidental. They are structural. Integration design cannot fix a domain-ownership problem after the fact; it must start by making ownership explicit.

## 7.5 Systems of record for key entities

One of the most important artifacts in enterprise commerce architecture is a system-of-record matrix. Representative entity choices include the following.

### 7.5.1 Customer

Customer identity is often split across systems:

- the storefront may own channel account credentials and preferences,
- ERP may own bill-to and ship-to parties for financial processing,
- CRM may own relationship context, lead lifecycle, and account history.

The correct design is usually not “one system owns everything.” Instead, different aspects of customer identity are mastered in different places and linked through governed identifiers.

### 7.5.2 Product

Product ownership is also frequently split:

- ERP may own procurement and internal stock-keeping information,
- a catalog or PIM layer may own rich digital merchandising content,

- analytics may derive behavioral product features.

The architecture must distinguish the commercial product definition from channel presentation and from analytical enrichment.

### 7.5.3 Price list

Price lists are especially sensitive in B2B environments. Negotiated terms, tax logic, discount hierarchies, and approval rules often belong closer to ERP or enterprise sales processes than to storefront presentation logic. If the storefront computes pricing independently while ERP maintains contractual prices, disputes are inevitable.

### 7.5.4 Order

The commerce platform often captures the initial order intent, but the enterprise question is whether ERP becomes the authoritative record for downstream operational and financial lifecycle. This must be decided explicitly. Some organizations treat commerce as the order system of record and export to ERP. Others treat the storefront order as a pre-order until ERP accepts and materializes it.

### 7.5.5 Invoice and payment reconciliation

Invoices, ledger effects, and payment reconciliation usually belong in ERP or finance-owned systems. Even if the storefront shows payment confirmation, the enterprise system remains authoritative for accounting truth.

### 7.5.6 Shipment

Shipment truth may belong in warehouse or logistics systems even if ERP or storefront mirrors the status. The architecture should distinguish customer-facing delivery status from the internal system that owns shipment execution and exception handling.

## 7.6 Master data management (MDM)

Master data management is the discipline of defining and governing shared core entities across systems. In commerce, the most common shared entities include customer, product, supplier, location, and price list.

### 7.6.1 Why MDM matters

Without MDM, organizations accumulate:

- duplicate customer records created by multiple channels,

- product records with inconsistent category or attribute definitions,
- seller or supplier entities that differ across procurement and commerce tools,
- incompatible codes used by analytics, ERP, and storefront systems.

### 7.6.2 Practical MDM principles

- Define authoritative identifiers and mapping rules.
- Separate creation rights from consumption rights.
- Document which attributes are mastered where.
- Maintain reconciliation workflows for conflicts and duplicates.
- Treat data stewardship as an operating responsibility, not an afterthought.

MDM is often perceived as bureaucratic. In reality, it is a scaling mechanism for preventing repeated operational confusion.

## 7.7 Order-to-cash in integrated commerce

Order-to-cash is one of the most important enterprise commerce value streams because it crosses customer experience, logistics, and finance.

### 7.7.1 Typical flow

A simplified integrated order-to-cash process may include:

1. customer places an order in the storefront,
2. payment is authorized or payment terms are validated,
3. order is persisted in commerce services,
4. order is exported or synchronized to ERP,
5. inventory is reserved or confirmed,
6. picking, packing, and shipment are triggered,
7. invoice is issued,
8. payment is captured or reconciled,
9. order and fulfillment statuses propagate back to customer-facing systems.

### 7.7.2 Architectural decisions inside the flow

Key design choices include:

- whether the storefront can confirm an order before ERP acceptance,
- whether stock is reserved in real time or reconciled later,
- whether ERP or commerce owns status transitions,
- how payment confirmation maps to invoicing and revenue recognition,
- which events trigger downstream systems such as warehouse or support tools.

### 7.7.3 Failure modes

Common failure modes include:

- the order is captured online but fails to post in ERP,
- payment succeeds but the order export fails,
- the storefront shows available stock that has already been consumed elsewhere,
- fulfillment completes but shipment status never propagates back to customer channels,
- invoice creation and refund processing become inconsistent during cancellations.

These are not edge cases. They are the expected complexity of distributed commercial workflows.

## 7.8 Inventory, fulfillment, and availability

Inventory integration is especially difficult because availability is both operationally critical and highly time-sensitive.

### 7.8.1 Inventory concepts that should not be conflated

- **On-hand inventory:** physical stock currently available in storage.
- **Reserved inventory:** stock allocated to pending commitments.
- **Available-to-promise:** stock that can still be sold considering current commitments and business rules.
- **Future availability:** stock expected from replenishment or transfer.

If the storefront treats all four as the same number, overselling and customer disappointment become likely.

### 7.8.2 Synchronization strategies

Common strategies include:

- near-real-time events for stock adjustments,
- short-lived availability caches,
- reservation services for high-demand items,
- periodic reconciliation against warehouse truth.

The right design depends on volume, latency needs, and oversell tolerance. Enterprise architecture should make these tradeoffs explicit rather than hiding them behind “inventory sync” language.

## 7.9 Returns and refunds

Post-purchase flows are often under-modeled in early architecture work, yet they are essential to both customer trust and accounting correctness.

### 7.9.1 Typical return flow

A return process may involve:

1. customer initiates a return,
2. policy eligibility is checked,
3. return merchandise authorization is created,
4. item is received or inspected,
5. stock is adjusted,
6. refund or credit note is issued,
7. customer and finance systems are updated.

### 7.9.2 Architectural risks

Returns create risks when:

- the storefront approves a return that finance policy rejects,
- inventory is restocked before inspection rules are applied,
- refunds are triggered twice through retries or manual intervention,

- credit notes in ERP do not align with customer-facing refund status.

These flows require the same rigor as order creation, including idempotency, state control, and clear ownership.

## 7.10 B2B pricing and account structures

B2B commerce introduces additional integration complexity because customer relationships are not purely transactional. They often involve negotiated pricing, approval chains, sales-assist workflows, credit limits, and account hierarchies.

### 7.10.1 Price lists and negotiated terms

B2B price lists may depend on:

- customer account,
- contract terms,
- region,
- product family,
- volume thresholds,
- approval status.

These rules are frequently managed in ERP or enterprise sales processes. The storefront may present the result, but it should not become a shadow source of truth for contractual pricing.

### 7.10.2 Account hierarchies

One B2B customer may represent:

- a legal entity,
- several branches,
- multiple buyers,
- different shipping destinations,
- separate billing relationships.

Architecture must therefore support organization-level identity and policy, not only individual user accounts.



## 7.11 Using Odoo as a reference ERP

Odoo is useful as a reference platform because it illustrates how enterprise workflows span modules rather than living in a single monolithic screen.

### 7.11.1 Illustrative Odoo-aligned mapping

- **Sales:** quotations, orders, commercial terms, B2B account workflows.
- **Inventory:** stock movements, reservations, warehouse updates.
- **Accounting:** invoices, payments, credit notes, reconciliation.
- **CRM:** opportunity context, customer engagement history, account management.

The point of using Odoo here is not vendor lock-in. It is to give students a concrete enterprise reference for how commerce-facing systems interact with structured back-office processes.

### 7.11.2 What should remain vendor-neutral

Even when Odoo is the teaching reference, the following concepts remain general:

- system-of-record analysis,
- event and API contracts,
- master-data governance,
- process ownership,
- enterprise reconciliation and auditability.

## 7.12 Integration risks and mitigation strategies

Representative risks in ERP/CRM/SCM integration include:

- **Duplicate sources of truth:** mitigated by explicit ownership and reconciliation.
- **Latency mismatch:** mitigated by choosing appropriate sync versus async boundaries.
- **Identifier mismatch:** mitigated by governed master identifiers and mapping tables.
- **Status divergence:** mitigated by explicit lifecycle models and state-mapping rules.
- **Operational blind spots:** mitigated by observability, error queues, and business-level monitoring.

Students should notice that these risks are architectural and organizational, not only technical.

### 7.13 Artifacts for this chapter

The expected architecture artifacts for ERP/CRM/SCM integration include:

- an integration context diagram showing commerce, ERP, CRM, warehouse, payments, and analytics boundaries,
- a system-of-record matrix for shared entities,
- a small set of data contracts for critical events or APIs,
- and a short narrative for order-to-cash and returns ownership.

### 7.14 Managerial implications

Enterprise integration decisions affect more than system design. They shape service reliability, customer trust, financial control, and the ability to scale channels. If commerce and ERP are tightly coupled in every user interaction, customer experience becomes hostage to back-office latency. If they are too loosely coupled without governance, the company loses commercial coherence.

The correct architectural aim is not maximal coupling or maximal decoupling. It is controlled coordination with explicit ownership and recoverable failure handling.

### 7.15 Multiple-choice questions (MCQs)

1. In an integrated commerce + ERP architecture, the *system of record* is best defined as:
  - (a) The system with the nicest UI
  - (b) The system that is most convenient for analytics
  - (c) The authoritative source of truth for a given entity (e.g., invoice) and its lifecycle
  - (d) The system that runs in the cloud
2. Which flow is most strongly associated with *order-to-cash*?
  - (a) Creating a product image embedding
  - (b) Quote/order → delivery → invoice → payment reconciliation
  - (c) Building a recommendation model
  - (d) Defining a Kafka topic
3. Master data management (MDM) is primarily concerned with:
  - (a) Training deep learning models for search
  - (b) Defining and governing shared entities (customer, product, supplier) and preventing duplicates/inconsistencies
  - (c) Choosing the best database index

- (d) Encrypting API requests
- 4. A common integration risk when connecting commerce with ERP is:
  - (a) Too many unit tests
  - (b) Duplicate sources of truth for prices, customers, or inventory
  - (c) Having an event schema
  - (d) Using idempotency keys
- 5. In practice, why do integrations often require both APIs and events?
  - (a) Because events are always cheaper than APIs
  - (b) Because some steps require immediate confirmation while others benefit from decoupling and retries
  - (c) Because ERP systems cannot expose APIs
  - (d) Because events guarantee zero duplicates without design effort

**Answer key**

- 1. (c)
- 2. (b)
- 3. (b)
- 4. (b)
- 5. (b)

**7.16 Exercises (short)**

- 1. For each entity, choose a system of record (commerce vs Odoo) and justify: customer, product, price list, inventory on-hand, order, invoice, payment, return.
- 2. Describe two failure modes of inventory sync and propose mitigations (e.g., eventual consistency, reservations, reconciliation jobs).
- 3. Draft a “data contract” for an `OrderCreated` event: required fields, optional fields, and versioning rules.

**7.17 Mini-case (Odoo-linked, vendor-neutral)**

**Scenario:** A hybrid company runs a headless storefront for B2C and a B2B portal for wholesale customers. Odoo is used for Sales, Inventory, Accounting, and CRM.

**Task:** Design the integration for three end-to-end processes:

- **Order-to-cash:** from checkout to invoice posting and payment reconciliation.
- **Returns/refunds:** from return request to stock adjustment and credit note.
- **B2B pricing:** negotiated price lists and approvals without duplicating sources of truth.

For each process, state: key events, required synchronous calls, idempotency strategy, and the minimum monitoring you would implement.

## Chapter 8

# Recommenders I: Classical Methods + Learning-to-Rank (Week 7)

### 8.1 Learning objectives

- Explain the business role of recommendation and ranking systems in e-commerce discovery.
- Distinguish memory-based collaborative filtering, matrix factorization, and learning-to-rank approaches.
- Understand the difference between explicit and implicit feedback, and why e-commerce is usually dominated by implicit signals.
- Evaluate recommendation systems using offline ranking metrics while recognizing the limits of offline evaluation.

### 8.2 Why recommenders matter in e-commerce

Recommendation systems are central to modern e-commerce because large catalogs create a matching problem. Customers rarely inspect all available items. The platform therefore needs mechanisms for deciding which products to surface, in what order, and in which context.

Recommendation is not only a technical personalization problem. It is also a business control layer for:

- product discovery,
- cross-sell and upsell,
- session continuation,
- re-engagement,
- merchandising efficiency,

- and marketplace supply exposure.

This chapter focuses on classical recommenders and ranking methods. These methods remain important because they are interpretable, computationally efficient, and often strong baselines even when deep-learning systems are eventually adopted.

## 8.3 Recommendation surfaces and tasks

Not all recommender problems are identical. A commerce platform usually contains several recommendation surfaces with different goals.

### 8.3.1 Common recommendation tasks

- **Homepage recommendations:** highlight likely-interest products for returning or anonymous users.
- **Product detail recommendations:** suggest similar, complementary, or substitute items.
- **Cart recommendations:** support bundle expansion and accessory attachment.
- **Email or push recommendations:** re-engage users outside the session.
- **Search ranking and re-ranking:** order retrieved items for a query.

### 8.3.2 Retrieval versus ranking

It is useful to distinguish two subproblems:

- **Candidate generation:** find a manageable subset of items from a very large catalog.
- **Ranking or re-ranking:** order the candidates for the current context.

This chapter emphasizes ranking-oriented thinking and classical candidate selection methods. Chapter 8 will extend this to deep retrieval and embedding-heavy systems.

## 8.4 Feedback in commerce recommendation

Recommendation systems learn from user behavior. The nature of that behavior strongly affects modeling choices.

### 8.4.1 Explicit feedback

Explicit feedback is directly provided by users, such as:

- star ratings,
- likes or dislikes,
- structured review scores,
- preference selections.

Explicit feedback is valuable because it is interpretable, but in many commerce settings it is sparse. Most customers do not rate enough items to support the full recommendation problem.

### 8.4.2 Implicit feedback

Implicit feedback is inferred from behavior:

- clicks,
- product views,
- add-to-cart events,
- purchases,
- dwell time,
- repeat visits,
- returns or cancellations as negative indicators.

E-commerce is largely an implicit-feedback domain. This makes modeling harder because a missing interaction does not necessarily mean dislike. It may simply mean the user never saw the item.

### 8.4.3 Signal strength

Not all implicit signals are equally informative. For example:

- a purchase usually signals stronger intent than a click,
- an add-to-cart event may indicate interest but not final preference,
- a return may weaken confidence in the original purchase signal,
- repeated category browsing may indicate exploration rather than commitment.

Feature and label design should therefore reflect business meaning, not just event counts.

## 8.5 Memory-based collaborative filtering

Classical collaborative filtering starts from the idea that similar users may like similar items, and similar items may be preferred by similar users.

### 8.5.1 User-based collaborative filtering

In user-based collaborative filtering, the system identifies users whose historical behavior resembles the target user and recommends items those similar users engaged with. Conceptually:

- define a user-item interaction matrix  $R = [r_{ui}]$ ,
- compute user-user similarity,
- aggregate neighbor preferences for unseen items.

This approach is intuitive but often scales poorly for large consumer systems with many users and rapidly changing behavior.

### 8.5.2 Item-based collaborative filtering

Item-based collaborative filtering compares items rather than users. If two items are frequently co-viewed, co-purchased, or similarly interacted with, one can recommend one when the other is observed.

Item-based methods are often more stable in commerce because items change more slowly than user sessions. They work well for:

- “similar items”,
- “customers also bought”,
- and complementary recommendation widgets.

### 8.5.3 Similarity measures

Common similarity measures include cosine similarity, Pearson correlation, and co-occurrence-based heuristics. The choice depends on data type and normalization assumptions. Architecturally, what matters is that similarity calculations must be based on well-defined interaction windows and item identities.

## 8.6 Matrix factorization

Memory-based methods are useful, but they struggle with sparsity and scale. Matrix factorization addresses this by learning low-dimensional latent representations for users and items.

### 8.6.1 Basic idea

Given an interaction matrix  $R$ , matrix factorization approximates it as:

$$\hat{r}_{ui} = \mathbf{p}_u^\top \mathbf{q}_i$$



where  $\mathbf{p}_u$  is a latent vector for user  $u$  and  $\mathbf{q}_i$  is a latent vector for item  $i$ .

These latent factors capture hidden structure such as taste patterns, style clusters, or price sensitivity without requiring manually defined categories.

### 8.6.2 Why factorization works well

Matrix factorization often performs strongly because:

- it compresses sparse interactions into learnable structure,
- it generalizes beyond direct co-occurrence,
- it supports efficient ranking once vectors are learned,
- it forms a bridge toward later embedding-based methods.

### 8.6.3 Limitations

Despite its strengths, factorization has limits:

- it handles cold-start users and items poorly without additional features,
- it may ignore rich session context,
- it can struggle when preferences shift quickly,
- it usually needs careful treatment of implicit feedback and negative sampling.

## 8.7 Implicit-feedback recommendation

Most commerce recommenders do not predict star ratings. Instead, they infer preference from implicit interactions. This changes both objective design and evaluation.

### 8.7.1 Positive-only observations

In implicit-feedback settings, observed interactions are often treated as positive evidence, while unobserved interactions are ambiguous. Training usually requires assumptions such as:

- sampling unobserved items as negatives,
- weighting confidence by interaction type,
- or optimizing pairwise ranking losses instead of rating prediction.

### 8.7.2 Confidence weighting

Different interactions can be assigned different strengths. For example:

- purchase > add-to-cart > click > impression,
- repeat purchase may strengthen confidence,
- return or refund events may weaken confidence.

This is not just a modeling choice. It is a business-rule decision and should align with how the organization interprets engagement quality.

## 8.8 Cold start and sparsity

Cold start is one of the defining challenges of recommendation systems. It occurs when the system has insufficient data about a new user, a new item, or a new marketplace participant.

### 8.8.1 User cold start

New users may have little or no history. Common mitigation strategies include:

- popularity-based defaults,
- category-entry onboarding questions,
- contextual recommendations from landing page or query behavior,
- session-based short-term signals.

### 8.8.2 Item cold start

New items may not yet have interaction data. Mitigation strategies include:

- using metadata such as category, brand, and price range,
- manual merchandising boosts,
- seller and catalog-quality priors,
- controlled exploration to gather initial interactions.

### 8.8.3 Marketplace and fairness considerations

In marketplaces, recommendation systems also influence seller exposure. If the system only amplifies already popular items, new or smaller sellers may never gather enough interaction data to compete. This creates a business and governance concern, not just a model-performance issue.

## 8.9 Learning-to-rank

Recommendation is often framed as a ranking problem rather than a rating-prediction problem. Learning-to-rank (LTR) methods optimize the order of items shown to users for a given context.

### 8.9.1 Pointwise, pairwise, and listwise views

LTR methods are often grouped into three families:

- **Pointwise:** predict a score or label for each item independently.
- **Pairwise:** learn which of two items should be ranked higher.
- **Listwise:** optimize over the ordering of whole ranked lists.

### 8.9.2 Why LTR is important in commerce

E-commerce surfaces almost always have limited slots. The platform is not merely predicting whether an item is relevant; it is deciding which few items to prioritize. LTR is therefore especially appropriate for:

- search result ordering,
- recommendation carousel ranking,
- category page re-ranking,
- email content selection.

### 8.9.3 Feature design for LTR

Classical ranking models often use hand-engineered or structured features such as:

- user affinity to category or brand,
- historical click-through rate,
- purchase propensity,
- price distance from user norms,
- popularity and freshness,
- inventory availability,
- promotion eligibility,
- query-item lexical match for search contexts.

This makes LTR highly practical because structured business features are often already available in commerce systems.

## 8.10 Offline evaluation

Recommendation systems need offline evaluation for rapid iteration, model comparison, and failure analysis before live deployment.

### 8.10.1 Common ranking metrics

Frequently used metrics include:

- $\text{precision@}k$ ,
- $\text{recall@}k$ ,
- mean average precision ( $\text{MAP@}k$ ),
- normalized discounted cumulative gain ( $\text{NDCG@}k$ ),
- mean reciprocal rank (MRR) in some search-style tasks.

### 8.10.2 What these metrics measure

- **Precision@ $k$ :** how many of the top- $k$  items are relevant.
- **Recall@ $k$ :** how much of the relevant set is recovered in the top- $k$ .
- **NDCG@ $k$ :** rewards placing more relevant items higher in the ranked list.

These metrics are useful because they reflect ranking quality rather than raw score accuracy.

### 8.10.3 Limits of offline evaluation

Offline metrics are not enough on their own. They cannot fully capture:

- causal business impact,
- position bias and feedback loops,
- changing user behavior under new experiences,
- interference with other surfaces such as promotions or search.

This is why offline evaluation should be seen as a filtering step, not the final decision criterion.

## 8.11 Online evaluation and business impact

Ultimately, recommender systems exist to improve user and business outcomes in production. Typical online measures include:

- click-through rate,
- add-to-cart rate,
- conversion rate,
- average order value,
- revenue per session,
- repeat engagement,
- return rate or dissatisfaction signals.

However, online KPI interpretation must be careful. A recommender can increase clicks while harming profitability or reducing diversity. Guardrail metrics are therefore important.

## 8.12 Biases and practical risks

Recommendation systems are affected by several structural biases.

### 8.12.1 Popularity bias

Popular items accumulate more exposure and therefore more interactions, reinforcing their future dominance. This can suppress niche discovery and reduce long-tail coverage.

### 8.12.2 Selection bias

Models only observe interactions for items that were shown. Unseen items do not generate evidence, so naive training pipelines may overestimate the quality of historically exposed items.

### 8.12.3 Position bias

Higher-ranked items receive more clicks partly because they are higher ranked, not only because they are more relevant. This complicates interpretation of click data.

### 8.12.4 Business-constraint tension

The “best” recommendation by predicted engagement may conflict with:

- margin objectives,
- inventory constraints,
- seller-fairness rules,
- compliance or policy requirements,
- user-experience diversity needs.

Recommendation is therefore not a pure prediction task. It is a constrained decision problem embedded in business policy.

## 8.13 A practical recommender architecture

Even for classical methods, recommender systems are not just models. They require an end-to-end architecture that includes:

- event collection and identity linking,
- feature generation,
- offline training,
- candidate generation,
- ranking logic,
- serving interfaces,
- monitoring and experimentation.

This chapter remains focused on the modeling layer, but students should keep the surrounding system in mind. The model only creates value when the architecture can serve it reliably and measure its effect.

## 8.14 Hands-on lab: a baseline recommender

The hands-on goal for this chapter is to build and evaluate a classical recommendation baseline.

### 8.14.1 Lab scenario

Students are given a synthetic or public interaction dataset containing user IDs, item IDs, timestamps, and event types such as view, add-to-cart, and purchase.

### 8.14.2 Lab tasks

1. Construct a user–item interaction matrix using implicit feedback.
2. Implement a simple popularity baseline and one collaborative filtering or matrix-factorization baseline.
3. Evaluate the models using  $\text{MAP}@K$  and  $\text{NDCG}@K$ .
4. Analyze cold-start and sparsity cases explicitly.
5. Propose one feature-based re-ranking rule that incorporates a business constraint such as inventory or diversity.

### 8.14.3 Expected learning outcome

By the end of the lab, students should be able to explain not only which model performed better offline, but why the evaluation setup and business constraints matter.

## 8.15 Managerial implications

Classical recommendation methods remain valuable in practice because they provide:

- strong baselines,
- fast iteration cycles,
- reasonable interpretability,
- lower infrastructure cost than many deep systems.

Organizations often make a mistake by jumping directly to deep-learning architectures without first establishing strong baselines, high-quality data, and trustworthy evaluation pipelines. For many commerce surfaces, the real bottleneck is not model sophistication but poor signal design or unclear business objectives.

## 8.16 Multiple-choice questions (MCQs)

1. In e-commerce recommendation, which statement best distinguishes *implicit* from *explicit* feedback?
  - (a) Implicit feedback is always more accurate than explicit feedback.
  - (b) Explicit feedback comes directly from user judgments such as ratings, whereas implicit feedback is inferred from behaviors such as clicks or purchases.
  - (c) Explicit feedback can only be used in search ranking.
  - (d) Implicit feedback means the platform has no user identifiers.

2. Item-based collaborative filtering is often attractive in commerce because:
  - (a) items are usually more stable than users and item-item similarities can support scalable recommendation widgets.
  - (b) it eliminates cold start entirely.
  - (c) it requires no interaction data.
  - (d) it guarantees causal lift.
3. In matrix factorization, the expression  $\hat{r}_{ui} = \mathbf{p}_u^\top \mathbf{q}_i$  is intended to represent:
  - (a) a network timeout policy.
  - (b) the predicted affinity between user  $u$  and item  $i$  using latent factors.
  - (c) the exact accounting value of an order.
  - (d) a lossless compression ratio for logs.
4. Which metric is most appropriate for evaluating the quality of a top- $k$  ranked recommendation list offline?
  - (a) NDCG@ $k$
  - (b) Mean absolute error on shipping cost
  - (c) CPU utilization
  - (d) Database replication lag
5. Why is learning-to-rank especially relevant in e-commerce?
  - (a) Because commerce systems never need candidate generation.
  - (b) Because the platform must decide which limited set of items to place at the top of a ranked surface.
  - (c) Because it replaces online experimentation.
  - (d) Because it removes the need for features.

### Answer key

1. (b)
2. (a)
3. (b)
4. (a)
5. (b)



## 8.17 Exercises (short)

1. Compare user-based, item-based, and matrix-factorization approaches for a medium-sized commerce catalog. State one strength and one weakness of each.
2. Define an implicit-feedback weighting scheme for view, add-to-cart, purchase, and return events. Explain the business reasoning behind your weights.
3. Design a simple learning-to-rank feature set for a “recommended for you” widget. Include at least one user feature, one item feature, one context feature, and one business-rule feature.

## 8.18 Mini-case (Odoo-linked, vendor-neutral)

**Scenario:** A company uses a headless storefront, Odoo for ERP, and a central data platform. The business wants recommendation widgets on the homepage, product pages, and cart page.

**Task:** Propose a classical recommender design:

- Define which interaction signals from storefront and Odoo-derived outcomes should be used for training.
- Choose a baseline method (item-based CF, matrix factorization, or feature-based ranking) for each surface and justify the choice.
- Identify two business constraints the ranking layer should respect, such as inventory, margin, seller exposure, or return risk.



## Chapter 9

# Recommenders II: Deep Learning, Embeddings, and Retrieval (Week 8)

### 9.1 Learning objectives

- Explain why deep learning methods are useful for large-scale recommendation and retrieval problems in e-commerce.
- Distinguish candidate generation from re-ranking and understand why modern recommenders are often multi-stage systems.
- Describe embeddings, two-tower retrieval, sequence models, and session-aware recommendation at a system level.
- Analyze production concerns such as cold start, diversity, latency, and bias in deep recommender architectures.

### 9.2 From classical recommenders to deep retrieval

Chapter 7 focused on classical collaborative filtering, matrix factorization, and learning-to-rank. Those methods remain important, but large-scale e-commerce platforms often face additional pressures:

- extremely large catalogs,
- rapidly changing inventories,
- sparse user histories,
- multi-modal item information such as text and images,
- session behavior that changes within minutes,
- and tight latency requirements at serving time.

Deep learning is useful in this setting because it can learn richer representations from heterogeneous data. Rather than relying only on manually designed similarity signals, deep systems can learn embeddings that organize users, items, queries, sessions, and contexts into a common representational space.

## 9.3 Modern recommender architecture: retrieval and ranking

Large recommender systems are rarely a single model. They are usually multi-stage decision systems.

### 9.3.1 Candidate generation

Candidate generation or retrieval narrows a huge catalog to a much smaller set of promising items. For a catalog containing millions of products, the system cannot fully rank every item in real time for every request. Retrieval therefore acts as a coarse but scalable selection layer.

### 9.3.2 Re-ranking

After candidate generation, a second-stage model or rule system orders the retrieved candidates more precisely. This stage can use richer features, business constraints, and context-aware logic because the item set is now small enough to handle with more expensive computation.

### 9.3.3 Why multi-stage design matters

The multi-stage architecture exists because of a systems tradeoff:

- retrieval prioritizes speed and coverage,
- re-ranking prioritizes precision and business control.

Students should understand that recommendation quality depends not only on model accuracy, but also on how these stages interact. A weak retrieval stage can permanently exclude relevant items before ranking ever sees them.

## 9.4 Embeddings as learned representations

Embeddings are dense vectors that represent entities such as users, items, queries, categories, or sessions. They generalize the latent-factor intuition of matrix factorization, but in a more flexible framework.

### 9.4.1 Why embeddings are powerful

Embeddings are useful because they:

- capture similarity in a continuous space,
- support nearest-neighbor retrieval,
- can combine multiple signals and modalities,
- transfer across tasks such as recommendation, search, and personalization.

### 9.4.2 What can be embedded

Common embedded entities in commerce include:

- users or customer accounts,
- products,
- search queries,
- brands or categories,
- sellers,
- sessions,
- product text and images.

### 9.4.3 Interpretation

An embedding dimension does not usually have a simple human-readable meaning. The value of the representation comes from geometry: items or users that are close in vector space are more likely to be related in terms of behavior or semantics.

## 9.5 Two-tower retrieval models

One of the most influential deep retrieval architectures is the two-tower model. It is widely used because it balances expressiveness with serving practicality.

### 9.5.1 Basic structure

A two-tower model learns:

- one tower that maps user or context features to an embedding,
- one tower that maps item features to an embedding.

The compatibility between a user/context and an item is then computed using a similarity function such as dot product or cosine similarity.

### 9.5.2 Why two towers are useful

Two-tower models are attractive because item embeddings can often be precomputed offline and indexed for fast approximate nearest-neighbor retrieval. At serving time, the system computes the user or session embedding and retrieves nearby items efficiently.

### 9.5.3 Typical inputs

The user/context tower may consume:

- recent interaction history,
- device and channel context,
- customer segment,
- geographic or temporal features,
- active search or page context.

The item tower may consume:

- item ID,
- category,
- brand,
- price bucket,
- title and description text,
- image-derived features,
- seller or inventory context.

## 9.6 Approximate nearest-neighbor retrieval

Deep retrieval only becomes operationally useful when the serving system can search the embedding space quickly. This is where approximate nearest-neighbor (ANN) indexing enters the architecture.

### 9.6.1 Why exact search is often too expensive

If the platform must compare each request embedding against millions of item embeddings, exact retrieval may be too slow for production latency budgets. ANN methods trade a small amount of exactness for major speed improvements.

### 9.6.2 Architectural implication

This means recommender design now depends on more than model training. It also depends on:

- embedding refresh cadence,
- index rebuild strategy,
- item insertion and deletion handling,
- consistency between item catalog state and retrieval index.

An item that is unavailable or unpublished should not remain in the active candidate index without explicit policy.

## 9.7 Sequence models and session-aware recommendation

Many commerce interactions are highly sequential. A user may move from broad exploration to narrow intent within a single session. Classical long-term preference models may miss this short-term drift.

### 9.7.1 Session behavior

Session-aware recommendation tries to capture recency and order, not only aggregate history. Examples of useful sequential signals include:

- last viewed products,
- order of category transitions,
- repeated brand examination,
- add-to-cart followed by substitution browsing,
- search reformulation patterns.

### 9.7.2 Sequence model families

Common sequence-oriented approaches include:

- recurrent models,
- convolutional sequence models,
- Transformer-style attention models,
- next-item prediction objectives.

### 9.7.3 Why sequence models help

These models help because they can distinguish:

- stable preference from temporary intent,
- exploration from purchase preparation,
- repeated exposure from genuine narrowing of interest.

This is especially useful for anonymous users, where long-term account history may be unavailable.

## 9.8 Multi-modal recommendation

Deep systems are particularly useful when products have rich unstructured content. E-commerce catalogs often contain:

- titles and descriptions,
- structured attributes,
- images,
- reviews,
- seller-generated content,
- taxonomy information.

### 9.8.1 Text and semantic understanding

Text encoders can help represent product meaning beyond exact attribute matching. For example, two products may be semantically similar even when they do not share the same surface keywords.

### 9.8.2 Image understanding

Visual features matter strongly in apparel, furniture, decor, and other style-driven categories. Image embeddings can support:

- visually similar item recommendation,
- cold-start support for new products,
- quality or anomaly checks for catalog content.



### 9.8.3 Fusion

Modern item representations often combine ID-based, attribute-based, text-based, and image-based features. This reduces dependence on raw interaction history alone and helps address sparse data problems.

## 9.9 Deep re-ranking

Retrieval is only the first stage. Deep models can also be applied in re-ranking after candidates are retrieved.

### 9.9.1 Why re-ranking differs from retrieval

The re-ranker can use richer cross-features because it evaluates a much smaller candidate set. For example, it may incorporate:

- user-item interaction history,
- session context,
- merchandising rules,
- margin, inventory, or delivery constraints,
- diversity adjustments,
- fraud or return-risk features.

### 9.9.2 Business role of re-ranking

The re-ranker is often where business policy becomes explicit. The retrieval stage finds promising candidates. The re-ranker decides which candidates best satisfy a mix of relevance and operational constraints.

## 9.10 Training data and objectives

Deep recommender quality depends heavily on how the training problem is defined.

### 9.10.1 Positive and negative examples

Training often requires:

- positive examples from clicks, carts, or purchases,

- sampled negatives from unseen or skipped items,
- sometimes hard negatives that were exposed but not chosen.

### 9.10.2 Objective choices

Common objectives include:

- next-item prediction,
- pairwise ranking losses,
- contrastive retrieval objectives,
- multi-task objectives combining click, cart, and purchase behavior.

### 9.10.3 Label quality

As in earlier chapters, the label design is a business interpretation problem. If a purchase is later returned, or if a click was caused mainly by position bias, the training data need careful handling. Deep models can amplify weak signals just as efficiently as strong ones.

## 9.11 Cold start in deep recommendation

Deep systems do not eliminate cold start. They only change the available tools for addressing it.

### 9.11.1 User cold start

For new users, deep models may rely on:

- session behavior,
- page context,
- acquisition source,
- demographic or geo context where appropriate,
- popularity and trend priors.

### 9.11.2 Item cold start

For new items, deep models can use metadata and content features more effectively than classical ID-only approaches. Text and image embeddings are especially helpful because they allow the item to be placed in a meaningful neighborhood before interaction data accumulate.

### 9.11.3 Catalog churn

Commerce catalogs change frequently. A useful architecture must define how new item embeddings are generated, validated, indexed, and exposed without waiting for a full model retrain in every case.

## 9.12 Diversity, exploration, and bias

A recommender that always shows the most probable item is not necessarily the best commerce system. The platform may need to balance relevance with diversity and exploration.

### 9.12.1 Diversity

Diversity can improve user experience by preventing repetitive lists and broadening discovery. It may also improve business robustness by reducing over-concentration on a few items or sellers.

### 9.12.2 Exploration

Exploration is necessary to learn about new items, uncertain user interests, and shifting trends. Without exploration, the platform over-commits to historical winners and limits future learning.

### 9.12.3 Bias

Deep systems can intensify:

- popularity bias,
- seller-exposure bias,
- feedback-loop bias,
- and representation bias if product metadata are uneven or low quality.

This means governance remains essential even when model quality appears strong offline.

## 9.13 Serving and latency constraints

Deep recommenders exist under production latency budgets. This makes serving architecture as important as modeling.

### 9.13.1 Key serving questions

- Which features are available in real time?
- Which embeddings can be precomputed?
- How often are indexes refreshed?
- What is the fallback plan if the retrieval service is unavailable?
- How are out-of-stock or restricted items removed before final ranking?

### 9.13.2 Caching and freshness

Caching can reduce latency, but stale caches can recommend unavailable or irrelevant items. Architecture must define freshness tolerances by surface. For example, homepage suggestions may tolerate more staleness than cart recommendations near checkout.

### 9.13.3 Fallback behavior

Production systems usually need fallback strategies such as:

- popularity-based lists,
- category best-sellers,
- recent-view heuristics,
- rule-based complements.

Fallbacks are not a sign of failure. They are part of robust recommender system design.

## 9.14 Evaluation beyond offline metrics

Deep recommendation systems are often evaluated offline with ranking metrics, but their real value appears only in production.

### 9.14.1 Why deep models can mislead offline

Deep models may improve offline ranking scores because they capture rich structure, yet still disappoint online if:

- latency increases and user experience degrades,
- popularity bias worsens,

- retrieval excludes important fresh inventory,
- serving features differ from training features,
- business constraints are ignored in ranking.

### 9.14.2 Online considerations

Relevant online metrics include:

- click-through and add-to-cart lift,
- incremental conversion,
- revenue per session,
- item and seller coverage,
- return rate,
- latency and availability,
- fairness or exposure guardrails.

## 9.15 A practical deep recommender stack

A practical deep recommender architecture often includes:

- event collection and identity stitching,
- a feature store or equivalent feature pipelines,
- model training jobs,
- embedding generation,
- ANN index building,
- retrieval serving,
- re-ranking serving,
- experimentation and monitoring layers.

This is why deep recommendation belongs not only to machine learning, but also to systems design and MLOps.

## 9.16 Hands-on lab: retrieval plus re-ranking

The goal of the practical work is to help students think beyond one-model recommendation.

### 9.16.1 Lab scenario

Students are given a dataset with user interaction histories, product metadata, and timestamps. They are asked to design a two-stage pipeline for recommendation.

### 9.16.2 Lab tasks

1. Train or simulate product embeddings from interaction and metadata signals.
2. Build a simple candidate retrieval mechanism using embedding similarity.
3. Add a lightweight re-ranking stage using business-aware features such as inventory or price sensitivity.
4. Define a serving plan that explains which representations are precomputed and which features are computed online.
5. Propose online guardrail metrics that would be monitored during an experiment.

### 9.16.3 Expected learning outcome

By the end of the lab, students should be able to explain why a high-quality recommender is a pipeline and not just a neural model.

## 9.17 Managerial implications

Deep recommendation systems can unlock major gains in relevance and catalog understanding, but they also introduce higher system complexity. Organizations should adopt them when the business problem justifies:

- richer representation learning,
- very large-scale retrieval,
- multi-modal item understanding,
- session-aware adaptation.

They should not adopt them merely because they sound more advanced. Without strong data foundations, business constraints, and serving discipline, deep recommenders can become expensive complexity with limited net value.

## 9.18 Multiple-choice questions (MCQs)

1. In a modern recommender stack, the main purpose of *candidate generation* is to:

- (a) finalize the exact business KPI improvement.
  - (b) reduce a very large catalog to a manageable subset before detailed ranking.
  - (c) eliminate the need for embeddings.
  - (d) replace experimentation.
2. A two-tower retrieval model is most accurately described as:
- (a) a model with one tower for pricing and one tower for fraud.
  - (b) separate encoders for user/context and item representations whose similarity supports retrieval.
  - (c) a rule engine with two fallback policies.
  - (d) a batch-only ERP integration pattern.
3. Why are embeddings useful in e-commerce recommendation?
- (a) They always eliminate cold start completely.
  - (b) They provide dense learned representations that support similarity and retrieval across users, items, and contexts.
  - (c) They make item metadata unnecessary.
  - (d) They are only relevant for payment systems.
4. Sequence models are especially valuable when:
- (a) user intent changes within a session and order of interactions matters.
  - (b) the catalog has only ten products.
  - (c) recommendation is based only on ERP invoices.
  - (d) latency does not matter.
5. Which is the best example of a production concern specific to deep retrieval systems?
- (a) Whether the catalog has chapter numbers.
  - (b) How to refresh embeddings and ANN indexes while respecting latency and availability constraints.
  - (c) Whether to use bullet points in documentation.
  - (d) Whether offline  $\text{MAP}@K$  is equal to  $\text{precision}@K$ .

**Answer key**

- 1. (b)
- 2. (b)
- 3. (b)
- 4. (a)
- 5. (b)

## 9.19 Exercises (short)

1. Compare a classical matrix-factorization recommender and a two-tower retrieval system for a large commerce catalog. State two reasons why the deep system may help and two costs it introduces.
2. Propose an embedding feature set for items in a fashion store. Include structured, textual, and visual signals.
3. Design a retrieval plus re-ranking pipeline for a cart recommendation widget. Specify where you would enforce inventory and diversity constraints.

## 9.20 Mini-case (Odoo-linked, vendor-neutral)

**Scenario:** A hybrid B2C/B2B commerce company uses a headless storefront, Odoo for ERP, and a central ML platform. The product catalog is large and changes frequently.

**Task:** Design a deep recommender architecture:

- Define the inputs to the user/context tower and the item tower.
- Explain how item embeddings would be refreshed when products, prices, or availability change.
- Identify two risks in using deep retrieval for B2B customers with negotiated catalogs or restricted assortments, and propose mitigations.



## Chapter 10

# Pricing and Promotions with ML + Optimization (Week 9)

### 10.1 Learning objectives

- Explain why pricing in e-commerce is both a prediction problem and an optimization problem.
- Distinguish demand forecasting, price elasticity estimation, and promotion uplift modeling.
- Analyze how business constraints such as inventory, margins, contracts, and competition shape pricing decisions.
- Design a pricing workflow that combines forecasting, optimization, and experimentation rather than treating price as a static input.

### 10.2 Why pricing is a core AI problem in commerce

Pricing is one of the most consequential decisions in e-commerce. It affects demand, conversion, margin, inventory movement, customer perception, and competitive position. Unlike many prediction tasks, pricing also changes the system it measures. When a platform changes price, user behavior changes, competitor response may change, and future training data are altered.

This is why pricing cannot be treated as “just another regression target.” The practical pricing problem combines:

- demand estimation,
- elasticity understanding,
- business-rule enforcement,
- optimization under constraints,
- and experimentation to measure causal impact.

In enterprise settings, pricing also interacts with ERP-owned price lists, promotions engines, finance controls, and B2B contractual terms.

## 10.3 Pricing objectives

Different organizations optimize different pricing objectives. Students should recognize that “best price” depends on the business question.

### 10.3.1 Common objectives

- **Revenue maximization:** increase gross sales volume.
- **Margin or contribution optimization:** emphasize profitability rather than top-line revenue.
- **Inventory balancing:** accelerate slow-moving stock or protect scarce inventory.
- **Customer acquisition or retention:** use pricing to influence long-term account value rather than immediate transaction value.
- **Competitive positioning:** maintain price perception within a market segment.

### 10.3.2 Why objective clarity matters

The same product may warrant different prices depending on whether the business prioritizes:

- clearing inventory,
- preserving premium brand position,
- growing category share,
- protecting margin under rising costs.

A pricing model trained without explicit objective alignment often produces analytically plausible but operationally poor recommendations.

## 10.4 Price, promotion, and discount structure

In e-commerce, price is rarely a single number. The platform may have:

- list price,
- sale price,
- coupon-adjusted price,

- customer-specific or B2B negotiated price,
- bundle price,
- shipping-inclusive effective price.

### 10.4.1 Promotions as structured interventions

Promotions are not merely “lower prices.” They can take many forms:

- percentage discount,
- fixed-amount coupon,
- buy-one-get-one structures,
- free shipping thresholds,
- loyalty-only offers,
- category-wide or seller-funded campaigns.

Each structure affects user behavior differently. For example, a coupon may increase conversion among already interested users, while a shipping threshold may increase basket size rather than item-level demand.

## 10.5 Demand forecasting

Pricing decisions require an estimate of how demand changes over time and under different commercial conditions.

### 10.5.1 What is being forecast

Forecast targets may include:

- unit demand by SKU and day,
- category demand,
- promotion-period lift,
- expected cart conversion,
- replenishment-sensitive demand.

### 10.5.2 Inputs to demand models

Typical features include:

- historical sales,
- current and historical prices,
- promotions,
- seasonality and holidays,
- catalog position and visibility,
- marketing traffic,
- competitor signals where available,
- inventory availability.

### 10.5.3 Why forecasting is not enough

Forecasting predicts what may happen. Pricing still requires a decision about what the system should do. A forecast without an optimization layer does not tell the platform which price to choose among feasible options.

## 10.6 Price elasticity

Elasticity describes how demand changes in response to price changes. It is one of the most important concepts in pricing analytics.

### 10.6.1 Intuition

If a small increase in price causes a large reduction in demand, the product is relatively price-sensitive. If demand changes little, the product is relatively price-insensitive.

### 10.6.2 Why elasticity is difficult to estimate

In live commerce systems, observed prices are not randomly assigned. They may reflect prior business policy, inventory conditions, promotions, or manual intervention. This means historical price-demand relationships are often confounded.

### 10.6.3 Practical considerations

Elasticity may vary by:

- category,
- season,
- customer segment,
- brand strength,
- competitive intensity,
- stock scarcity,
- whether the user arrived organically or through a promotion.

Students should therefore avoid assuming a single global elasticity coefficient is sufficient for realistic commerce settings.

## 10.7 Promotions and uplift

Promotion modeling asks not only “what happened under promotion?” but “what additional effect did the promotion cause?”

### 10.7.1 Observed uplift versus causal uplift

If conversion rises during a campaign, that increase may reflect:

- genuine causal impact from the promotion,
- seasonal demand shifts,
- traffic-source changes,
- selection bias in who received the offer.

True uplift modeling attempts to estimate incremental effect rather than raw observed response.

### 10.7.2 Why uplift matters

Promotions can be expensive. If the platform discounts heavily for users who would have purchased anyway, the business sacrifices margin without creating incremental value. Targeted promotion strategies therefore require causal thinking, not just descriptive analytics.

## 10.8 Optimization under business constraints

After forecasting and elasticity estimation, the pricing problem becomes an optimization problem.

### 10.8.1 Typical constraints

Real pricing decisions are constrained by:

- minimum margin thresholds,
- inventory availability,
- supplier or brand pricing rules,
- MAP or contractual restrictions,
- competitive guardrails,
- fairness or anti-discrimination rules,
- B2B price-list obligations,
- operational simplicity limits on how often prices can change.

### 10.8.2 Why unconstrained optimization is unrealistic

A mathematically optimal price from a simple model may be unusable if it:

- violates margin policy,
- creates stockouts,
- damages brand trust,
- conflicts with contractual pricing,
- or overwhelms operations with constant price updates.

Optimization in enterprise commerce therefore means “best feasible decision,” not “best unconstrained number.”

## 10.9 Dynamic pricing

Dynamic pricing adjusts prices over time in response to changing conditions. It is widely discussed, but it should not be treated as a universal default.

### 10.9.1 When dynamic pricing can help

Dynamic pricing is often useful when:

- demand fluctuates quickly,
- inventory is perishable or time-sensitive,
- competitive prices move frequently,
- the catalog is large enough that manual pricing is infeasible.

### 10.9.2 When dynamic pricing can hurt

It can also create problems:

- customer trust may decline if prices feel unstable or unfair,
- operations may struggle with frequent updates,
- historical comparison becomes harder,
- legal or contractual issues may arise in some markets.

The correct architecture therefore distinguishes between categories that justify dynamic treatment and those that require more stable commercial rules.

## 10.10 Competition-aware pricing

In many categories, prices are influenced by competitor behavior. Yet competitor-aware pricing is not the same as “match the lowest price.”

### 10.10.1 Competitive signals

Useful signals may include:

- competitor list price,
- promotion intensity,
- assortment overlap,
- shipping speed comparisons,
- seller reputation and availability differences.

### 10.10.2 Strategic caution

A competitor may use a loss leader or temporary campaign that should not be copied directly. Blind matching can destroy margin without improving long-term position. The architecture should therefore support competitor monitoring while still respecting internal pricing objectives and constraints.

## 10.11 B2C versus B2B pricing

Pricing design differs sharply between B2C and B2B commerce.

### 10.11.1 B2C

B2C pricing usually emphasizes:

- high-volume SKU-level decisions,
- campaign responsiveness,
- public price perception,
- consumer conversion behavior.

### 10.11.2 B2B

B2B pricing often involves:

- negotiated price lists,
- account-specific discounts,
- credit terms,
- approval workflows,
- contract validity windows,
- product-assortment restrictions.

This means many B2B pricing decisions belong closer to ERP and enterprise sales workflows than to the storefront alone.



## 10.12 Feature design for pricing models

Pricing models often require features from multiple systems. Representative feature groups include:

- **Product features:** category, brand, cost, lifecycle stage, substitutes.
- **Demand features:** recent sales velocity, conversion rate, traffic.
- **Promotion features:** campaign type, discount depth, channel.
- **Inventory features:** on-hand stock, days of supply, replenishment lead time.
- **Customer features:** segment, loyalty status, contract class.
- **Market features:** competitor price, holiday periods, macro signals.

### 10.12.1 Feature governance

Pricing is high impact, so feature governance matters. If cost, inventory, or contract data are delayed or inconsistent, the resulting recommendations may be financially harmful. This ties pricing directly to the data-foundation and systems-of-record questions from earlier chapters.

## 10.13 Experimentation in pricing

Offline evaluation is useful, but pricing decisions have strong causal and behavioral effects. Therefore online experimentation is especially important.

### 10.13.1 Why pricing experiments are sensitive

Pricing tests differ from many UI experiments because they:

- affect revenue directly,
- may trigger competitor response,
- can influence customer trust,
- may have spillover across channels or segments,
- often require tighter legal and policy review.

### 10.13.2 What to measure

Useful experiment metrics may include:

- conversion rate,

- unit sales,
- revenue,
- contribution margin,
- basket size,
- repeat purchase,
- return rate,
- customer support complaints,
- price-perception surveys where available.

### **10.13.3 Guardrails**

Pricing experiments should include guardrails such as:

- margin floor protection,
- inventory depletion limits,
- customer fairness checks,
- B2B account or contractual exclusions.

## **10.14 A practical pricing architecture**

A practical pricing system often includes:

- demand and elasticity data pipelines,
- forecasting models,
- optimization logic,
- a pricing decision service,
- a promotions engine,
- ERP integration for governed price lists and approvals,
- experimentation and monitoring controls.

### 10.14.1 Decision flow

A simplified pricing flow may be:

1. gather product, demand, inventory, and market signals,
2. forecast demand under candidate price scenarios,
3. optimize within business constraints,
4. publish recommended price or promotional action,
5. validate through workflow or approval gates,
6. serve price to commerce channels,
7. measure downstream outcomes.

### 10.14.2 Why architecture matters

Pricing is not just a model output. It is a governed operational decision that affects customers, suppliers, finance, and brand perception. The architecture must therefore support auditability and rollback, not only prediction accuracy.

## 10.15 Common failure modes

Pricing systems often fail in predictable ways:

- forecasting ignores stockouts and mistakes absence of sales for lack of demand,
- elasticity estimates are biased by historical policy,
- optimization recommends prices that violate margin or contractual constraints,
- promotions are evaluated on revenue alone while margin collapses,
- competitor signals are stale or incomparable,
- different channels display conflicting prices due to synchronization delays.

These failures are reminders that pricing is a socio-technical control problem, not just an analytics exercise.

## 10.16 Hands-on lab: forecast plus constrained optimization

The practical goal of this chapter is to connect prediction and decision-making.

**10.16.1 Lab scenario**

Students are given historical SKU sales, prices, promotion flags, inventory levels, and seasonality indicators. They are asked to produce a pricing recommendation for a small catalog segment.

**10.16.2 Lab tasks**

1. Build a baseline demand forecast for a selected product or category.
2. Estimate how demand changes across a small set of candidate prices.
3. Define a margin floor and one inventory-related constraint.
4. Choose a recommended price or promotion and justify it.
5. Specify how you would validate the decision using an online experiment.

**10.16.3 Expected learning outcome**

By the end of the lab, students should be able to explain why pricing requires a pipeline of assumptions, constraints, and measurement rather than a single predictive model.

**10.17 Managerial implications**

Pricing systems can create large business value, but they can also create large business damage when poorly governed. Organizations should therefore invest in:

- clear pricing objectives,
- data quality around cost, stock, and catalog attributes,
- explicit separation between descriptive analytics and prescriptive decisions,
- and controlled experimentation before scaling automated pricing.

The most effective pricing systems are not those that change the most prices. They are those that align predictive intelligence with disciplined commercial control.

**10.18 Multiple-choice questions (MCQs)**

1. Why is pricing in e-commerce best viewed as both a prediction and an optimization problem?
  - (a) Because prices are always random.
  - (b) Because the platform must estimate demand response and then choose the best feasible action under business constraints.

- (c) Because forecasting and optimization are mathematically identical.
  - (d) Because optimization eliminates the need for data.
2. Which statement best describes price elasticity?
- (a) It is the speed of the API that serves prices.
  - (b) It describes how demand responds when price changes.
  - (c) It is the number of SKUs in a category.
  - (d) It guarantees causal inference from any historical dataset.
3. In promotion analytics, *uplift* is most closely concerned with:
- (a) the additional causal effect of a promotion beyond what would have happened otherwise.
  - (b) the size of the ERP database.
  - (c) the number of products in a campaign.
  - (d) the compression ratio of log files.
4. Which is the best example of a realistic pricing constraint?
- (a) The model should only use purple charts.
  - (b) The recommended price must preserve a minimum margin and respect contractual rules.
  - (c) The system should never store timestamps.
  - (d) The experiment should avoid all monitoring.
5. Why are pricing experiments especially sensitive?
- (a) Because prices have no effect on customer behavior.
  - (b) Because pricing directly affects revenue, margin, customer trust, and sometimes legal or contractual obligations.
  - (c) Because they can only be run offline.
  - (d) Because recommendations and pricing are unrelated.

**Answer key**

- 1. (b)
- 2. (b)
- 3. (a)
- 4. (b)
- 5. (b)

## 10.19 Exercises (short)

1. Compare revenue-maximizing and margin-maximizing pricing decisions for one product category. Explain when they may diverge.
2. Design a feature set for estimating promotional uplift on a category page campaign. Include traffic, inventory, price, and seasonality signals.
3. Propose a constrained pricing policy for a B2B account segment where negotiated price lists exist in ERP but promotional flexibility is allowed for selected categories.

## 10.20 Mini-case (Odoo-linked, vendor-neutral)

**Scenario:** A hybrid B2C/B2B company uses a headless storefront, Odoo for ERP and accounting, and a data platform for forecasting. The company wants better pricing decisions for seasonal inventory and account-specific promotions.

**Task:** Design a pricing workflow:

- Define which signals come from storefront behavior, which come from Odoo, and which are external.
- Explain how you would separate public promotional pricing from ERP-governed B2B price lists.
- Identify two experiment guardrails and two operational constraints that must be enforced before automated pricing recommendations are exposed.

# Chapter 11

## Fraud, Risk, and Trust & Safety (Week 10)

### 11.1 Learning objectives

- Identify major categories of fraud and abuse in e-commerce, including payment fraud, account takeover, promotion abuse, and refund abuse.
- Explain why fraud detection is a cost-sensitive, adversarial, and highly imbalanced decision problem.
- Compare supervised learning, anomaly detection, graph-based methods, and rule+ML hybrid approaches.
- Design a trust-and-safety workflow that combines scoring, thresholds, review operations, and business guardrails.

### 11.2 Why fraud and trust & safety matter

E-commerce systems are open economic systems. They make it easy for legitimate customers to browse, buy, receive promotions, request returns, and manage accounts. Unfortunately, the same accessibility also creates opportunities for malicious behavior.

Fraud and abuse affect more than chargeback loss. They influence:

- payment costs,
- customer trust,
- seller trust in marketplaces,
- operational review burden,
- promotional efficiency,

- regulatory and compliance exposure.

For this reason, fraud prevention is best understood as part of a broader *trust & safety* capability rather than a narrow payment model. The system must protect revenue while preserving a smooth user experience for legitimate customers.

## 11.3 Fraud and abuse categories

Fraud in e-commerce is heterogeneous. Different abuse patterns require different signals and response strategies.

### 11.3.1 Payment fraud

Payment fraud often involves stolen cards, synthetic identities, mule accounts, or coordinated abuse of payment flows. The platform must distinguish legitimate but unusual behavior from malicious transactions.

### 11.3.2 Account takeover

Account takeover occurs when an attacker gains control of an existing customer account. This can lead to unauthorized purchases, loyalty theft, stored-value abuse, and privacy harm.

### 11.3.3 Promotion and coupon abuse

Promotions are common attack surfaces because they create direct economic value. Abuse patterns include:

- creating multiple accounts to redeem new-user offers,
- coordinated coupon sharing,
- automated checkout attempts,
- exploiting referral programs or cashback loops.

### 11.3.4 Refund and return abuse

Refund abuse includes false non-delivery claims, empty-box returns, repeated wardrobing behavior, and manipulation of return policies. This category often requires joining transaction, logistics, support, and account-history signals.



### 11.3.5 Marketplace and seller abuse

In multi-seller environments, trust & safety also includes:

- counterfeit listings,
- fake reviews,
- seller collusion,
- policy evasion,
- manipulated performance metrics.

## 11.4 Fraud as a machine learning problem

Fraud detection is not a standard balanced classification task. It has several defining properties.

### 11.4.1 Class imbalance

Fraudulent events are typically a small fraction of all transactions. This means a naive high-accuracy classifier may still be useless if it mostly predicts “not fraud.”

### 11.4.2 Adversarial behavior

Fraudsters adapt. Once a platform changes a rule or scoring model, attackers may alter devices, payment instruments, shipping behavior, or timing patterns. This makes fraud a moving-target problem.

### 11.4.3 Label delay and label noise

Ground truth often arrives late. Chargebacks, disputes, and investigations may occur days or weeks after the original transaction. Even then, labels can be incomplete or noisy.

### 11.4.4 Asymmetric costs

False negatives lose money and trust. False positives block legitimate customers and harm conversion. The decision threshold must therefore reflect business costs rather than abstract accuracy alone.

## 11.5 Signals and features for fraud detection

Fraud models typically depend on a broad feature space spanning account, transaction, device, and network behavior.

### 11.5.1 Transaction signals

- order value,
- payment method,
- basket composition,
- discount depth,
- shipping speed,
- billing-shipping mismatch,
- repeated authorization attempts.

### 11.5.2 Account signals

- account age,
- prior successful orders,
- return history,
- password reset activity,
- loyalty balance behavior,
- contact-information change frequency.

### 11.5.3 Device and network signals

- device fingerprint,
- IP reputation,
- geolocation inconsistency,
- emulator or automation indicators,
- velocity across devices or sessions.

### 11.5.4 Behavioral and sequence signals

- login and checkout timing,
- session depth,
- rapid account creation followed by high-value purchase,
- repeated failed payments,
- navigation patterns inconsistent with normal shopping behavior.

These signals are only useful when identity linkage and event quality are sound. Fraud modeling therefore depends directly on the data-foundation chapter and on integration with operational systems.

## 11.6 Rule-based systems

Fraud prevention historically relied heavily on rules. Rules remain important because they are fast, transparent, and easy to deploy for known abuse patterns.

### 11.6.1 Examples of useful rules

- block transactions from known bad cards or IPs,
- trigger review if order value exceeds a threshold and account age is very low,
- limit coupon redemption frequency by account, device, and payment instrument,
- hold high-risk shipping destinations when signals conflict.

### 11.6.2 Strengths and weaknesses

Rules are valuable for:

- deterministic policy enforcement,
- emergency response,
- explainability in operations.

They are weak when:

- abuse patterns evolve quickly,
- interactions among many weak signals matter,
- manual rule maintenance becomes operationally expensive.

## 11.7 Supervised learning

Supervised models learn from labeled examples of fraud and legitimate behavior. They are widely used because they can combine many signals more flexibly than hand-written rules.

### 11.7.1 Typical model families

Common model families include:

- logistic regression,
- tree-based ensembles,
- gradient boosting,
- neural models where scale and feature complexity justify them.

### 11.7.2 Why supervised learning helps

Supervised learning helps by:

- aggregating many weak predictors,
- ranking risk more smoothly than rigid rules,
- enabling threshold tuning by cost and segment,
- adapting to new feature sets over time.

### 11.7.3 Operational caution

Supervised models are only as good as:

- the quality of labels,
- the recency of training data,
- the alignment between training and serving features,
- the monitoring for drift and policy changes.

## 11.8 Anomaly detection

Some abuse patterns are rare or novel enough that labeled examples are limited. In these settings, anomaly detection can be useful.

### 11.8.1 What anomaly methods look for

Anomaly-oriented methods try to identify behavior that differs sharply from established norms, such as:

- unusual transaction timing,
- atypical device usage,
- unexpected country or route changes,
- extreme coupon redemption behavior,
- sudden seller or account behavior shifts.

### 11.8.2 Strengths and weaknesses

Anomaly methods are useful for surfacing suspicious patterns and complementing supervised systems. However, anomaly does not automatically mean fraud. Legitimate edge cases can also appear unusual. This means anomaly scores usually work best as part of a broader review or hybrid decision pipeline.

## 11.9 Graph-based fraud analysis

Fraud is often relational. Accounts, devices, cards, addresses, merchants, and products form networks rather than isolated rows of data.

### 11.9.1 Why graph structure matters

Graph methods help detect:

- many accounts linked to one device,
- many cards shipping to related addresses,
- collusive seller or review networks,
- abuse rings coordinating promotions or refunds.

### 11.9.2 Graph features and graph ML

Practical graph-based approaches may include:

- connected-component and degree features,
- shared-entity counts,
- neighborhood risk propagation,
- graph neural networks where scale and maturity justify them.

Graph analysis is powerful because fraudsters often try to hide in transactional noise while still reusing infrastructure or identity fragments.

## 11.10 Hybrid rules + ML systems

In production e-commerce systems, the strongest approach is often hybrid rather than purely rule-based or purely model-based.

### 11.10.1 Why hybrids work

Rules are good for hard policy constraints and urgent pattern blocking. ML is good for scoring complex combinations of soft signals. Combined systems can:

- enforce mandatory blocks,
- assign risk scores,
- route borderline cases to review,
- and learn from review outcomes.

### 11.10.2 Example decision flow

A practical decision flow may look like:

1. apply deterministic policy rules,
2. compute risk score from supervised and anomaly signals,
3. adjust using graph or network intelligence,
4. compare against thresholds for approve, review, or reject,
5. log explanations and outcomes for feedback.

## 11.11 Thresholds and decision economics

Fraud systems do not end at model scoring. The business must choose action thresholds.

### 11.11.1 Approve, review, reject

Many systems use three zones:

- **Approve:** low-risk transactions proceed automatically.
- **Review:** medium-risk cases go to human review or secondary checks.
- **Reject or challenge:** high-risk cases are blocked or subjected to stronger verification.

### 11.11.2 Cost-sensitive tuning

Thresholds should reflect:

- fraud loss,
- customer lifetime value,
- review-team capacity,
- conversion impact,
- chargeback and compliance costs.

This means threshold selection is not a pure ML exercise. It is an economic operating decision.

## 11.12 Human-in-the-loop review

Not all fraud decisions should be fully automated. Human review remains important for edge cases, novel fraud patterns, appeals, and policy-sensitive decisions.

### 11.12.1 What reviewers need

Review tools should expose:

- transaction history,
- device and network context,
- linked entities,
- model score and contributing signals,
- policy reason codes,
- prior outcomes on related entities.

### 11.12.2 Feedback loop

Review outcomes should feed back into:

- rule refinement,
- model retraining,
- investigation workflows,
- reporting on false positives and operational load.

## 11.13 Fraud evaluation metrics

Accuracy alone is misleading in fraud detection because the class distribution is highly imbalanced.

### 11.13.1 Useful metrics

Practical evaluation often includes:

- precision,
- recall,
- precision-recall curves,
- ROC-AUC with caution,
- false-positive rate,
- fraud-dollar capture,
- review rate,
- cost-weighted utility.

### 11.13.2 Business interpretation

A model with strong recall but excessive false positives may block good customers. A model with low review cost but weak fraud capture may lose large sums. Evaluation must therefore align with operational and economic goals.

## 11.14 Drift, adaptation, and monitoring

Fraud systems need continuous monitoring because attacker behavior evolves.

### 11.14.1 What to monitor

- score distribution shifts,
- feature drift,
- approval and decline rates,
- review queue volume,
- chargeback lag trends,
- abuse-type mix over time.



### 11.14.2 Why monitoring is essential

Even a well-trained model can degrade if:

- fraud patterns shift,
- product flows change,
- new promotions alter normal behavior,
- identity or logging systems are modified.

Fraud detection is therefore inseparable from MLOps and operational analytics.

## 11.15 Trust & safety beyond payment fraud

Trust & safety is broader than payment authorization. It includes keeping the entire commerce environment healthy.

### 11.15.1 Examples beyond transactions

- review authenticity,
- marketplace seller compliance,
- abuse of support or refund workflows,
- harmful or prohibited catalog content,
- bot-driven scraping or stock hoarding.

### 11.15.2 Why this matters architecturally

These problems span domains. A trust & safety capability often depends on:

- commerce events,
- ERP refund records,
- CRM support interactions,
- identity and authentication systems,
- marketplace governance workflows.

This reinforces the enterprise view that trust & safety is an operating system for commerce, not a single checkout model.

## 11.16 Hands-on lab: risk scoring and thresholding

The practical goal of this chapter is to connect fraud modeling with decision policy.

### 11.16.1 Lab scenario

Students are given a synthetic transaction dataset containing payment, account, device, and label information, plus estimated fraud-loss and false-positive costs.

### 11.16.2 Lab tasks

1. Build a baseline risk model using structured transaction and account features.
2. Compute precision, recall, and a cost-sensitive evaluation summary.
3. Define approve, review, and reject thresholds.
4. Propose one anomaly signal and one graph-style linkage feature to improve the system.
5. Explain how a human review queue would be integrated into the workflow.

### 11.16.3 Expected learning outcome

By the end of the lab, students should be able to explain fraud detection as a staged decision system rather than a single prediction score.

## 11.17 Managerial implications

Fraud and trust & safety systems require balancing protection with customer experience. Overly strict controls reduce conversion and frustrate legitimate users. Overly weak controls invite direct loss and operational abuse.

The best systems therefore combine:

- clear policy,
- strong data quality,
- hybrid detection methods,
- review operations,
- and continuous monitoring.

Organizations that treat fraud as an isolated technical model usually underinvest in the workflow and governance layers that actually determine business outcomes.

## 11.18 Multiple-choice questions (MCQs)

1. Why is fraud detection in e-commerce often difficult from an ML perspective?
  - (a) Because fraud labels are always complete and immediate.
  - (b) Because the problem is highly imbalanced, adversarial, and often affected by delayed labels.
  - (c) Because fraud can only be solved with static rules.
  - (d) Because customer behavior never changes.
2. Which is the best example of *promotion abuse*?
  - (a) A customer browsing many products before purchase.
  - (b) Creating many accounts or identities to repeatedly exploit a first-order coupon.
  - (c) A warehouse updating shipment status.
  - (d) A user selecting a shipping address.
3. Why are graph-based methods useful in fraud detection?
  - (a) Because fraud is often relational and linked through shared devices, addresses, accounts, or payment instruments.
  - (b) Because graphs eliminate the need for labels.
  - (c) Because they only work for product recommendation.
  - (d) Because they remove the need for monitoring.
4. In a hybrid fraud system, rules are especially useful for:
  - (a) hard policy enforcement and fast blocking of known bad patterns.
  - (b) replacing all human review.
  - (c) guaranteeing perfect recall.
  - (d) making embeddings unnecessary in every domain.
5. Why are thresholds in fraud systems business decisions as well as technical decisions?
  - (a) Because thresholds determine the tradeoff between fraud loss, false positives, review cost, and customer experience.
  - (b) Because thresholds only affect documentation style.
  - (c) Because all thresholds are set randomly.
  - (d) Because supervised learning does not produce scores.

### Answer key

1. (b)
2. (b)
3. (a)
4. (a)
5. (a)

### 11.19 Exercises (short)

1. Compare rule-based detection, supervised learning, and anomaly detection for one fraud scenario of your choice. State one advantage and one limitation of each.
2. Design a minimal feature set for detecting account takeover. Include account, device, sequence, and review-history signals.
3. Define a three-zone approve/review/reject threshold policy for a checkout fraud model and explain how review-team capacity changes your choice.

### 11.20 Mini-case (Odoo-linked, vendor-neutral)

**Scenario:** A hybrid commerce company uses a headless storefront, Odoo for orders and refunds, and a central data platform. The business is seeing rising chargebacks, coupon abuse, and suspicious refund patterns.

**Task:** Design a trust & safety workflow:

- Define which events and records must be linked across storefront, payment, Odoo, and support systems.
- Propose a hybrid detection design using rules, ML scoring, and at least one graph-style signal.
- Specify the criteria for automatic approval, manual review, and automatic rejection in one high-risk flow such as refund approval or checkout authorization.

## Chapter 12

# Customer Service Automation with LLMs (Week 11)

### 12.1 Learning objectives

- Explain how LLMs change the design space of customer service automation in e-commerce.
- Distinguish when to use rules, classical NLP, retrieval, and LLM-based systems.
- Design a retrieval-augmented generation (RAG) assistant with tool use, workflow orchestration, and escalation paths.
- Evaluate LLM service systems using operational, quality, and risk-oriented criteria rather than only surface fluency.

### 12.2 Why customer service is a natural LLM use case

Customer service in e-commerce involves large volumes of language-centered work:

- answering policy questions,
- helping customers track orders,
- explaining returns and refunds,
- resolving account issues,
- supporting product questions,
- handling multilingual or multi-channel interactions.

Traditional service automation relied on rules, scripted chat flows, FAQ search, and intent classification. These tools are still useful, but LLMs change what is possible by supporting more flexible language understanding and generation.

However, flexibility is not enough. Customer support is operationally sensitive. A fluent but incorrect answer about order status, refund policy, or account action can create real business harm. This means LLM systems must be designed as controlled service workflows, not just chat interfaces.

## 12.3 Customer service tasks in e-commerce

Not all service tasks are equally suited to LLMS. An architectural view should distinguish several task types.

### 12.3.1 Informational tasks

Examples include:

- store policy questions,
- shipping timelines,
- payment methods,
- warranty coverage,
- how to initiate a return.

These are often strong candidates for retrieval-augmented LLM assistance because the main challenge is interpreting and presenting known information.

### 12.3.2 Transactional tasks

Examples include:

- checking order status,
- canceling an order,
- changing a shipping address,
- opening a support case,
- issuing a return label.

These tasks require interaction with operational systems. The LLM cannot safely invent outcomes. It must use tools or APIs to read or modify real system state.

### 12.3.3 Judgment-heavy tasks

Examples include:

- refund exceptions,
- suspicious behavior review,
- policy edge cases,
- emotionally complex complaints,
- high-value account escalations.

These tasks often require human oversight even if an LLM assists with summarization, drafting, or routing.

## 12.4 Rules, classical NLP, and LLMS

A common mistake is to assume that LLMS should replace all earlier automation methods. That is rarely the correct design.

### 12.4.1 When rules are best

Rules remain useful when:

- the workflow is deterministic,
- the action must be tightly controlled,
- compliance requires explicit logic,
- the trigger patterns are narrow and stable.

Examples include authentication checks, eligibility thresholds, or exact routing conditions.

### 12.4.2 When classical NLP may still be enough

Classical NLP or simple classifiers can still be appropriate for:

- intent routing,
- spam detection,
- language identification,
- sentiment tagging,
- FAQ ranking baselines.

### 12.4.3 When LLMs are most valuable

LLMs are strongest when the task requires:

- flexible language interpretation,
- synthesis from multiple knowledge sources,
- paraphrasing and explanation,
- conversation continuity,
- summarization of long contexts,
- structured tool selection under guidance.

The right architecture therefore combines all three approaches rather than treating them as mutually exclusive.

## 12.5 Retrieval-augmented generation (RAG)

RAG is one of the most important patterns for customer service with LLMs. Instead of asking the model to rely on parametric memory alone, the system retrieves relevant documents or records and asks the model to ground its answer in them.

### 12.5.1 Why RAG matters

RAG helps because:

- support knowledge changes frequently,
- policy and catalog information must remain current,
- grounded answers are easier to audit,
- hallucination risk is reduced when good retrieval is present.

### 12.5.2 Knowledge sources

Typical RAG sources in e-commerce support include:

- return and shipping policies,
- help-center content,
- seller or marketplace rules,
- product documentation,



- order and shipment records,
- CRM case history,
- internal support playbooks.

### 12.5.3 Retrieval quality

RAG only works well if retrieval quality is strong. If the wrong policy document or stale order record is retrieved, the LLM may produce a polished but misleading answer. This is why retrieval architecture matters as much as model prompting.

## 12.6 Tool use and operational actions

Customer service systems often need more than retrieved documents. They need access to operational tools.

### 12.6.1 Examples of service tools

- get order status,
- get shipment tracking,
- verify return eligibility,
- create support ticket,
- generate return label,
- check account verification state,
- retrieve CRM notes.

### 12.6.2 Why tools are necessary

Without tools, the LLM can only talk *about* the workflow. With tools, it can participate in the workflow while remaining grounded in current system state.

### 12.6.3 Action control

Not all tools should be available with the same level of autonomy. Some are read-only and low risk. Others change customer state and should require:

- authentication confirmation,
- policy checks,

- explicit user confirmation,
- or human approval.

## 12.7 Workflow orchestration

A customer service assistant is rarely just a single prompt to a model. It is more often a workflow with distinct steps.

### 12.7.1 Typical service flow

A service flow may include:

1. classify the issue or infer likely task,
2. retrieve relevant policy or account context,
3. call operational tools if needed,
4. generate a grounded response,
5. decide whether to resolve, request clarification, or escalate,
6. log the interaction and outcome.

### 12.7.2 Why orchestration matters

Without orchestration, an LLM assistant may:

- skip required policy checks,
- fail to collect missing information,
- call tools unnecessarily,
- give inconsistent answers across turns,
- or fail to escalate when confidence is low.

Orchestration is therefore what turns a language model into a service system.

## 12.8 Conversation memory and context

Customer support conversations often span multiple turns and sometimes multiple channels. The system must decide what context to preserve.

### 12.8.1 Useful forms of context

- current conversation history,
- prior support interactions,
- customer account tier,
- recent orders and return events,
- prior escalations or unresolved issues.

### 12.8.2 Context hygiene

More context is not automatically better. The system should avoid:

- pulling irrelevant records,
- mixing incompatible customer identities,
- feeding stale or sensitive data without justification,
- overloading prompts with low-signal content.

Good service design uses enough context to ground the answer without degrading precision or privacy.

## 12.9 Grounding, citations, and answer control

When the system answers policy or status questions, it should ideally show the basis for the answer.

### 12.9.1 Why citations help

Citations or source references can:

- increase operator trust,
- help reviewers audit output,
- make it easier to detect retrieval failures,
- support handoff to human agents.

### 12.9.2 Controlled answer styles

Service responses often need formatting and policy constraints such as:

- concise answer first,
- explicit next step,
- no unsupported promises,
- no invented refund or delivery guarantees,
- handoff language when the policy is uncertain.

This is one reason prompt design alone is insufficient. The surrounding system must constrain what the model is allowed to say and do.

## 12.10 Escalation and human handoff

An LLM service assistant should not aim to resolve every case autonomously. Escalation is a core capability, not a fallback afterthought.

### 12.10.1 Escalation triggers

Escalation may be appropriate when:

- the policy is ambiguous,
- the required action is high risk,
- the customer is frustrated or high value,
- fraud or compliance concerns are present,
- the model lacks sufficient evidence,
- repeated attempts have failed.

### 12.10.2 Good handoff design

A good handoff should include:

- summary of customer intent,
- relevant retrieved documents,
- tool outputs already collected,

- recommended next action,
- uncertainty or policy flags.

This can reduce agent workload while preserving human accountability.

## 12.11 Evaluation of LLM service systems

Evaluation must go beyond subjective fluency. In customer service, a helpful-sounding answer can still be wrong or risky.

### 12.11.1 Evaluation dimensions

Useful dimensions include:

- factual correctness,
- policy compliance,
- grounding quality,
- task completion success,
- escalation appropriateness,
- customer satisfaction,
- average handling time,
- hallucination or unsafe-action rate.

### 12.11.2 Offline versus online evaluation

Offline evaluation may use test conversations, policy QA sets, or synthetic workflows. Online evaluation may track:

- containment rate,
- escalation rate,
- resolution quality,
- recontact rate,
- CSAT,
- and downstream business costs.

### 12.11.3 Human review in evaluation

Because support quality is nuanced, human review remains important. Reviewers can assess:

- whether the response was grounded,
- whether the tone was appropriate,
- whether the assistant acted within policy,
- whether escalation should have occurred sooner.

## 12.12 Guardrails and safety

Customer service assistants interact with customers directly, so safety controls are essential.

### 12.12.1 Common risk areas

- fabricated order or refund status,
- privacy leakage across customer accounts,
- unsafe or unauthorized operational actions,
- policy violations,
- manipulative or inappropriate tone,
- prompt injection through retrieved content or user messages.

### 12.12.2 Guardrail mechanisms

Typical guardrails include:

- scoped retrieval and identity verification,
- tool permissions by action type,
- policy filters and post-generation checks,
- refusal or defer behavior under uncertainty,
- human approval for sensitive workflows,
- full logging and audit trails.

Guardrails should be treated as first-class design components, not optional moderation add-ons.

## 12.13 Prompt injection and retrieval risks

When the assistant uses retrieved documents or conversation text, it can be exposed to malicious or misleading content.

### 12.13.1 Examples of injection risk

- a seller document includes hidden instructions not meant for the assistant,
- a support message tries to override policy behavior,
- external or user-generated content contains adversarial text,
- stale or conflicting documents appear in retrieval results.

### 12.13.2 Mitigations

Mitigation strategies include:

- restricting retrieval sources,
- separating system instructions from retrieved content,
- validating tool actions independently of model output,
- ranking and curating trusted knowledge bases.

## 12.14 A practical architecture for LLM support

A practical LLM support stack often includes:

- a conversation layer,
- retrieval over curated support knowledge,
- connectors to order, CRM, and return systems,
- policy-aware orchestration,
- evaluation and monitoring,
- escalation into human service tools.

### 12.14.1 Read and write separation

It is often useful to separate:

- **read tools:** retrieve order status, policy, case history,
- **write tools:** cancel order, create refund request, open ticket.

Write tools usually deserve stricter safeguards because they change operational state.

### 12.14.2 Role of CRM and ERP integration

An LLM support system may need:

- CRM for case history and customer context,
- ERP or commerce systems for order and refund truth,
- warehouse or shipping systems for delivery state,
- policy repositories for standardized answers.

This means the assistant is only as good as its enterprise integration design.

## 12.15 Hands-on lab: a grounded FAQ and service assistant

The practical goal of this chapter is to move from generic chat to controlled support automation.

### 12.15.1 Lab scenario

Students are given a small local knowledge base containing return policy, shipping policy, warranty terms, and a mock order-status data source.

### 12.15.2 Lab tasks

1. Build a retrieval-based assistant over the local policy documents.
2. Add source citations to every policy answer.
3. Add one read-only tool such as mock order-status lookup.
4. Define escalation rules for low-confidence or high-risk requests.
5. Design an evaluation rubric for correctness, grounding, and escalation behavior.



### 12.15.3 Expected learning outcome

By the end of the lab, students should be able to explain why a useful LLM support assistant is a policy-aware workflow, not just a conversational model.

## 12.16 Managerial implications

LLM-based service systems can reduce repetitive workload and improve responsiveness, but only if they are embedded in disciplined operational processes. Organizations should resist the temptation to optimize for containment rate alone. A service assistant that resolves many tickets badly can be more damaging than one that escalates appropriately.

The correct managerial objective is controlled leverage:

- automate routine, well-grounded service work,
- support human agents on harder cases,
- maintain clear accountability for sensitive actions,
- and measure quality continuously.

## 12.17 Multiple-choice questions (MCQs)

1. Why is retrieval-augmented generation (RAG) especially valuable for e-commerce customer service?
  - (a) Because it removes the need for any support content.
  - (b) Because it helps ground answers in current policies, records, and support knowledge rather than relying only on model memory.
  - (c) Because it guarantees every answer is legally correct.
  - (d) Because it eliminates escalation.
2. Which is the best example of a task that usually requires tool use rather than pure text generation?
  - (a) Summarizing a return policy.
  - (b) Checking real-time order status for a specific customer.
  - (c) Rewriting a help-center article title.
  - (d) Translating a greeting.
3. Why should escalation be treated as a core feature in LLM customer service systems?
  - (a) Because the goal is to avoid any automated resolution.
  - (b) Because some requests are ambiguous, sensitive, or high risk and require human review or approval.

- (c) Because escalation improves embedding speed.
  - (d) Because rules cannot coexist with LLMs.
4. Which is the best example of a support-system guardrail?
- (a) Allowing the model to invent missing refund policies to sound helpful.
  - (b) Requiring grounded sources, scoped tool permissions, and human approval for sensitive actions.
  - (c) Removing all logs to protect latency.
  - (d) Hiding all uncertainty from the customer.
5. Why is fluency alone an insufficient evaluation criterion for LLM service assistants?
- (a) Because customer service is only about grammar.
  - (b) Because a response can sound natural while still being wrong, ungrounded, or policy-violating.
  - (c) Because evaluation is unnecessary after deployment.
  - (d) Because retrieval systems cannot be tested offline.

### Answer key

- 1. (b)
- 2. (b)
- 3. (b)
- 4. (b)
- 5. (b)

## 12.18 Exercises (short)

- 1. Compare a rule-based FAQ system, a classical intent-routing system, and an LLM+RAG assistant for customer support. State one strength and one limitation of each.
- 2. Design a tool schema for a read-only order-status lookup that could be safely used by a support assistant. Include required inputs, expected outputs, and at least two safety checks.
- 3. Propose an escalation policy for refund-related requests. Define when the assistant can answer directly, when it can create a request, and when it must hand off to a human agent.

## 12.19 Mini-case (Odoo-linked, vendor-neutral)

**Scenario:** A hybrid B2C/B2B company uses a headless storefront, Odoo for orders and invoices, a CRM for support history, and a local knowledge base for policies. The company wants an LLM-based support assistant for chat and agent assistance.

**Task:** Design the service architecture:

- Define which data should be retrieved from policy documents, which should come from tools, and which should remain human-only.
- Explain how the assistant would handle order-status questions, return eligibility questions, and refund exceptions.
- Identify two guardrails and two evaluation metrics that must be in place before broad deployment.



## Chapter 13

# MLOps + DataOps for E-Commerce (Week 12)

### 13.1 Learning objectives

- Explain why production ML in e-commerce requires disciplined operational practices beyond model training.
- Distinguish MLOps concerns from DataOps concerns while understanding their overlap in end-to-end commerce systems.
- Analyze training-serving skew, feature consistency, monitoring, drift, and deployment risk in production ML workflows.
- Design an ML system artifact set that includes reproducibility, observability, and governance requirements.

### 13.2 Why MLOps and DataOps matter in e-commerce

E-commerce organizations often move quickly from a notebook proof-of-concept to a production dependency. Models may begin as experiments for recommendation, fraud, pricing, or service automation, but once deployed they affect revenue, customer experience, and operational load.

At that point, the technical question is no longer “does the model score well offline?” It becomes:

- How is the model trained repeatedly?
- How are features computed consistently?
- How is the model deployed safely?
- How is degradation detected?
- How is the system rolled back when behavior becomes harmful?

MLOps addresses the operational lifecycle of models. DataOps addresses the operational lifecycle of the data systems those models depend on. In e-commerce, the two are tightly coupled because behavior, catalog state, pricing, inventory, and fraud patterns all change continuously.

## 13.3 From experimentation to production systems

Many teams can build a model. Far fewer can operate one reliably.

### 13.3.1 The production gap

The production gap often appears when a promising model meets real constraints such as:

- missing or delayed features,
- inconsistent identifiers across systems,
- changes in business process that alter label meaning,
- latency requirements,
- lack of rollback or approval workflow,
- weak monitoring once the model is live.

### 13.3.2 Why e-commerce is especially operational

E-commerce ML systems are unusually exposed to changing conditions:

- product assortments change,
- promotions alter behavior,
- demand shifts seasonally,
- fraudsters adapt,
- support policies evolve,
- customers interact across multiple channels.

This means that a model that looked strong last month may become weak or unsafe if the surrounding system changes.

## 13.4 What MLOps covers

MLOps is the discipline of making ML systems buildable, deployable, observable, and governable.

### 13.4.1 Typical MLOps responsibilities

- versioning code, data references, and model artifacts,
- training pipelines,
- model packaging and deployment,
- experiment tracking,
- model registry and release control,
- performance monitoring,
- rollback and incident response.

### 13.4.2 What DataOps covers

DataOps focuses more on:

- data pipeline reliability,
- data contracts and validation,
- data freshness and lineage,
- reproducible transformations,
- quality checks across analytical and operational data movement.

In practice, production ML depends on both. Poor DataOps can silently break a model. Poor MLOps can deploy a model without control or visibility.

## 13.5 Reproducibility

Reproducibility is a foundation of trustworthy ML operations. If the team cannot explain how a model was trained and with which data, governance and debugging become weak.

### 13.5.1 What should be reproducible

A strong production workflow should make it possible to recover:

- source code version,
- feature definitions,
- training data snapshot or reference,
- hyperparameters,

- evaluation outputs,
- model artifact,
- deployment configuration.

### 13.5.2 Why this matters in commerce

When a recommender starts pushing low-margin items, or a fraud model begins over-blocking legitimate customers, teams need to determine whether the cause was:

- new data,
- a model change,
- a feature transformation bug,
- a serving mismatch,
- or a business-rule change upstream.

Without reproducibility, that diagnosis becomes slow and politically difficult.

## 13.6 Pipelines and orchestration

Production ML systems depend on pipelines rather than ad hoc manual execution.

### 13.6.1 Typical pipeline stages

A simplified pipeline may include:

1. ingest and validate raw data,
2. transform features,
3. train models,
4. evaluate and compare candidates,
5. register approved artifacts,
6. deploy to staging or production,
7. monitor outcomes and trigger retraining or rollback when needed.



### 13.6.2 Batch and streaming coexistence

E-commerce ML often requires both:

- batch pipelines for training and aggregate features,
- streaming or near-real-time pipelines for serving-time context such as session behavior, current inventory, or recent fraud indicators.

The system design must make clear which features belong to each path.

## 13.7 Feature consistency and training-serving skew

One of the most important production risks in ML systems is training-serving skew. This occurs when the features or transformations used at serving time do not match those used in training.

### 13.7.1 How skew happens

Typical causes include:

- using different code paths offline and online,
- stale or delayed online feature values,
- inconsistent joins or identifiers,
- business-rule changes applied only in one environment,
- time-window leakage in offline training that is impossible online.

### 13.7.2 Why skew is dangerous

Skew can degrade a model even when the model artifact itself is correct. For example:

- a fraud model may receive delayed account-age features online,
- a recommender may use a stale inventory flag,
- a pricing model may be trained on clean margin data but served with inconsistent cost feeds.

### 13.7.3 Feature stores as a concept

Feature stores are often introduced to reduce this risk by standardizing feature definitions and access patterns across training and serving. The technology choice matters less than the architectural goal: one governed definition of each important feature.

## 13.8 Model registry and release management

Production ML systems need a controlled way to manage candidate and approved models.

### 13.8.1 Why a registry matters

A model registry or equivalent control plane helps track:

- which model versions exist,
- which one is approved,
- which environment each version is deployed in,
- what evaluation evidence supported release,
- how to roll back if needed.

### 13.8.2 Release strategies

Common release strategies include:

- shadow deployment,
- canary rollout,
- staged traffic ramp-up,
- champion-challenger comparison,
- full rollback on guardrail breach.

These practices are especially useful in e-commerce where small degradations can have immediate commercial consequences.

## 13.9 Monitoring production ML systems

Monitoring must cover more than CPU usage and uptime. A live model can be operationally healthy and still commercially harmful.

### 13.9.1 Technical monitoring

Useful technical signals include:

- latency,

- throughput,
- error rate,
- feature freshness,
- missing-value spikes,
- tool or dependency failures in inference paths.

### 13.9.2 Model monitoring

Useful model-focused signals include:

- score distribution shifts,
- calibration changes,
- prediction drift,
- confidence anomalies,
- slice-level performance changes.

### 13.9.3 Business monitoring

Production evaluation should also include business outcomes such as:

- click-through and conversion for recommenders,
- fraud-dollar capture and false-positive burden,
- margin and inventory effects for pricing models,
- containment and recontact rate for service assistants.

This is what turns monitoring from infrastructure tracking into operational accountability.

## 13.10 Data drift and concept drift

Drift is unavoidable in e-commerce ML. Students should understand several forms of drift.

### 13.10.1 Data drift

Data drift occurs when the distribution of inputs changes. Examples include:

- seasonal traffic changes,
- new product categories,
- device mix shifts,
- changed promotion patterns,
- new fraud tactics.

### 13.10.2 Concept drift

Concept drift occurs when the relationship between inputs and outcomes changes. For example:

- the same browsing behavior stops indicating purchase intent after a UI redesign,
- fraud patterns evolve so past indicators weaken,
- pricing response changes because a competitor enters the market.

### 13.10.3 Response to drift

Possible responses include:

- alerting,
- retraining,
- threshold adjustment,
- feature revision,
- temporary fallback to safer baselines.

## 13.11 CI/CD for ML

Continuous integration and deployment for ML differs from classical software delivery because the deployed artifact is influenced by data as well as code.

### 13.11.1 What CI for ML should validate

Reasonable validation steps include:

- unit tests for transformations,
- schema and feature checks,
- reproducibility checks,
- training pipeline success,
- offline evaluation against baselines,
- safety or policy checks for sensitive outputs.

### 13.11.2 What CD for ML should consider

Deployment should consider:

- compatibility with serving infrastructure,
- readiness of online features,
- staged rollout strategy,
- alert thresholds,
- rollback readiness.

In e-commerce, deployment speed matters, but so does release discipline.

## 13.12 Experiment tracking and model comparison

Teams need a reliable record of what was tried and what worked.

### 13.12.1 What to track

Useful tracked metadata include:

- dataset reference,
- feature set version,
- model type,
- hyperparameters,

- training date and window,
- evaluation metrics,
- release decision.

### 13.12.2 Why this matters

Without structured experiment tracking, teams can neither compare models honestly nor explain why a deployed model was chosen. This becomes particularly important when several business units or use cases share the same ML platform.

## 13.13 Governance and approvals

Not every model should be promoted automatically. Some e-commerce use cases have a high enough impact that governance gates are necessary.

### 13.13.1 Examples of governed models

- fraud and abuse scoring,
- dynamic pricing,
- customer-service agents with action tools,
- high-visibility recommendation systems,
- risk or compliance-related models.

### 13.13.2 Typical governance artifacts

Reasonable artifacts may include:

- model cards,
- risk assessment notes,
- monitoring plans,
- evaluation summaries,
- rollback criteria,
- ownership and escalation paths.

Governance is not meant to slow all experimentation equally. It is meant to scale decision quality for systems that matter.

## 13.14 Incident response for ML systems

Production ML systems need incident handling just like other production services.

### 13.14.1 Examples of ML incidents

- a recommender starts surfacing unavailable items,
- a fraud model causes a spike in false declines,
- a service assistant cites the wrong policy version,
- a pricing model recommends below-cost prices after a cost feed failure.

### 13.14.2 Preparedness

Good preparedness includes:

- clear runbooks,
- fallback modes,
- alert ownership,
- frozen-safe baselines,
- audit logs for recent model and data changes.

This is one reason operational simplicity is valuable. Complex ML systems without response plans tend to fail expensively.

## 13.15 Artifacts for this chapter

The central artifact deliverable for this chapter is an ML system design package containing:

- pipeline overview,
- feature and data dependencies,
- deployment path,
- monitoring plan,
- drift and rollback strategy,
- governance and ownership summary.

## 13.16 Hands-on artifacts

Students should produce:

- an ML system design document,
- a monitoring plan with technical, model, and business signals,
- and a short incident playbook for one selected use case.

## 13.17 Managerial implications

Organizations often underestimate the operational cost of ML. The model may be only a small fraction of the total system effort. Most production pain comes from weak data contracts, poor observability, inconsistent features, and unclear ownership.

Well-run MLOps and DataOps allow the business to benefit from ML repeatedly rather than episodically. They convert isolated model wins into an operational capability.

## 13.18 Multiple-choice questions (MCQs)

1. What is *training-serving skew*?
  - (a) A difference between the fonts used in notebooks and dashboards.
  - (b) A mismatch between how features or transformations are produced in training and how they are produced at serving time.
  - (c) A pricing discount given only during model training.
  - (d) A mandatory property of all neural networks.
2. Why is a model registry or equivalent release control useful?
  - (a) It guarantees no model will ever fail.
  - (b) It helps track approved model versions, deployment state, evaluation evidence, and rollback options.
  - (c) It replaces monitoring entirely.
  - (d) It only matters for academic experiments.
3. Which is the best example of a *business* monitoring signal for an e-commerce ML system?
  - (a) CPU temperature
  - (b) Fraud-dollar capture or revenue-per-session lift
  - (c) Container image size
  - (d) Disk inode count
4. Data drift differs from concept drift because:



- (a) data drift is about input-distribution change, while concept drift is about change in the relationship between inputs and outcomes.
  - (b) concept drift is only about network outages.
  - (c) data drift only happens in recommendation systems.
  - (d) there is no real difference in practice.
5. Why do production ML systems need incident playbooks?
- (a) Because model failures can affect real business outcomes and require fast, structured response.
  - (b) Because playbooks improve offline accuracy.
  - (c) Because incidents only happen in ERP systems.
  - (d) Because monitoring removes the need for rollback.

**Answer key**

- 1. (b)
- 2. (b)
- 3. (b)
- 4. (a)
- 5. (a)



## Chapter 14

# Security, Privacy, Compliance, and Responsible AI (Week 13)

### 14.1 Learning objectives

By the end of this chapter, students should be able to:

- explain how security, privacy, compliance, and responsible AI interact in e-commerce platforms;
- identify major risks across payments, customer accounts, APIs, data pipelines, models, and LLM-enabled interfaces;
- apply practical control patterns such as least privilege, data minimization, encryption, audit logging, and human oversight;
- distinguish between technical controls, process controls, and governance controls;
- design a risk assessment and operating model for AI features used in commerce.

### 14.2 Why this chapter matters

Digital commerce systems sit at the intersection of money, identity, behavioral data, operational processes, and automated decision-making. That combination makes them attractive to attackers and highly visible to regulators. A mature commerce organization therefore does not treat security as a narrow infrastructure problem, or privacy as a legal afterthought, or responsible AI as a public-relations layer. These are architectural concerns that shape how systems are modeled, integrated, deployed, monitored, and governed.

The stakes are especially high when AI is introduced into search, recommendation, pricing, fraud scoring, customer service, and content generation. These systems can scale decisions quickly, but they can also amplify mistakes quickly. A poorly secured recommendation service can expose user data. A careless LLM assistant can leak sensitive information. A biased risk model can block legitimate customers. A weak retention policy can keep personal data longer than justified. In

production environments, these are not isolated failures. They often cascade across teams, vendors, and business processes.

### 14.3 Security as a business capability

Security in e-commerce is often discussed in technical language, but managers experience it through business outcomes: uptime, payment acceptance, customer trust, regulatory exposure, and loss prevention. A security program therefore needs to be framed as a business capability with clear ownership, metrics, and operating procedures.

At a minimum, the enterprise must protect:

- customer identities, credentials, sessions, and account recovery channels;
- payment data, refund flows, stored tokens, and settlement operations;
- order data, addresses, invoices, and operational records;
- APIs, integrations, partner connections, and event streams;
- training data, feature tables, prompts, model endpoints, and evaluation artifacts;
- internal administrative tools used by support, finance, merchandising, and operations.

Security should be understood as defense in depth. There is no single control that makes a commerce platform safe. The real objective is to create overlapping layers so that one control failure does not automatically become a business crisis.

### 14.4 Threat modeling for e-commerce platforms

Threat modeling is the discipline of asking what can go wrong before an incident proves it the hard way. In commerce systems, threat modeling is more useful when it is tied to business flows rather than abstract infrastructure diagrams. A practical starting point is to map critical journeys such as sign-up, login, browse, cart, checkout, payment, refund, support interaction, and seller onboarding.

For each journey, teams should ask:

- what assets are touched;
- which actors can initiate or alter the flow;
- what trust boundaries are crossed;
- what external services or vendors are involved;
- which failures would lead to financial loss, service disruption, data leakage, or compliance exposure.

Common threat categories in commerce include account takeover, credential stuffing, card testing, promo abuse, API scraping, inventory manipulation, refund abuse, business email compromise, insider misuse, and denial of service. AI-enabled components introduce additional risks such as prompt injection, sensitive context leakage, unsafe tool use, training data poisoning, and automated decision errors at scale.

#### 14.4.1 Threat modeling example: intelligent customer support

Consider an LLM-powered support assistant with retrieval over order history, policy documents, and customer profile data. This system may expose several attack paths:

- a user attempts to retrieve another customer's order details through prompt manipulation;
- a malicious knowledge-base document contains adversarial instructions that alter model behavior;
- the assistant triggers refund or coupon actions without sufficient verification;
- logs capture sensitive personal data and become accessible to staff who do not need it;
- agents over-trust generated responses and execute harmful actions without review.

This example shows why security and responsible AI must be co-designed. Traditional application controls are still required, but they are no longer sufficient by themselves.

### 14.5 Identity, authentication, and access control

Most major incidents in digital systems involve identity in some form. Attackers rarely begin by breaking encryption; they usually exploit weak authentication, poor session handling, excessive privileges, or flawed recovery flows.

For customer-facing systems, core controls include:

- strong password policies or passwordless flows;
- multi-factor authentication for high-risk actions;
- rate limiting and bot mitigation for login and recovery endpoints;
- device, session, and anomaly-based risk checks;
- secure handling of tokens, cookies, and session expiration.

For internal systems, the priority is least privilege. Staff should access only the data and actions necessary for their role. Support agents may need order history, but not full payment details. Merchandising teams may update catalog content, but not financial ledgers. Data scientists may query curated analytical datasets, but not unrestricted production tables with raw personal data.

Role-based access control is often a useful starting point, but large enterprises usually need finer-grained policies. Attribute-based rules, approval workflows, just-in-time access, and strong audit logs become important as the operating environment grows more complex.

## 14.6 Payment security and transactional integrity

Payments are a special case because they combine fraud risk, compliance obligations, vendor dependencies, and customer trust. Even when a firm outsources payment processing to a specialized provider, it still owns the architecture around payment initiation, status handling, refunds, dispute workflows, and reconciliation.

Key design principles include:

- minimize the platform's direct exposure to sensitive card data;
- tokenize wherever possible rather than storing primary account details;
- isolate payment services from broader application concerns;
- treat refund and chargeback processes as security-sensitive workflows;
- reconcile payment state carefully across commerce, ERP, and finance systems.

Transactional integrity matters beyond card data. If an order is captured twice, a refund is processed without verification, or a payment callback is replayed, the platform can suffer both operational and reputational harm. Idempotency keys, signed callbacks, event replay protection, and well-defined state transitions are therefore central security mechanisms, not merely engineering niceties.

## 14.7 API and integration security

Modern commerce platforms depend heavily on APIs and integrations: payment gateways, tax engines, shipping carriers, marketplaces, ERP systems, recommendation services, identity platforms, and analytics pipelines. Every integration is also a security boundary.

Secure integration design typically requires:

- strong service authentication and secret management;
- explicit authorization scopes for each API consumer;
- request signing or mutual trust mechanisms for sensitive channels;
- replay protection, schema validation, and rate limiting;
- network segmentation and environment isolation;
- monitoring of abnormal patterns in partner or service-to-service traffic.

The hardest integration problems are often organizational rather than cryptographic. Teams forget who owns a credential, fail to rotate keys, grant overly broad access to vendor systems, or let undocumented data exports proliferate. Mature integration security requires inventory discipline: every data flow, secret, endpoint, and external dependency should have an owner.

## 14.8 Data protection and privacy-by-design

Privacy-by-design means that systems are structured to reduce unnecessary data exposure before policies are written or incidents occur. The goal is not to eliminate data use. Commerce organizations need data for fulfillment, analytics, personalization, forecasting, and service quality. The goal is to align data collection and processing with clear purpose, proportionality, retention, and control.

Core privacy design principles include:

- collect only the data required for a justified purpose;
- separate operational necessity from analytical convenience;
- define retention windows and deletion mechanisms early;
- avoid copying personal data into uncontrolled spreadsheets and side systems;
- document lawful basis, consent dependencies, and processing purpose where relevant;
- design customer-facing controls for access, correction, deletion, and preference management.

In practice, privacy failures often emerge from data sprawl rather than from a single dramatic breach. Data is replicated into logs, caches, notebooks, feature stores, exports, and test environments. Once copied, it becomes much harder to track, govern, or delete. Data architecture therefore matters directly to privacy outcomes.

### 14.8.1 Data minimization in AI systems

AI teams are often tempted to keep every field because future modeling value is uncertain. This creates tension with privacy principles. A more disciplined approach is to distinguish:

- features required for the current decision use case;
- features that are merely easy to collect;
- sensitive attributes that require special scrutiny or exclusion;
- derived features whose predictive power may act as proxies for protected or private attributes.

The question is not only whether data is available, but whether it should be used. This is where privacy, ethics, and model governance intersect.

## 14.9 Compliance in the commerce environment

Compliance is not a single global checklist. It is a layered set of obligations shaped by geography, sector, customer base, payment methods, and product category. Some obligations are clearly legal or regulatory. Others are contractual, such as merchant, payment-provider, or marketplace requirements. Still others come from internal policy commitments made to customers or partners.

From an architectural perspective, compliance work usually translates into:

- classification of data types and processing activities;
- control mapping to systems, teams, and vendors;
- evidence generation through logs, approvals, and documented processes;
- periodic reviews of access, retention, change management, and incident handling;
- traceability between declared policy and operational reality.

A common mistake is to treat compliance as a documentation layer added after deployment. That approach produces weak evidence and brittle controls. If deletion, approval, consent handling, auditability, or override workflows are not built into the system design, policy language will not compensate for the gap.

## 14.10 Responsible AI in e-commerce

Responsible AI is the practice of building and operating AI systems in ways that are fair, accountable, transparent, safe, and governable in their business context. In e-commerce, this is not abstract philosophy. AI decisions influence which products customers see, what prices they face, whether transactions are flagged, how support cases are resolved, and what marketing messages are delivered.

Several properties matter:

- fairness: comparable users should not be systematically harmed without justification;
- explainability: teams should be able to explain the basis of important outcomes at the appropriate level;
- accountability: a business owner should be responsible for model behavior and overrides;
- auditability: inputs, outputs, versions, and decision paths should be reviewable after the fact;
- human oversight: some decisions should remain reviewable or reversible by trained staff;
- safety: systems should resist harmful content generation, unsafe actions, and inappropriate automation.

Responsible AI does not imply that every model must be fully interpretable in the same way. Rather, it means the level of explanation, testing, monitoring, and human control should match the risk of the use case.



### 14.10.1 High-risk and low-risk AI use cases

Not all AI applications deserve the same governance burden. Product-description generation and internal demand forecasting may be relatively low risk if the outputs are reviewed and their harms are limited. Fraud scoring, credit-like decisions, seller ranking, account suspension, dynamic pricing, and automated refund denial are more sensitive because they can materially affect users and business relationships.

A risk-based governance model helps organizations avoid two failures:

- under-governing important systems, which creates real harm;
- over-governing trivial systems, which creates bureaucracy without proportional benefit.

## 14.11 Bias, proxies, and unintended discrimination

Bias in commerce AI does not only arise from explicitly protected attributes. It often emerges indirectly through proxies, historical patterns, and feedback loops. Postal code, browsing device, purchase history, refund patterns, and engagement metrics may all encode social or economic differences. If these signals are used carelessly, the resulting model may reproduce or amplify unequal treatment.

Examples include:

- a fraud model disproportionately challenging transactions from certain neighborhoods or regions;
- a promotion optimization engine steering premium offers toward already advantaged segments;
- a ranking model reducing visibility for small sellers because historic click-through favors large brands;
- a support triage model assigning slower service to users whose writing style or language differs from the majority training set.

Mitigating such risks requires more than deleting sensitive columns. Teams need dataset review, outcome analysis by segment, challenge testing, and business discussion about acceptable tradeoffs. Some use cases may require policy constraints that override pure predictive optimization.

## 14.12 Explainability, contestability, and customer trust

Commerce organizations often focus on model performance but neglect user recourse. Yet customers and sellers care deeply about contestability: if a decision affects them, can they understand it at a useful level, and can they challenge or appeal it?

In practice, explainability has several audiences:

- engineers need diagnostic explanations to debug systems;
- business owners need performance explanations to manage tradeoffs;
- auditors need traceable evidence of control operation;
- customers need plain-language explanations and escalation paths when outcomes affect them materially.

An explanation does not always need to expose model internals. Often it is enough to explain the category of factors considered, the policy boundary involved, and the next steps available to the affected party. What matters is that the explanation is honest, actionable, and aligned with the actual system behavior.

### 14.13 LLM-specific risks and controls

Large language models create a distinct control challenge because they are probabilistic, context-dependent, and often connected to tools or internal knowledge bases. In commerce settings, they are typically used in support automation, seller assistance, catalog enrichment, policy summarization, and internal productivity workflows.

Important risk categories include:

- hallucinated answers about refunds, shipping, warranties, or policy;
- prompt injection through customer messages, seller content, or retrieved documents;
- leakage of personal or confidential data through prompts or logs;
- unsafe action execution when the model is connected to tools;
- over-automation that removes human review from sensitive decisions.

Controls for LLM systems should include:

- scoped retrieval and access filtering;
- prompt and tool guardrails;
- red-team testing with adversarial scenarios;
- response policies for sensitive domains such as refunds, legal questions, and account access;
- human escalation thresholds;
- monitoring of failure modes, not just average response quality.

The most important design choice is to decide what the model is allowed to do. Generating a draft is not the same as approving a refund. Summarizing a case is not the same as suspending a seller. Tool access should match the business risk.

## 14.14 Logging, monitoring, and auditability

Security, privacy, and responsible AI all depend on good records. If an enterprise cannot reconstruct what data was accessed, what model version produced an outcome, what prompt was sent, or which human approved an override, then investigation becomes guesswork.

Good auditability requires:

- structured logs with consistent identifiers across systems;
- versioning of models, prompts, rules, and datasets;
- records of approvals, overrides, and policy exceptions;
- separation of operational logs from broad user-accessible analytics;
- retention policies that preserve evidence without creating unnecessary data exposure.

Auditability is also essential for organizational learning. Incident reviews become much more valuable when teams can trace how a decision was made and which assumptions were embedded at the time.

## 14.15 Incident response and crisis handling

No system is perfectly secure or perfectly governed. The question is therefore not whether incidents will occur, but whether the organization can detect, contain, investigate, communicate, and recover effectively.

An e-commerce incident response plan should define:

- detection channels and severity criteria;
- ownership across engineering, security, legal, operations, and communications;
- containment steps for account abuse, data leakage, fraudulent activity, or harmful model behavior;
- evidence preservation requirements;
- decision rules for customer notification, regulator engagement, and vendor escalation;
- post-incident review and control improvement workflows.

AI incidents require particular care because the root cause may span model behavior, training data, prompts, product design, or organizational incentives. A model rollback alone may not address the underlying governance failure.

## 14.16 Operating model: who owns what

The recurring failure in governance programs is unclear ownership. If everyone is responsible, no one is responsible. A practical commerce operating model usually assigns:

- product owners to define acceptable use and business outcomes;
- engineering teams to implement security and reliability controls;
- data and ML teams to manage datasets, models, evaluation, and monitoring;
- security and privacy specialists to set standards, review risks, and support investigations;
- legal and compliance teams to interpret obligations and approve control approaches;
- frontline operations to handle escalations, appeals, and exception workflows.

This does not imply strict handoffs. Good governance is collaborative. But decision rights and approval paths need to be explicit, especially for launching or materially changing AI features.

## 14.17 A practical risk assessment template

Before launching a new AI-enabled e-commerce feature, teams should be able to answer a concise set of questions:

1. What decision or assistance function does the system perform?
2. Who is affected by the output: customer, seller, employee, or partner?
3. What data is used, and which fields are sensitive?
4. Could the system materially affect access, price, eligibility, safety, or reputation?
5. What failure modes are plausible, and how severe are they?
6. What human review, override, or appeal path exists?
7. What logs, metrics, and alerts will reveal misuse or drift?
8. What policy, compliance, or contractual obligations apply?
9. Which owner is accountable after go-live?

This kind of template creates discipline without forcing every project through the same heavyweight process. It also gives instructors a concrete bridge between technical architecture and governance practice.

## 14.18 Managerial implications

Executives often ask whether stronger controls slow innovation. In reality, the absence of controls slows scaling. Teams that do not know which data they can use, which actions models may take, which approvals are required, or how incidents will be handled tend to move either recklessly or not at all.

Well-designed governance increases execution quality in three ways. First, it reduces avoidable risk. Second, it clarifies decision rights and speeds launch readiness. Third, it improves trust with customers, partners, regulators, and internal stakeholders. In commerce, trust is not an abstract virtue. It affects conversion, retention, brand value, and partnership durability.

## 14.19 Hands-on lab: assessing an AI returns assistant

Consider an online retailer that wants to deploy an AI assistant for post-purchase support. The assistant can answer policy questions, summarize order status, and propose next-step actions for returns and exchanges. The business wants lower support cost and faster resolution times.

Students should evaluate the feature across four dimensions:

- security: can the assistant expose the wrong order or trigger unauthorized actions?
- privacy: what personal data is retrieved, logged, or retained?
- compliance: are policy promises, consumer rights, and escalation requirements reflected in the workflow?
- responsible AI: could the assistant generate misleading answers, treat some users unfairly, or deny effective recourse?

A good student output would include:

- a simple architecture diagram;
- a list of assets and trust boundaries;
- a risk register with severity and mitigations;
- proposed controls for authentication, retrieval filtering, tool permissions, logging, and human escalation;
- launch criteria and post-launch monitoring metrics.

## 14.20 Multiple-choice questions (MCQs)

1. Which statement best captures privacy-by-design in e-commerce?
  - (a) collect all potentially useful data and decide later what to retain;

- (b) minimize data collection and retention based on justified business purpose;
  - (c) anonymize all customer data immediately after checkout;
  - (d) prohibit analytics on operational systems entirely.
2. Which of the following is the clearest example of a trust boundary in a commerce architecture?
- (a) a local function calling another function in the same process;
  - (b) a service crossing into a third-party payment gateway;
  - (c) a variable being reused inside a notebook;
  - (d) a chart rendered in an internal dashboard.
3. Why is least privilege important for internal commerce tools?
- (a) it guarantees zero security incidents;
  - (b) it ensures every employee can inspect all customer records when needed;
  - (c) it limits damage from misuse or compromised accounts by restricting access to necessary functions only;
  - (d) it removes the need for audit logs.
4. Which risk is most specific to LLM-enabled systems?
- (a) warehouse stockouts;
  - (b) prompt injection via retrieved or user-supplied content;
  - (c) currency conversion latency;
  - (d) barcode duplication.
5. Which governance approach is generally most practical for AI in commerce?
- (a) identical controls for all use cases regardless of risk;
  - (b) no controls until a major incident occurs;
  - (c) risk-based governance with stronger controls for higher-impact uses;
  - (d) full manual review of every model output in all systems.
6. What is the main reason auditability matters?
- (a) it makes models mathematically perfect;
  - (b) it allows teams to reconstruct decisions, investigate incidents, and demonstrate control operation;
  - (c) it eliminates the need for security monitoring;
  - (d) it replaces compliance documentation.
7. Which example best illustrates proxy discrimination risk?
- (a) a model uses a random seed during training;
  - (b) a ranking system uses historical clicks that systematically favor dominant sellers;
  - (c) a dashboard refreshes every hour instead of every minute;
  - (d) a support script includes a greeting.

8. What is the best interpretation of contestability?
- (a) the ability of customers or sellers to understand and challenge significant outcomes;
  - (b) the requirement that every model be open-source;
  - (c) the practice of hiding decision logic to reduce gaming;
  - (d) the replacement of support teams with automation.

**Answer key**

- 1. b
- 2. b
- 3. c
- 4. b
- 5. c
- 6. b
- 7. b
- 8. a





## Chapter 15

# Capstone: Enterprise-Grade AI E-Commerce Architecture (Week 14)

### Capstone scenario (agreed)

- Scenario: end-to-end “AI commerce + Odoo” reference architecture for a single company.
- Goal: demonstrate EA alignment, ERP integration, and measurable AI value with operational readiness.
- Business model: hybrid (B2C + B2B).
- Commerce front-end: both (compare Odoo Website/eCommerce vs a separate headless storefront integrating to Odoo).
- AI stack: both (ML for recs/fraud/pricing + LLMs for support/ops).

### 15.1 Scope (placeholder)

- End-to-end architecture: channels, services, data platform, ML/LLM layer, ERP integration.
- Evaluation: business metrics, technical metrics, and operational readiness.

### 15.2 Capstone deliverables (placeholder)

- Architecture pack (diagrams), backlog, data contracts, and evaluation plan.



## Appendix A

# Datasets and Synthetic Data for Teaching (placeholder)



## Appendix B

### Template: System Design Document (placeholder)



## Appendix C

### Template: Data Contract (placeholder)





## Appendix D

### Template: Model Card and Risk Assessment (placeholder)