

AI-Driven E-Commerce (Master's Level)
Book Structure & 14-Week Plan (Placeholder)

Faculty of Computers and Information Sciences
Mansoura University, Egypt

March 23, 2026

Contents

1	Weekly Plan Overview	21
2	An Introduction to E-Commerce with a Focus on Machine Learning, Deep Learning, and Artificial Intelligence	23
2.1	Introduction to E-Commerce	23
2.1.1	E-Commerce Business Models and Value Chain	24
2.2	Foundations of AI, Machine Learning, and Deep Learning	24
2.2.1	Definitions and Relationships	24
2.2.2	Types of Learning and E-Commerce Examples	25
2.3	Machine Learning Across the E-Commerce Lifecycle	25
2.3.1	Personalized Recommendations	25
2.3.2	Search and Ranking	25
2.3.3	Segmentation, churn, and CLV	25
2.3.4	Dynamic Pricing and Promotion Optimization	25
2.3.5	Demand Forecasting and Inventory Management	26
2.3.6	Fraud Detection and Transaction Security	26
2.3.7	Customer Service and Conversational Commerce	26
2.4	Deep Learning in E-Commerce	26
2.4.1	Neural Recommendation Systems	26
2.4.2	NLP for Product Text and Reviews	26
2.4.3	Computer Vision and Visual Search	26
2.5	Data, Architecture, and Deployment Considerations	26
2.6	Challenges and Responsible AI	27

2.7	Multiple-choice questions (MCQs)	27
3	Platforms, Marketplaces, and Business Models (Week 2)	29
3.1	Learning objectives	29
3.2	From storefront to platform	29
3.3	Value creation, value capture, and growth	29
3.3.1	Value creation and matching	29
3.3.2	Value capture models	30
3.3.3	Network effects	30
3.4	Governance and trust as system requirements	30
3.5	Business model patterns (and architectural implications)	30
3.5.1	B2C retail	30
3.5.2	B2B commerce	31
3.5.3	Marketplace	31
3.5.4	Digital services (subscriptions)	31
3.6	Where AI fits (system-level view)	31
3.7	Hands-on (placeholder)	31
3.8	Case studies (placeholder)	31
3.9	Multiple-choice questions (MCQs)	32
4	E-Commerce Data Foundations: Events, Experimentation, and Metrics (Week 3)	33
4.1	Learning objectives	33
4.2	Why data foundations matter	33
4.3	From business process to event model	34
4.3.1	Events versus state	34
4.3.2	Event taxonomies	35
4.4	Event schema design	35
4.4.1	Required design principles	35
4.4.2	Example event record	36
4.5	Identity resolution and entity modeling	36

4.5.1	Identifiers are not interchangeable	37
4.5.2	Identity stitching	37
4.5.3	Privacy-aware identity design	37
4.6	Data quality as a product requirement	38
4.6.1	Data contracts	38
4.6.2	Validation and observability	38
4.7	Attribution and its limits	39
4.7.1	Common attribution models	39
4.7.2	Selection bias and measurement bias	39
4.8	Metrics layers and decision quality	40
4.8.1	North Star and supporting metrics	40
4.8.2	Metric dimensions	40
4.8.3	Retention and cohort logic	41
4.9	Experimentation in e-commerce	41
4.9.1	A/B testing fundamentals	41
4.9.2	Common experimental pitfalls	41
4.9.3	Bandits and adaptive experimentation	42
4.10	From analytics to machine learning features	42
4.11	Hands-on lab: build a minimal metrics layer	42
4.11.1	Lab scenario	43
4.11.2	Lab tasks	43
4.11.3	Expected learning outcome	43
4.12	Managerial and architectural implications	43
4.13	Multiple-choice questions (MCQs)	44
4.14	Exercises (short)	45
4.15	Mini-case (Odoo-linked, vendor-neutral)	45
5	Enterprise Architecture for E-Commerce (Week 4)	47
5.1	Learning objectives	47
5.2	Why enterprise architecture matters in e-commerce	47

5.3	Enterprise architecture as a set of views	48
5.4	Capability maps	48
5.4.1	Core e-commerce capabilities	49
5.4.2	Why capability maps are useful	49
5.5	Value streams	49
5.5.1	Typical e-commerce value streams	50
5.5.2	Why value streams matter architecturally	50
5.6	Domain boundaries and bounded contexts	50
5.6.1	Typical commerce domains	50
5.6.2	Good boundaries versus bad boundaries	51
5.7	Systems of record and systems of engagement	51
5.7.1	Typical patterns	51
5.7.2	Systems of engagement	52
5.8	Reference architecture for AI-driven e-commerce	52
5.8.1	Where AI fits	52
5.9	Current state, target state, and transition architecture	53
5.9.1	Current-state analysis	53
5.9.2	Target architecture	53
5.9.3	Transition architecture	53
5.10	Architectural concerns specific to AI	54
5.11	Operating model and team topology	54
5.11.1	Typical team patterns	54
5.12	Artifacts for students and architects	54
5.13	Hands-on artifacts	55
5.14	Managerial implications	55
5.15	Multiple-choice questions (MCQs)	56
5.16	Exercises (short)	57
5.17	Mini-case (Odoo-linked, vendor-neutral)	57

6.1	Learning objectives	59
6.2	Why integration patterns matter	59
6.3	Integration styles	60
6.3.1	Synchronous request-response	60
6.3.2	Asynchronous events	60
6.3.3	Batch and file-based integration	60
6.4	Choosing the right integration style	61
6.4.1	Questions to ask	61
6.4.2	Typical choices in commerce	61
6.5	API-first integration	62
6.5.1	Good API design principles	62
6.5.2	Commands versus queries	62
6.6	Idempotency	62
6.6.1	Why idempotency is required	62
6.6.2	Idempotency keys	62
6.7	Events and event-driven architecture	63
6.7.1	Domain events versus technical events	63
6.7.2	Benefits of event-driven integration	63
6.7.3	Costs of event-driven integration	63
6.8	Reliability patterns in distributed workflows	64
6.8.1	Retries and backoff	64
6.8.2	Dead-letter handling	64
6.8.3	Timeouts and circuit breakers	64
6.8.4	Reconciliation jobs	64
6.9	The outbox pattern	64
6.9.1	How the outbox pattern works	65
6.9.2	Why it matters in commerce	65
6.10	Message contracts and schema governance	65
6.10.1	What a contract should define	65
6.10.2	Backward compatibility	66

6.11	Observability for integrations	66
6.11.1	Minimum signals	66
6.11.2	Correlation and traceability	66
6.12	iPaaS, ESB, brokers, and workflow orchestration	67
6.12.1	iPaaS and ESB concepts	67
6.12.2	Event brokers and streams	67
6.12.3	Workflow orchestration	67
6.13	Integration anti-patterns	67
6.14	Hands-on lab: modeling an order lifecycle	68
6.14.1	Lab scenario	68
6.14.2	Lab tasks	68
6.14.3	Expected learning outcome	68
6.15	Managerial and architectural implications	68
6.16	Multiple-choice questions (MCQs)	69
6.17	Exercises (short)	70
6.18	Mini-case (Odoo-linked, vendor-neutral)	70
7	ERP/CRM/SCM Integration in Commerce (Week 6)	71
7.1	Learning objectives	71
7.2	Why ERP/CRM/SCM integration is central to commerce	71
7.3	The enterprise systems perspective	72
7.3.1	ERP	72
7.3.2	CRM	72
7.3.3	SCM and warehouse systems	72
7.4	Commerce versus enterprise ownership	72
7.4.1	Common ownership patterns	73
7.4.2	Why ownership must be explicit	73
7.5	Systems of record for key entities	73
7.5.1	Customer	73
7.5.2	Product	73

7.5.3	Price list	74
7.5.4	Order	74
7.5.5	Invoice and payment reconciliation	74
7.5.6	Shipment	74
7.6	Master data management (MDM)	74
7.6.1	Why MDM matters	74
7.6.2	Practical MDM principles	75
7.7	Order-to-cash in integrated commerce	75
7.7.1	Typical flow	75
7.7.2	Architectural decisions inside the flow	76
7.7.3	Failure modes	76
7.8	Inventory, fulfillment, and availability	76
7.8.1	Inventory concepts that should not be conflated	76
7.8.2	Synchronization strategies	77
7.9	Returns and refunds	77
7.9.1	Typical return flow	77
7.9.2	Architectural risks	77
7.10	B2B pricing and account structures	78
7.10.1	Price lists and negotiated terms	78
7.10.2	Account hierarchies	78
7.11	Using Odoo as a reference ERP	79
7.11.1	Illustrative Odoo-aligned mapping	79
7.11.2	What should remain vendor-neutral	79
7.12	Integration risks and mitigation strategies	79
7.13	Artifacts for this chapter	80
7.14	Managerial implications	80
7.15	Multiple-choice questions (MCQs)	80
7.16	Exercises (short)	81
7.17	Mini-case (Odoo-linked, vendor-neutral)	81

8	Recommenders I: Classical Methods + Learning-to-Rank (Week 7)	83
8.1	Scope (placeholder)	83
8.2	Hands-on lab (placeholder)	83
9	Recommenders II: Deep Learning, Embeddings, and Retrieval (Week 8)	85
9.1	Scope (placeholder)	85
9.2	Hands-on lab (placeholder)	85
10	Pricing and Promotions with ML + Optimization (Week 9)	87
10.1	Scope (placeholder)	87
10.2	Hands-on lab (placeholder)	87
11	Fraud, Risk, and Trust & Safety (Week 10)	89
11.1	Scope (placeholder)	89
11.2	Hands-on lab (placeholder)	89
12	Customer Service Automation with LLMs (Week 11)	91
12.1	Scope (placeholder)	91
12.2	Hands-on lab (placeholder)	91
13	MLOps + DataOps for E-Commerce (Week 12)	93
13.1	Scope (placeholder)	93
13.2	Artifacts (placeholder)	93
14	Security, Privacy, Compliance, and Responsible AI (Week 13)	95
14.1	Scope (placeholder)	95
14.2	Discussion (placeholder)	95
15	Capstone: Enterprise-Grade AI E-Commerce Architecture (Week 14)	97
15.1	Scope (placeholder)	97
15.2	Capstone deliverables (placeholder)	97
A	Datasets and Synthetic Data for Teaching (placeholder)	99

<i>CONTENTS</i>	11
B Template: System Design Document (placeholder)	101
C Template: Data Contract (placeholder)	103
D Template: Model Card and Risk Assessment (placeholder)	105

How to Use This Draft

This document is a *structure-only* placeholder for an AI-driven, practical Master's-level e-commerce book. Each chapter will be written later; for now, sections contain brief bullets stating scope, key skills, and deliverables.

Intended Audience and Prerequisites

- Audience: Master's students in Information Systems / Computer & Information Sciences.
- Prereqs (recommended): databases, basic programming, web fundamentals, and introductory ML.
- Tooling (suggested): Python, notebooks, APIs, cloud services, and an ERP/EA modeling tool (as available).

Hands-on Track (Agreed)

- Approach: both tracks (lighter each).
- Track A (Enterprise/ERP): Odoo-focused integration labs (Sales/Inventory/Accounting flows).
- Track B (AI/ML): ML-heavy notebooks using e-commerce datasets (recsys, pricing, fraud, support).

Course Outcomes (Draft)

- Design end-to-end e-commerce systems with enterprise architecture (EA) and integration concerns.
- Apply ML/DL/AI to personalization, search, pricing, fraud, operations, and customer service.
- Integrate e-commerce with ERP/CRM/SCM (reference: Odoo) via APIs, events, and modern integration patterns.
- Evaluate systems for security, privacy, governance, ethics, and measurable business value.

Chapter 1

Weekly Plan Overview

Course emphasis (agreed)

- Overall emphasis: balanced (EA/ERP + ML/DL/AI).
- ERP reference stack: Odoo (concepts remain vendor-neutral where possible).

Week	Topic (maps to the corresponding chapter)
1	Foundations of modern e-commerce systems and AI-driven product thinking
2	Digital platforms, marketplaces, and business models (B2C/B2B/D2C, multi-sided platforms)
3	E-commerce data foundations: events, tracking, data quality, and experimentation
4	Enterprise architecture for e-commerce: capabilities, domains, and reference architectures
5	Enterprise integration: APIs, iPaaS, ESB, events/streaming, and integration patterns
6	ERP/CRM/SCM integration: order-to-cash, procure-to-pay, inventory, and master data
7	Recommenders I: classical + learning-to-rank for personalization and search
8	Recommenders II: deep learning, embeddings, sequence models, and retrieval
9	Pricing and promotion analytics: forecasting, elasticity, and optimization (ML + OR)
10	Fraud, risk, and trust & safety: anomaly detection, graph ML, and AML patterns
11	Customer service automation: LLMs, RAG, chatbots, and agentic workflows
12	MLOps + DataOps for e-commerce: deployment, monitoring, drift, and governance
13	Security, privacy, compliance, and responsible AI in commerce
14	Capstone: enterprise-grade AI e-commerce solution architecture + evaluation

Chapter 2

An Introduction to E-Commerce with a Focus on Machine Learning, Deep Learning, and Artificial Intelligence

Abstract

Electronic commerce (e-commerce) has evolved from static online catalogs into highly dynamic, data-driven ecosystems. This transformation has been powered largely by advances in artificial intelligence (AI), particularly machine learning (ML) and deep learning (DL). This chapter introduces the foundations of e-commerce and explains how AI techniques are embedded across the e-commerce value chain—from customer acquisition and product discovery to pricing, logistics, and fraud prevention. The discussion is aimed at master’s students in computer and information sciences, and therefore emphasizes conceptual clarity, system-level thinking, and the mapping between theoretical models and real-world e-commerce applications.

Keywords: e-commerce, artificial intelligence, machine learning, deep learning, recommendation systems, dynamic pricing, fraud detection, personalization, logistics optimization

2.1 Introduction to E-Commerce

E-commerce refers to the buying and selling of goods and services via electronic networks, primarily the internet [?]. It encompasses a broad set of transactional models: business-to-consumer (B2C) online retail, business-to-business (B2B) procurement platforms, consumer-to-consumer (C2C) marketplaces, and consumer-to-business (C2B) models such as influencer platforms and freelance marketplaces [?].

From a computing perspective, an e-commerce platform is not just a web front-end. It is a socio-technical system composed of:

- User interfaces (web, mobile, conversational)

- Application logic (catalog, cart, checkout, account management)
- Payment and risk engines
- Logistics and fulfillment systems
- Data infrastructure and analytics pipelines
- AI/ML components that drive personalization, prediction, and automation

A key characteristic of e-commerce is its data richness. Every user interaction—page views, searches, clicks, scrolls, wishlists, add-to-carts, purchases, returns, and even time-to-decision—can be captured as fine-grained behavioral logs. Combined with product metadata, prices, promotions, and external contextual data, this creates an ideal environment for AI and ML [? ?].

2.1.1 E-Commerce Business Models and Value Chain

Common e-commerce models include:

- **B2C retail:** Online stores selling directly to consumers.
- **Marketplaces:** Platforms mediating between many sellers and buyers.
- **B2B platforms:** Portals for wholesale ordering, procurement, and supply chain integration.
- **Digital services:** Software subscriptions, digital content, and online education.

Across these models, the **e-commerce value chain** typically includes:

1. Customer acquisition and traffic generation (marketing, ads, SEO)
2. Product discovery and decision support (search, recommendations, reviews)
3. Conversion and transaction processing (checkout, payments, risk checks)
4. Order fulfillment and logistics (inventory, warehousing, routing)
5. Post-purchase engagement (support, returns, loyalty, re-activation)

AI, ML, and DL are now present at each stage of this chain [?].

2.2 Foundations of AI, Machine Learning, and Deep Learning

2.2.1 Definitions and Relationships

- **Artificial Intelligence (AI):** Systems that exhibit behaviors considered “intelligent” (perception, reasoning, learning, decision-making).

- **Machine Learning (ML):** Algorithms that learn patterns from data and improve with experience.
- **Deep Learning (DL):** ML methods based on multi-layer neural networks that learn representations from large-scale data.

Thus, deep learning $\subset$ machine learning $\subset$ artificial intelligence [? ? ?].

2.2.2 Types of Learning and E-Commerce Examples

1. **Supervised learning:** learn $f : X \rightarrow Y$ from labeled examples (x_i, y_i) . Examples: conversion prediction, fraud classification, demand forecasting.
2. **Unsupervised learning:** discover structure in unlabeled data. Examples: customer segmentation, anomaly detection, and learning embeddings from interactions.
3. **Reinforcement learning:** learn a policy to maximize long-term reward. Examples: recommendation policies and dynamic pricing under constraints.
4. **Representation learning:** learn features $\phi(x)$ automatically (often via DL). Examples: user/item embeddings; Transformer-based session representations.

2.3 Machine Learning Across the E-Commerce Lifecycle

2.3.1 Personalized Recommendations

Recommendation systems use past behavior and product attributes to suggest items a user is likely to buy [? ?]. Key evaluation ideas include ranking quality (precision@ k , recall@ k , NDCG) and business impact (incremental revenue) [?].

2.3.2 Search and Ranking

Learning-to-rank uses user feedback (including clicks) to optimize product ranking in response to queries [?]. Modern systems increasingly incorporate semantic retrieval and NLP components [?].

2.3.3 Segmentation, churn, and CLV

Segmentation (e.g., via clustering) and predictive modeling (e.g., churn/CLV) connect behavioral data to marketing and retention actions [?].

2.3.4 Dynamic Pricing and Promotion Optimization

Pricing combines prediction (demand response) with optimization under constraints; evaluation typically requires controlled online experimentation [? ?].

2.3.5 Demand Forecasting and Inventory Management

Forecasting methods range from statistical time-series models to ML sequence models; forecasts drive replenishment and inventory decisions [? ?].

2.3.6 Fraud Detection and Transaction Security

Fraud detection is an imbalanced classification problem with non-stationarity (attackers adapt), and must balance false positives against fraud loss [? ?].

2.3.7 Customer Service and Conversational Commerce

Customer support automation mixes information retrieval, intent/entity extraction, and increasingly LLM-based interaction [?].

2.4 Deep Learning in E-Commerce

DL is especially useful for unstructured and high-dimensional data (text, images, sequences) [?].

2.4.1 Neural Recommendation Systems

Deep recommender architectures model non-linear user–item interactions and session dynamics [? ?].

2.4.2 NLP for Product Text and Reviews

NLP supports query understanding, semantic search, review analysis, and conversational interfaces [?].

2.4.3 Computer Vision and Visual Search

Vision models support category classification, attribute extraction, and similarity search via embeddings [?].

2.5 Data, Architecture, and Deployment Considerations

E-commerce ML requires robust pipelines (data quality, governance), reliable deployment (latency, availability), and continuous evaluation (monitoring and experiments) [?].

2.6 Challenges and Responsible AI

Key challenges include bias, privacy/security, explainability for high-impact decisions, and the organizational processes needed for safe deployment [? ?].

2.7 Multiple-choice questions (MCQs)

1. Which relationship is correct?
 - (a) $AI \subset ML \subset DL$
 - (b) $DL \subset ML \subset AI$
 - (c) $ML \subset DL \subset AI$
 - (d) AI, ML, and DL are disjoint fields
2. In e-commerce, which is **most** naturally framed as a *ranking* problem?
 - (a) Choosing the best warehouse location for a new facility
 - (b) Ordering products in a search results page for a query
 - (c) Estimating the total demand for next quarter
 - (d) Balancing accounting entries for an invoice
3. Which metric is **most** aligned with offline evaluation of a top- k recommender?
 - (a) $NDCG@k$
 - (b) Mean squared error (MSE) on prices
 - (c) Page load time (ms)
 - (d) Uptime percentage
4. Fraud detection is often difficult because:
 - (a) Fraud labels are always perfectly accurate
 - (b) The problem is typically extremely imbalanced and attackers adapt over time
 - (c) Fraud has no economic cost, only technical cost
 - (d) The best approach is always rule-based
5. Which statement best describes why A/B testing is important in e-commerce ML?
 - (a) It guarantees the model is unbiased
 - (b) It measures causal impact on business KPIs under real user behavior
 - (c) It replaces the need for offline evaluation entirely
 - (d) It eliminates the need for monitoring after deployment

Answer key

1. (b)
2. (b)
3. (a)
4. (b)
5. (b)

Chapter 3

Platforms, Marketplaces, and Business Models (Week 2)

3.1 Learning objectives

- Distinguish e-commerce *channels* (storefronts) from *platforms* (ecosystems) and *marketplaces* (multi-seller).
- Analyze platform economics: network effects, multi-homing, and pricing of participation.
- Translate business model choices into architectural requirements (identity, trust, data, integration, ERP).

3.2 From storefront to platform

- **Storefront:** a single merchant selling to buyers (B2C/B2B) via a web/app channel.
- **Platform:** a system that enables interactions among multiple participant groups (e.g., buyers, sellers, advertisers, logistics providers).
- **Marketplace:** a platform where multiple sellers list items and transactions are mediated by platform policies and infrastructure.

3.3 Value creation, value capture, and growth

3.3.1 Value creation and matching

Platforms create value by reducing search and transaction costs, increasing variety, and improving trust; AI frequently strengthens these effects by improving matching and decision support.

3.3.2 Value capture models

- Take rate (commission on completed transactions)
- Subscription (seller or buyer membership tiers)
- Listing fees and value-added services (e.g., promotions, analytics)
- Advertising (sponsored listings, retail media)
- Fulfillment fees (warehousing, last-mile delivery, returns handling)

3.3.3 Network effects

- **Direct network effects:** more users attract more users (e.g., social commerce communities).
- **Indirect network effects:** more buyers attract more sellers and vice versa (classic marketplace dynamic).
- **Cold start:** bootstrapping supply and demand when network effects are weak.
- **Multi-homing:** sellers/buyers participate in multiple platforms; affects differentiation strategy.

3.4 Governance and trust as system requirements

Marketplace governance is not “policy only”: it becomes a set of product and system requirements.

- Seller onboarding, verification, and compliance workflows (KYC/KYB where relevant).
- Catalog integrity: listing policies, moderation, and counterfeit detection.
- Reviews/ratings, dispute resolution, refunds/returns policy enforcement.
- Trust & safety operations: rule+ML detection, case management, and human-in-the-loop escalation.

3.5 Business model patterns (and architectural implications)

3.5.1 B2C retail

- Key systems: PIM/catalog, promotions, payments, fulfillment, customer support.
- ERP/CRM integration: order-to-cash, inventory, accounting, returns.

3.5.2 B2B commerce

- Quotes, negotiated pricing, approval workflows, invoicing, credit limits.
- Integration-heavy: procurement, supplier management, and EDI/API flows.

3.5.3 Marketplace

- Split catalog ownership (platform vs sellers), seller tools, payout/settlement.
- Additional “control plane”: seller policies, listing moderation, abuse prevention.

3.5.4 Digital services (subscriptions)

- Recurring billing, entitlements, and churn/retention analytics.
- Usage-based pricing requires event design and metering.

3.6 Where AI fits (system-level view)

- **Matching layer:** search, recommendations, ranking.
- **Trust layer:** fraud/abuse detection, content moderation.
- **Operations layer:** forecasting, replenishment, routing.
- **Experience layer:** personalization and conversational commerce.

3.7 Hands-on (placeholder)

- Create a one-page *platform blueprint*: participant groups, value exchange, pricing model, and required capabilities.
- Map capabilities to an EA view and identify which capabilities are ERP-owned vs commerce-owned.

3.8 Case studies (placeholder)

- Modern marketplace patterns (generic; no vendor lock-in).

3.9 Multiple-choice questions (MCQs)

1. Which statement best distinguishes a *storefront* from a *marketplace*?
 - (a) A storefront is always B2B; a marketplace is always B2C.
 - (b) A storefront has a single merchant of record; a marketplace mediates transactions among multiple sellers and buyers.
 - (c) A storefront cannot use AI; a marketplace must use AI.
 - (d) A storefront requires ERP integration; a marketplace does not.
2. *Indirect network effects* in marketplaces most directly mean:
 - (a) The platform has a single user group and growth is linear.
 - (b) The value to buyers increases as more sellers join, and the value to sellers increases as more buyers join.
 - (c) Users prefer multiple platforms (multi-homing) regardless of differentiation.
 - (d) The platform must subsidize logistics to succeed.
3. In B2B commerce, which requirement is **most typical** compared to B2C?
 - (a) Session-based recommendations
 - (b) Negotiated pricing and approval workflows
 - (c) Visual search
 - (d) Social reviews and influencer marketing
4. *Multi-homing* refers to:
 - (a) Hosting the platform in multiple cloud regions.
 - (b) Sellers or buyers participating in multiple platforms simultaneously.
 - (c) Using multiple payment gateways for redundancy.
 - (d) Maintaining multiple warehouses for the same SKU.
5. Which is the **best** example of a marketplace *governance* capability?
 - (a) Re-ranking items using embeddings
 - (b) Seller onboarding verification and dispute resolution policies
 - (c) Demand forecasting for replenishment
 - (d) Offline model training with hyperparameter tuning

Answer key

1. (b)
2. (b)
3. (b)
4. (b)
5. (b)

Chapter 4

E-Commerce Data Foundations: Events, Experimentation, and Metrics (Week 3)

4.1 Learning objectives

- Design an event-driven data model for storefront, marketplace, and mobile commerce scenarios.
- Distinguish identifiers, identities, sessions, and entities across operational and analytical systems.
- Explain how attribution, data quality, and metric definitions affect decision-making and ML performance.
- Plan controlled experiments and understand when adaptive approaches such as bandits are appropriate.

4.2 Why data foundations matter

AI systems in e-commerce are only as reliable as the data substrate beneath them. Recommendation models, pricing systems, fraud detectors, and executive dashboards all depend on the same underlying facts: who did what, when, in which context, and with what downstream outcome. If those facts are delayed, duplicated, ambiguously defined, or disconnected across systems, both analytics and machine learning become fragile.

This chapter treats data foundations as a *system design* problem rather than a reporting problem. The goal is not merely to “collect more data,” but to create a coherent measurement layer that can support:

- Product analytics and funnel analysis
- Operational reporting across commerce and ERP systems

- Reliable feature generation for ML models
- Controlled online experimentation
- Governance, auditing, and cross-functional alignment

4.3 From business process to event model

E-commerce platforms contain several interacting process layers:

- **Customer interaction layer:** impressions, searches, clicks, scrolls, add-to-cart, checkout steps, support interactions.
- **Commercial transaction layer:** order placement, payment authorization, invoicing, shipment, cancellation, return, refund.
- **Operational layer:** inventory adjustments, picking, packing, supplier replenishment, carrier status changes.
- **Experimentation layer:** exposure to treatments, eligibility checks, metric assignment, and holdouts.

Each layer emits events, but those events serve different purposes. A page view event captures human behavior. An order-created event captures a business transaction. A shipment-delivered event captures operational state. An experiment-exposed event captures causal assignment context. Strong data design begins by separating these concerns while still linking them with stable keys.

4.3.1 Events versus state

Students often confuse *events* with *tables of current state*. The distinction is fundamental:

- **State** answers “what is true now?” Examples: current inventory, current product price, current loyalty tier.
- **Events** answer “what happened?” Examples: product viewed, cart updated, payment failed, order shipped.

State is necessary for operations, but events are essential for causality, sequence modeling, root-cause analysis, and reconstructing past behavior. For example, a fraud model may need the sequence of failed payment attempts before an order was placed, not just the final order record. Similarly, retention analysis depends on when a user returned, not merely whether the user is currently “active.”

4.3.2 Event taxonomies

A useful event taxonomy groups events into classes that are stable even when the user interface changes. For a commerce platform, a practical taxonomy may include:

- **Discovery events:** homepage viewed, category viewed, search submitted, filter applied, result clicked.
- **Consideration events:** product viewed, review expanded, comparison opened, wishlist updated.
- **Intent events:** add-to-cart, remove-from-cart, coupon applied, shipping option selected.
- **Transaction events:** checkout started, payment submitted, order completed, order canceled.
- **Post-purchase events:** shipment delivered, return requested, refund issued, support ticket opened.

This style of taxonomy makes analytics more robust than mirroring UI element names directly. If a button label changes from “Buy Now” to “Checkout,” the business meaning should still map to a stable event family rather than creating an entirely new measurement definition.

4.4 Event schema design

An event is only analytically useful if its schema is consistent. At minimum, most commerce events should include:

- **Event name:** stable business-oriented label such as `product_viewed` or `checkout_started`
- **Event time:** when the event actually occurred
- **Ingestion time:** when the event reached the data platform
- **Actor identifiers:** user ID, guest ID, device ID, session ID, seller ID where relevant
- **Entity identifiers:** product ID, SKU ID, cart ID, order ID, payment ID
- **Context fields:** channel, app version, locale, campaign, experiment assignment, page type
- **Properties:** quantity, price, discount, position in ranking, search query, payment method

4.4.1 Required design principles

- **Business semantics before UI semantics:** schemas should reflect domain meaning, not implementation accidents.
- **Stable field names:** renaming fields casually breaks dashboards, pipelines, and models.
- **Explicit units:** store price currency, quantity units, and time zones unambiguously.

- **Immutable raw events:** avoid rewriting the historical log except for tightly governed correction procedures.
- **Versioned schemas:** consumers must know when a producer changed meaning or structure.

4.4.2 Example event record

The following conceptual record shows how a product view may be represented:

```

event_name      = product_viewed
event_time      = 2026-03-23T14:03:11Z
ingestion_time  = 2026-03-23T14:03:13Z
user_id         = null
guest_id        = g_1842
session_id      = s_88fa
product_id      = p_501
sku_id          = sku_501_bl_m
page_type       = product_detail
channel         = mobile_app
campaign_id     = spring_push_7
rank_position   = 4
experiment_id   = rec_widget_v3
variant         = treatment

```

The example illustrates an important practice: authenticated and anonymous identifiers may coexist. This supports both guest behavior analysis and later identity stitching when a guest eventually logs in or purchases.

4.5 Identity resolution and entity modeling

One of the hardest problems in e-commerce analytics is deciding what a “customer” means across channels and systems. The same person may appear as:

- an anonymous web browser cookie,
- a mobile-app installation,
- a logged-in storefront account,
- a CRM contact,
- an ERP customer or partner record,
- a marketplace buyer account,
- or multiple records created through operational errors.

4.5.1 Identifiers are not interchangeable

Different identifiers serve different roles:

- **User/account ID:** application-level identity for authentication and personalization.
- **Customer/partner ID:** commercial identity used in ERP, billing, and support.
- **Session ID:** short-lived grouping of interactions.
- **Device or browser ID:** technical identifier for anonymous continuity.
- **Order ID:** transaction-level identity.
- **Household or organization ID:** aggregation identity for B2B or family accounts.

If analysts casually join these identifiers as though they were equivalent, metrics become unstable. For example, “active customers” may mean active accounts in one report but active billing entities in another. That inconsistency can distort conversion rates, retention cohorts, and lifetime value estimates.

4.5.2 Identity stitching

Identity stitching is the process of linking records that likely belong to the same real-world actor. Common rules include:

- linking a guest cart to a logged-in account after authentication,
- attaching a mobile device history to a user after sign-in,
- associating storefront orders with an ERP customer once order export completes,
- maintaining a separate unresolved pool when confidence is low.

The core design principle is to preserve raw evidence and build linkage logic transparently. Many downstream problems arise when teams overwrite raw IDs with guessed “master” identities and can no longer audit how the match occurred.

4.5.3 Privacy-aware identity design

Identity resolution must also respect privacy and compliance constraints. A technically convenient identifier is not automatically a legally or ethically appropriate one. Good practice includes:

- minimizing collection of directly identifying information,
- separating analytical IDs from operational personal data where possible,
- documenting retention periods,
- and ensuring that consent and deletion workflows can propagate to analytical systems.

4.6 Data quality as a product requirement

Data quality problems in commerce are rarely abstract. They directly alter revenue reporting, experiment reads, and ML labels. Representative failure modes include:

- **Missing events:** checkout steps are undercounted, making conversion funnels look healthier than reality.
- **Duplicate events:** add-to-cart counts inflate intent metrics and derived model features.
- **Late events:** daily dashboards and training jobs disagree because the data arrive after aggregation windows close.
- **Semantic drift:** an event name remains unchanged while the product team silently changes its meaning.
- **Broken joins:** product IDs in clickstream logs no longer match product IDs in the catalog or ERP.

4.6.1 Data contracts

An effective way to reduce these failures is to define *data contracts* between producers and consumers. A data contract specifies:

- event name and schema,
- field definitions and allowed values,
- freshness expectations,
- ownership and escalation path,
- versioning policy,
- and validation rules.

This shifts governance left. Instead of discovering breakage after dashboards fail, producers must acknowledge downstream dependencies before changing event payloads.

4.6.2 Validation and observability

Data platforms should monitor more than pipeline uptime. Useful checks include:

- row-count anomalies by event family,
- null-rate spikes in required fields,
- duplicate-rate thresholds by event key,

- cross-system reconciliation between orders, payments, and shipments,
- and schema drift alerts when producers add or remove fields.

In mature teams, these checks are treated as production monitoring. Broken analytics should be considered an operational incident, not a minor inconvenience.

4.7 Attribution and its limits

Attribution asks which touchpoints influenced an outcome such as purchase, repeat visit, or subscription renewal. In e-commerce, this sounds straightforward but is often confounded by user behavior, cross-device journeys, channel interactions, and pre-existing intent.

4.7.1 Common attribution models

- **Last-touch attribution:** credits the final recorded touchpoint before conversion.
- **First-touch attribution:** credits the earliest touchpoint in the conversion path.
- **Multi-touch heuristics:** divide credit across several interactions.
- **Incrementality-based approaches:** estimate causal lift rather than merely assigning credit.

Heuristic attribution is useful for operational reporting, but it should not be confused with causal inference. A campaign that appears near the end of many journeys may receive last-touch credit even if it did not truly create incremental demand.

4.7.2 Selection bias and measurement bias

Attribution systems are affected by:

- **Selection bias:** users exposed to an experience may differ systematically from users who are not.
- **Survivorship bias:** only recorded users or channels appear in the analysis.
- **Cross-device loss:** a user researches on one device and purchases on another without a reliable link.
- **Channel interference:** paid and organic channels interact, making isolated credit assignments misleading.

This is why attribution reporting is a descriptive tool, whereas experimentation is the primary causal tool.

4.8 Metrics layers and decision quality

Executives often ask for “the conversion rate” as though it were a single obvious number. In reality, metric definitions encode business rules. Without a shared metrics layer, each team may compute the same KPI differently.

4.8.1 North Star and supporting metrics

Many commerce organizations choose a high-level outcome metric such as gross merchandise value, contribution margin, or retained active buyers. That metric must then be decomposed into leading indicators:

- traffic volume,
- product-detail engagement,
- add-to-cart rate,
- checkout completion rate,
- repeat purchase rate,
- return rate,
- and support burden.

For educational purposes, it is useful to connect these metrics to event definitions explicitly. For instance:

$$\text{Add-to-cart rate} = \frac{\text{\#sessions with add_to_cart}}{\text{\#sessions with product_view}}$$

This formula is simple, but each term still depends on sessionization rules, bot filtering, and event validity criteria.

4.8.2 Metric dimensions

Commerce metrics usually need to be sliced by:

- channel (web, app, store-assisted, marketplace),
- customer segment (new, returning, VIP, B2B account tier),
- geography and locale,
- device class,
- product category and seller,
- experiment cohort.

A robust metrics layer therefore requires both canonical definitions and dimensional consistency. It should be possible to explain why a number changes, not just to report the number itself.

4.8.3 Retention and cohort logic

Retention metrics are especially sensitive to definition. Possible questions include:

- Did the user return to browse?
- Did the user make another purchase?
- Did the organization place another order from any employee account?
- Did the customer remain active despite returns or failed payments?

Each version may be legitimate, but each reflects a different product and business question. Good analytical practice is to name retention metrics with explicit behavioral criteria rather than using ambiguous shorthand.

4.9 Experimentation in e-commerce

Offline evaluation is valuable, but it cannot fully predict live business impact. Ranking changes alter user behavior. Promotions affect demand patterns. Fraud rules shift attacker behavior. For these reasons, e-commerce systems rely heavily on online experiments.

4.9.1 A/B testing fundamentals

An A/B test compares a treatment against a control by random assignment. The point of randomization is not to guarantee improvement; it is to make the groups comparable in expectation so that observed differences can be interpreted causally.

Key design choices include:

- **Unit of randomization:** user, session, device, seller, or region.
- **Exposure logic:** where and when eligibility is checked.
- **Primary KPI:** the main outcome the experiment is intended to improve.
- **Guardrails:** metrics that protect against harmful side effects such as latency, return rate, or support contacts.
- **Analysis window:** how long outcomes are observed after exposure.

4.9.2 Common experimental pitfalls

- **Contamination:** the same user sees both variants through multiple devices or sessions.
- **Novelty effects:** short-term excitement masks weak long-term performance.
- **Interference:** one user's treatment influences another user, seller, or market.

- **Metric leakage:** the treatment changes how events are logged rather than how users behave.
- **Underpowered tests:** sample sizes are too small to detect meaningful differences.

These pitfalls illustrate why experimentation is partly a data-engineering problem. If exposure logs and outcome metrics do not align, the test result is not credible.

4.9.3 Bandits and adaptive experimentation

Multi-armed bandits are useful when the platform wants to balance learning with near-term reward. Compared with standard A/B tests, bandits typically allocate more traffic to arms that appear to perform better. This can be attractive for:

- promotions on fast-moving catalogs,
- ranking widgets with many interchangeable treatments,
- high-traffic surfaces where opportunity cost is large.

However, bandits are not a universal replacement for A/B testing. They complicate interpretation, can interact badly with delayed rewards, and are less appropriate when decision-makers need a clean causal estimate for a fixed policy choice. In practice, many organizations use A/B tests for clear comparisons and adaptive methods for optimization on mature surfaces.

4.10 From analytics to machine learning features

The same event stream that supports dashboards often becomes the basis for ML features. Examples include:

- recent category views for recommendation,
- failed login or payment attempts for fraud detection,
- coupon usage history for promotion targeting,
- time since last order for churn models,
- return behavior and order defects for seller-quality models.

This creates an important dependency: if event semantics drift, model features drift as well. A model can degrade not only because user behavior changes, but because the measurement pipeline changed underneath it. For this reason, feature definitions should be traceable to governed event definitions.

4.11 Hands-on lab: build a minimal metrics layer

The aim of the lab is to help students bridge raw event logs and business-facing KPIs.

4.11.1 Lab scenario

Assume a storefront emits five core events: `session_started`, `product_viewed`, `add_to_cart`, `checkout_started`, and `order_completed`. Students are given a small synthetic log containing timestamps, session IDs, customer IDs where available, product IDs, and order IDs.

4.11.2 Lab tasks

1. Create a canonical event table with required fields, event validation flags, and ingestion metadata.
2. Build daily metrics for sessions, product views, add-to-cart rate, checkout-start rate, purchase conversion rate, and day-7 repeat purchase.
3. Define at least three data-quality checks and show how failing checks affect one KPI.
4. Create a simple cohort table for first purchase date and compute repeat purchase behavior by cohort.
5. Simulate an experiment exposure field and compute the primary KPI by control versus treatment.

4.11.3 Expected learning outcome

By the end of the lab, students should be able to explain why a metric is not just an aggregation, but the result of event design choices, identity assumptions, and validation rules.

4.12 Managerial and architectural implications

Data foundations are not owned by analysts alone. They require coordination across product, engineering, data, and enterprise systems teams. Architecturally, the following questions matter:

- Which identifiers are authoritative, and where are they mastered?
- Which events belong in the clickstream versus operational event streams?
- How are storefront events linked to ERP objects such as orders, invoices, and shipments?
- What happens when a schema changes in the storefront while ERP integrations remain stable, or vice versa?
- Who approves metric definitions that drive executive dashboards and model targets?

These are governance questions, but they also influence software boundaries. Weak ownership at this layer creates confusion throughout the commerce stack.

4.13 Multiple-choice questions (MCQs)

1. In event-based analytics for e-commerce, which statement is most accurate?
 - (a) Events should be stored without timestamps to save space.
 - (b) Events are most useful when they have a consistent schema, a timestamp, and stable identifiers.
 - (c) Only purchase events matter; click and view events are noise.
 - (d) The best schema changes frequently to match UI updates.
2. Which is an example of a *data quality* problem that can directly harm ML models?
 - (a) A model uses a neural network instead of linear regression.
 - (b) Missing or duplicated events for key actions (e.g., add-to-cart) create biased labels/features.
 - (c) Using SQL instead of Python for ETL.
 - (d) Choosing a smaller batch size during training.
3. Why are offline metrics often insufficient for evaluating an e-commerce ranking model?
 - (a) Offline metrics are always higher than online metrics.
 - (b) Offline metrics cannot measure causal impact under real user behavior and feedback loops.
 - (c) Offline evaluation cannot be computed from logs.
 - (d) Offline metrics are not used in industry.
4. In an A/B test, the primary purpose of randomization is to:
 - (a) Increase model complexity.
 - (b) Ensure the treatment and control groups are comparable in expectation.
 - (c) Eliminate the need for monitoring.
 - (d) Guarantee the KPI improves.
5. A common *identity resolution* challenge in e-commerce is:
 - (a) Mapping product IDs to SKU IDs.
 - (b) Linking the same person across devices/sessions while respecting privacy constraints.
 - (c) Computing NDCG@ k efficiently.
 - (d) Selecting a cloud region.

Answer key

1. (b)
2. (b)
3. (b)
4. (b)
5. (b)

4.14 Exercises (short)

1. Propose an event taxonomy for a storefront: list at least 10 events (search, view, add-to-cart, checkout steps, purchase, return) and for each event specify 3–5 key fields.
2. Define a North Star metric for an e-commerce product and break it into 3–5 leading indicators. Explain how you would compute each from event logs.
3. Design a simple A/B test for a recommendation widget. Specify: unit of randomization, primary KPI, guardrail metrics, and an example of a bias/confounder to watch for.

4.15 Mini-case (Odoo-linked, vendor-neutral)

Scenario: A hybrid B2C/B2B company uses Odoo (Sales, Inventory, Accounting, CRM). The storefront and mobile app generate clickstream events.

Task: Design a minimal data model that supports both analytics and ML:

- Define how you will link clickstream identities to Odoo customers/partners (when possible) and how you handle guests.
- Propose which entities should have stable IDs across systems (customer, product, order, shipment) and which are channel-specific.
- Specify 5 “data contracts” (producer $\rightarrow$ consumer) that reduce breakage when the UI or ERP changes.

Chapter 5

Enterprise Architecture for E-Commerce (Week 4)

5.1 Learning objectives

- Explain how enterprise architecture (EA) helps align business goals, operating models, and technology decisions in e-commerce.
- Use capability maps, value streams, domain boundaries, and system-of-record thinking to structure a commerce platform.
- Distinguish current-state architecture from target architecture and transition planning.
- Analyze how AI capabilities fit into enterprise architecture without turning the architecture into a collection of isolated point solutions.

5.2 Why enterprise architecture matters in e-commerce

E-commerce systems evolve quickly. New channels, promotions, AI features, logistics partnerships, and regulatory constraints are introduced faster than many enterprise systems were originally designed to handle. Without architectural discipline, organizations often accumulate fragmented storefronts, duplicated customer records, inconsistent pricing logic, and analytics pipelines that are disconnected from operational truth.

Enterprise architecture is useful because it provides a way to reason across these layers at once. It connects:

- business strategy,
- organizational capabilities,
- application boundaries,
- data ownership,

- integration patterns,
- and technology constraints.

In an AI-driven commerce setting, the architectural question is not only “which model should we use?” It is also:

- Which systems generate the data that the model depends on?
- Which domain owns the business decision that the model influences?
- Which workflow must remain reliable when the model is unavailable or uncertain?
- How are predictions turned into operational actions in storefront, CRM, ERP, or support systems?

5.3 Enterprise architecture as a set of views

An architecture description is easier to manage when it is divided into views. For e-commerce, four views are especially practical:

- **Business architecture:** capabilities, value streams, organization roles, policies, and business outcomes.
- **Application architecture:** major systems, services, platform modules, ownership boundaries, and interactions.
- **Data architecture:** entities, identifiers, systems of record, analytical pipelines, and governance rules.
- **Technology architecture:** hosting environment, integration infrastructure, observability stack, identity/security, and platform services.

These views should not be treated as isolated documents. They are projections of one system. For example, a business capability such as pricing must connect to application services, authoritative data sources, and runtime controls.

5.4 Capability maps

A capability map describes what an organization must be able to do, independent of how the capability is implemented today. That distinction matters because teams, products, and vendor choices can change while core business abilities remain relatively stable.

5.4.1 Core e-commerce capabilities

An AI-enabled commerce organization often needs capabilities such as:

- customer acquisition and campaign orchestration,
- product information and catalog management,
- search, discovery, and merchandising,
- personalization and recommendation,
- pricing, promotions, and offer management,
- cart, checkout, and payment orchestration,
- order management and fulfillment coordination,
- returns, refunds, and post-purchase support,
- customer profile and loyalty management,
- analytics, experimentation, and decision support,
- security, privacy, fraud, and governance.

5.4.2 Why capability maps are useful

Capability maps help teams avoid jumping directly into solution design. For example, if an organization says “we need a recommendation engine,” the architectural question should be expanded:

- Is personalization a standalone capability, or is it part of a broader discovery capability?
- Which parts of the recommendation workflow belong to the storefront experience layer, and which belong to a shared data platform?
- Which existing operational capabilities must change to make recommendations effective, such as promotions, inventory awareness, or content quality?

Capability maps also expose imbalance. A company may invest heavily in customer acquisition and storefront design while underinvesting in returns, data governance, or seller operations. Architecture should reveal such weaknesses before they become scaling problems.

5.5 Value streams

Capabilities describe *what* the organization can do. Value streams describe *how value moves* through the business from trigger to outcome.

5.5.1 Typical e-commerce value streams

Important value streams include:

- **Browse-to-buy:** discovery, consideration, cart, checkout, payment, confirmation.
- **Order-to-cash:** order capture, validation, fulfillment, invoicing, settlement, financial posting.
- **Procure-to-stock:** replenishment planning, supplier ordering, receiving, storage, inventory updates.
- **Return-to-resolution:** return request, approval, pickup/drop-off, inspection, refund or exchange.
- **Lead-to-account:** acquisition, registration, qualification, first purchase, retention.

5.5.2 Why value streams matter architecturally

Value streams expose cross-domain handoffs. For example, “browse-to-buy” may touch search, catalog, pricing, promotions, checkout, payment, fraud, order management, and analytics. If the stream depends on unclear ownership or duplicated business rules, failures appear at handoff points rather than inside a single system.

Value-stream mapping is especially valuable in hybrid commerce environments where storefront systems and ERP systems divide responsibility. Students should learn to ask not only which system executes a step, but which system is accountable for the business meaning of that step.

5.6 Domain boundaries and bounded contexts

Large commerce platforms become difficult to scale when everything is treated as one application or one shared database. Architecture therefore needs domain boundaries with explicit ownership.

5.6.1 Typical commerce domains

Common domains include:

- **Catalog:** products, attributes, categories, media, content enrichment.
- **Pricing and promotions:** price lists, discount rules, coupon logic, campaign offers.
- **Inventory and availability:** stock position, reservations, safety stock, availability promises.
- **Cart and checkout:** session basket state, shipping options, tax estimation, checkout workflow.
- **Payments:** payment intents, authorization, capture, refund orchestration.
- **Orders:** order lifecycle, status management, amendments, cancellations.

- **Customer and identity:** profile, authentication, preferences, loyalty, segmentation.
- **Fulfillment and logistics:** warehouse tasks, shipment creation, delivery events, returns routing.
- **Analytics and experimentation:** event collection, metric definitions, experiment assignment.

5.6.2 Good boundaries versus bad boundaries

A good boundary reduces coupling and clarifies ownership of behavior and data. A poor boundary creates ambiguity such as:

- pricing rules duplicated in storefront, ERP, and marketing tools,
- customer attributes updated independently in CRM and commerce applications,
- order status interpreted differently by warehouse, finance, and support teams,
- product IDs that differ between catalog, analytics, and inventory systems.

Architectural boundaries should therefore be tested against real business scenarios. If a business change requires editing logic in five systems with no clear owner, the boundary is likely wrong.

5.7 Systems of record and systems of engagement

Not every system should be authoritative for every type of data. One of the most important EA concepts in commerce is deciding the *system of record* for each key entity.

5.7.1 Typical patterns

- The **commerce layer** often owns high-velocity interaction data such as carts, storefront sessions, and recommendation exposures.
- The **ERP layer** often owns invoices, financial postings, stock movements, supplier records, and contractual B2B customer data.
- A **PIM or catalog service** may own enriched product content while ERP owns procurement-oriented product master fields.
- A **data platform** may own analytical aggregates and features but should not silently become the operational source of truth.

5.7.2 Systems of engagement

Storefronts, mobile apps, support consoles, and seller portals are often systems of engagement rather than systems of record. They present, collect, and orchestrate information, but they may not be authoritative for final commercial truth. Confusing these roles leads to reconciliation failures.

For example, a storefront may display an order as “confirmed” while ERP has not yet accepted it due to credit rules or stock constraints. Architecture must define which state is customer-facing, which state is financially binding, and how transitions are synchronized.

5.8 Reference architecture for AI-driven e-commerce

There is no single universal architecture, but a practical reference pattern often includes the following layers:

- **Experience layer:** web, mobile, agent interfaces, seller/admin portals.
- **Commerce services layer:** catalog, search, pricing, cart, checkout, order orchestration, customer services.
- **Enterprise systems layer:** ERP, CRM, WMS, finance, procurement, supplier management.
- **Integration layer:** APIs, event streams, workflow orchestration, master-data synchronization.
- **Data and AI layer:** event collection, warehouse/lakehouse, feature pipelines, experimentation, model training, model serving.
- **Cross-cutting controls:** identity, security, privacy, audit, observability, governance.

5.8.1 Where AI fits

AI should be mapped to business capabilities rather than inserted as an isolated layer that “does intelligence.” Examples include:

- search ranking and recommendations within discovery,
- fraud scoring within checkout and payments,
- demand forecasting within inventory and replenishment,
- case summarization and response assistance within support,
- anomaly detection within operations and finance.

This mapping matters because it keeps accountability clear. The discovery domain owns business outcomes for search and recommendation even if the underlying models are developed by a centralized data team.

5.9 Current state, target state, and transition architecture

Architecture is not useful if it only documents the present. Organizations need a target state that guides investment and sequencing.

5.9.1 Current-state analysis

Current-state architecture should identify:

- duplicated systems and overlapping responsibilities,
- brittle integrations and manual reconciliation steps,
- data ownership conflicts,
- missing capabilities,
- and operational pain points such as latency, outages, or inconsistent KPIs.

5.9.2 Target architecture

Target architecture is a future-state design that expresses intended business and technical structure. It is not a vendor brochure and not a line-by-line implementation plan. It should answer:

- Which domains and systems will own which capabilities?
- Which integrations will be synchronous, asynchronous, or batch?
- Which entities will be mastered where?
- Which platforms are shared across business units?
- Which AI capabilities will be productized versus embedded locally?

5.9.3 Transition architecture

Most organizations cannot move directly from current state to target state. Transition architecture defines intermediate steps. Examples include:

- separating catalog and pricing ownership before replacing checkout,
- introducing a shared event stream before centralizing experimentation,
- establishing a customer identity service before rolling out cross-channel personalization.

This is important in enterprise settings because a theoretically elegant target architecture may fail if the migration path is not operationally realistic.

5.10 Architectural concerns specific to AI

AI adds architectural concerns beyond model selection. Representative concerns include:

- **Training-serving consistency:** features must mean the same thing offline and online.
- **Latency budgets:** recommendation or fraud decisions may need sub-second responses.
- **Fallback behavior:** the commerce workflow must continue safely if a model is unavailable.
- **Observability:** systems must track not only uptime but also prediction quality, drift, and business impact.
- **Governance:** model changes may alter business policy and therefore require approval, auditability, and rollback procedures.

In many organizations, the failure is not that models are weak. The failure is that model outputs are not architecturally integrated into business workflows with clear ownership and control points.

5.11 Operating model and team topology

Architecture is influenced by team structure. A domain with clear technical boundaries but no accountable product owner often performs badly in practice.

5.11.1 Typical team patterns

- **Domain-aligned product teams:** own specific business capabilities such as checkout or catalog.
- **Platform teams:** provide shared infrastructure such as identity, observability, experimentation, or data platforms.
- **Enterprise systems teams:** maintain ERP, CRM, accounting, and operational master-data platforms.
- **Central AI or data teams:** develop shared methods, governance, and model platforms while collaborating with domain teams on use cases.

An architecture is more likely to succeed when ownership in the diagrams matches ownership in the organization. If a system is “owned by everyone,” it is usually owned by no one.

5.12 Artifacts for students and architects

A minimal EA artifact set for an e-commerce program may include:

- a capability map,
- one or more value stream maps,
- a domain map with ownership boundaries,
- a system-context diagram,
- a system-of-record matrix for key entities,
- and a target-state architecture note with transition priorities.

These artifacts should be consistent with each other. If the domain map says pricing is a shared domain but the system-of-record matrix shows three authoritative pricing systems, the architecture still contains a contradiction.

5.13 Hands-on artifacts

The key deliverable for this chapter is a compact but defensible architecture package:

- an e-commerce capability map,
- a domain model showing clear ownership,
- and a short target-architecture note that explains systems of record and integration choices.

5.14 Managerial implications

Enterprise architecture is sometimes dismissed as documentation overhead. In e-commerce, that is a mistake. Architectural ambiguity produces real costs:

- slow product delivery because every change triggers cross-team negotiation,
- financial reconciliation issues from conflicting commercial truth,
- duplicated AI initiatives that solve the same problem inconsistently,
- and weak governance when no domain clearly owns a high-impact decision.

Good architecture does not eliminate complexity. It makes complexity explicit, owned, and therefore manageable.

5.15 Multiple-choice questions (MCQs)

1. In enterprise architecture, a *capability map* primarily describes:
 - (a) A list of specific microservices and their endpoints.
 - (b) *What* the business does (stable abilities), independent of *how* it is implemented.
 - (c) The physical network topology of the data center.
 - (d) A Gantt chart of the project plan.
2. A good domain boundary (e.g., for a “Catalog” domain) typically aims to:
 - (a) Maximize shared database tables across all teams.
 - (b) Minimize coupling and define clear ownership of data and behaviors.
 - (c) Ensure all operations are synchronous.
 - (d) Avoid having APIs.
3. Which artifact is most suitable for describing how business value is delivered end-to-end?
 - (a) Value stream map
 - (b) Entity-relationship diagram (ERD)
 - (c) Source code repository layout
 - (d) TLS configuration file
4. “Target architecture” is best described as:
 - (a) The current-state system diagram.
 - (b) A future-state design that guides decisions and sequencing of change.
 - (c) A vendor product brochure.
 - (d) A test plan for unit tests.
5. In an AI-driven e-commerce context, which is the best example of an *architectural concern* (not a model choice)?
 - (a) Whether to use XGBoost or logistic regression
 - (b) How to ensure training-serving consistency and monitor drift
 - (c) Whether to use k -means or DBSCAN
 - (d) Whether to use SGD or Adam

Answer key

1. (b)
2. (b)
3. (a)
4. (b)
5. (b)

5.16 Exercises (short)

1. Draft a capability map for an AI-enabled e-commerce organization. Include at least: customer acquisition, discovery, pricing/promotions, order management, fulfillment/returns, customer support, data/analytics, and governance/security.
2. Choose one value stream (e.g., “order-to-cash” or “browse-to-buy”). Identify 5–8 steps and list the owning domain/system for each step (commerce vs ERP vs logistics vs payment).
3. Propose domain boundaries for: catalog, pricing, orders, payments, and customer profiles. For each boundary, name the system-of-record for key data entities.

5.17 Mini-case (Odoo-linked, vendor-neutral)

Scenario: Your company is hybrid B2C/B2B. Odoo is used for core ERP flows (Sales, Inventory, Accounting, CRM). A separate commerce layer handles web/mobile experiences and AI features.

Task: Produce a short “target architecture” note:

- Draw a capability-to-system mapping (capabilities → commerce layer vs Odoo vs data platform).
- Define the *system of record* for customer, product, price list, order, invoice, and shipment.
- Identify 3 integration risks (data consistency, latency, duplicate sources of truth) and propose mitigations.

Chapter 6

Enterprise Integration Patterns for E-Commerce (Week 5)

6.1 Learning objectives

- Compare synchronous APIs, asynchronous events, and batch integrations in enterprise commerce environments.
- Explain why reliability patterns such as idempotency, retries, dead-letter handling, and the outbox pattern are essential in distributed commerce systems.
- Design integration contracts that reduce coupling between storefront, ERP, payment, logistics, and data platforms.
- Evaluate integration choices in terms of latency, resilience, consistency, and operational complexity.

6.2 Why integration patterns matter

E-commerce systems rarely operate as a single application. Even a modest commerce stack may involve a storefront, search service, promotions engine, payment provider, ERP, warehouse system, shipping partner, CRM, and analytics platform. These systems must exchange data continuously while still tolerating delays, retries, outages, and partial failures.

The architectural problem is not merely how to “connect systems.” The deeper question is how to coordinate business processes when no single runtime boundary contains the entire workflow. An order may be accepted in the storefront, authorized by a payment gateway, validated by fraud controls, persisted in commerce services, exported to ERP, and later updated by warehouse and carrier events. Each handoff introduces failure modes that integration design must anticipate.

6.3 Integration styles

Integration patterns can be grouped into a small number of common styles.

6.3.1 Synchronous request-response

In synchronous integration, one system calls another and waits for a response. Typical examples include:

- payment authorization during checkout,
- tax estimation,
- retrieving customer-specific pricing,
- checking address validation services,
- fetching live delivery options.

The main advantage is immediacy. The caller receives a result within the user flow and can make a direct decision. The main disadvantage is coupling. The caller now depends on the callee's availability, latency, contract stability, and error behavior.

6.3.2 Asynchronous events

In asynchronous integration, a producer publishes an event and does not wait for every downstream consumer to finish. Examples include:

- `OrderCreated`,
- `InvoicePosted`,
- `ShipmentDispatched`,
- `InventoryAdjusted`,
- `CustomerRegistered`.

This style supports loose coupling and fan-out. Multiple consumers can react to the same business event without the producer needing to know all of them in advance. However, event-driven systems trade immediacy for eventual consistency and demand stronger operational discipline.

6.3.3 Batch and file-based integration

Batch integration remains common in enterprise settings despite modern API and event tooling. Examples include nightly product exports, finance reconciliation files, bulk supplier updates, and large historical loads into analytical platforms.

Batch is often appropriate when:

- latency requirements are low,
- partner systems are legacy or externally controlled,
- very large data volumes must be moved efficiently,
- or the business process is administrative rather than interactive.

Students should not assume batch is “obsolete.” It is often the right pattern for the wrong place if teams are not explicit about where it belongs.

6.4 Choosing the right integration style

No single integration style is universally best. The right choice depends on business semantics.

6.4.1 Questions to ask

When selecting a pattern, architects should ask:

- Does the user or operator need an immediate answer?
- What is the tolerance for stale data or delayed propagation?
- What happens if the downstream system is unavailable?
- Can the workflow continue in a degraded mode?
- Is the interaction command-like, query-like, or event-like?
- How many downstream consumers must react to the information?

6.4.2 Typical choices in commerce

- **Checkout payment authorization:** usually synchronous because the customer needs immediate confirmation.
- **Order export to ERP:** often asynchronous because the storefront should not block on internal back-office processing.
- **Inventory availability updates:** often event-driven with cache refreshes and reconciliation support.
- **Financial reporting loads:** often batch.
- **Shipment tracking updates:** usually asynchronous because carriers and fulfillment processes are naturally event-driven.

6.5 API-first integration

APIs are essential in modern commerce because they define explicit service boundaries. An API-first mindset means the contract is treated as a product artifact rather than an afterthought.

6.5.1 Good API design principles

- Model APIs around business resources and actions, not internal database tables.
- Make required fields, validation rules, and error semantics explicit.
- Use stable identifiers and versioning policies.
- Design for failure and retries rather than assuming a perfect network.
- Document authorization requirements and rate limits.

6.5.2 Commands versus queries

Many integration problems become clearer when teams distinguish:

- **Queries:** read current information, such as “get available shipping options.”
- **Commands:** request a state change, such as “create order” or “capture payment.”

Commands are especially sensitive to retry behavior. If a request is repeated because of a timeout, the system must avoid creating duplicate side effects. This leads directly to idempotency design.

6.6 Idempotency

Idempotency means that repeating the same operation does not create additional unintended effects. This is one of the most important patterns in distributed order processing.

6.6.1 Why idempotency is required

Networks fail in ambiguous ways. A caller may send a command, the server may execute it, and the response may be lost. If the caller retries without an idempotency mechanism, duplicate orders, duplicate refunds, or duplicate invoice creation may occur.

6.6.2 Idempotency keys

For high-value commands such as order creation or payment capture, the client or upstream service often sends an idempotency key. The receiving service stores the key with the outcome of the first successful execution and returns the same result for subsequent retries with the same key.

Good practice includes:

- making the key stable for a business attempt,
- scoping it to the operation type,
- storing it long enough to cover realistic retry windows,
- and linking it to request payload validation to prevent misuse.

6.7 Events and event-driven architecture

An event represents something that already happened in the business domain. Examples include order creation, payment success, stock decrement, and shipment delivery.

6.7.1 Domain events versus technical events

Teams should distinguish:

- **Domain events:** meaningful business facts such as `OrderCanceled`.
- **Technical events:** operational notifications such as cache refresh or job completion.

Domain events are more reusable because multiple consumers can act on them without knowing the producer's implementation details. If an event stream is dominated by low-level technical signals, consumers become tightly coupled to internals rather than business meaning.

6.7.2 Benefits of event-driven integration

- Producers and consumers can evolve more independently.
- Multiple downstream systems can react to the same fact.
- Workloads can be buffered and retried.
- Slow consumers do not necessarily block user-facing systems.

6.7.3 Costs of event-driven integration

- Debugging becomes harder because workflows are distributed across time and services.
- Consumers must tolerate reordering, duplication, and replay.
- Monitoring and tracing become mandatory rather than optional.
- Business users must understand eventual consistency.

6.8 Reliability patterns in distributed workflows

Integration reliability is not achieved by a single technology choice. It is achieved by combining patterns that reduce ambiguity under failure.

6.8.1 Retries and backoff

Transient failures are normal. Consumers and callers should usually retry with bounded exponential backoff rather than failing immediately. However, retries without idempotency or circuit breaking can amplify incidents.

6.8.2 Dead-letter handling

If a message repeatedly fails processing, it should not block the entire flow indefinitely. Dead-letter queues or equivalent quarantine mechanisms allow operators to inspect problematic messages, repair data if needed, and replay safely.

6.8.3 Timeouts and circuit breakers

Synchronous integrations should use explicit timeouts. If a downstream system becomes slow or unavailable, circuit breakers can stop a cascade of blocked requests. In commerce, this is often paired with degraded behavior such as:

- temporarily disabling a non-critical recommendation widget,
- falling back to cached delivery promises,
- or routing support cases for manual review.

6.8.4 Reconciliation jobs

Not all inconsistencies can be prevented in real time. Periodic reconciliation is a practical safeguard, especially for orders, invoices, stock, and refunds. Reconciliation does not replace good design, but it provides a controlled recovery mechanism when distributed flows drift apart.

6.9 The outbox pattern

One of the most common distributed data bugs is the dual-write problem. For example, a service may write a new order to its database and then publish an event. If the database write succeeds but the event publish fails, downstream consumers never learn about the order. If the event is published but the database write fails, consumers act on a fact that is not actually committed.

6.9.1 How the outbox pattern works

The outbox pattern addresses this by:

- writing the business state change and an “outbox record” in the same local transaction,
- then having a separate relay process publish the outbox record to the event broker,
- marking the record as sent only after successful publish.

This pattern does not magically create exactly-once delivery across the whole system. What it does is make producer-side event publication reliable relative to the local transaction boundary.

6.9.2 Why it matters in commerce

Commerce workflows are especially sensitive to dual-write failures because duplicated or missing side effects can affect revenue, customer trust, and financial records. Orders, refunds, shipment creation, and invoice triggers are all candidates for outbox-style publishing.

6.10 Message contracts and schema governance

Events and APIs require contracts. Without governed contracts, integration platforms become brittle collections of undocumented assumptions.

6.10.1 What a contract should define

A useful contract specifies:

- message or API name,
- required and optional fields,
- field semantics and units,
- identifier rules,
- allowed status transitions,
- versioning and deprecation policy,
- ownership and escalation path.

6.10.2 Backward compatibility

Contract evolution should aim for backward compatibility whenever possible. Common practices include:

- adding optional fields rather than repurposing existing ones,
- preserving existing semantics across versions,
- introducing new event names or explicit versions for breaking changes,
- and validating consumer readiness before removing old fields.

In enterprise commerce, breaking contract changes are expensive because downstream systems often include ERP integrations, partner processes, and compliance reporting.

6.11 Observability for integrations

Integration systems need observability that spans services and transport layers. Logs alone are insufficient.

6.11.1 Minimum signals

An integration operating model should include:

- **Logs:** structured records with correlation IDs, idempotency keys, and error codes.
- **Metrics:** request rates, latency, failure rates, retry counts, dead-letter counts, consumer lag.
- **Traces:** end-to-end flow visibility across API calls, event publication, and downstream processing.
- **Business counters:** order export success, invoice posting delay, shipment-update freshness, reconciliation mismatch rates.

6.11.2 Correlation and traceability

Cross-system workflows should carry shared identifiers such as order ID, correlation ID, and experiment or campaign context where relevant. Without this, incident response becomes guesswork. In distributed commerce systems, the ability to answer “what happened to order X?” is an operational requirement, not a nice-to-have feature.

6.12 iPaaS, ESB, brokers, and workflow orchestration

Enterprise integration discussions often involve platform choices such as ESBs, iPaaS products, event brokers, and orchestration tools. Students should understand their roles conceptually rather than treating them as interchangeable buzzwords.

6.12.1 iPaaS and ESB concepts

- **ESB-style platforms:** historically emphasized centralized mediation, transformation, and routing.
- **iPaaS platforms:** emphasize managed connectors, workflow automation, and cloud integration convenience.

These tools can accelerate delivery, especially where enterprise applications and partner ecosystems are involved. But they can also become hidden concentrations of logic if every business rule is embedded in a central integration layer.

6.12.2 Event brokers and streams

Message brokers and streaming platforms provide transport for asynchronous communication. Their value lies not just in moving messages, but in enabling buffering, replay, scaling, and decoupling. However, the transport does not solve data modeling or ownership problems by itself.

6.12.3 Workflow orchestration

Some workflows are easier to manage with explicit orchestration, especially long-running processes such as returns handling, supplier exceptions, or manual fraud review. Orchestration can improve clarity, but over-centralization can also produce a fragile control tower if every process depends on one orchestrator.

6.13 Integration anti-patterns

Common anti-patterns in commerce integration include:

- direct database-to-database coupling across domains,
- duplicating business rules independently in multiple systems,
- publishing events without stable schemas or owners,
- assuming exactly-once delivery where only at-least-once semantics are realistic,
- using synchronous calls for every workflow step,

- and treating reconciliation as proof that real-time design can be ignored.

Students should be able to identify these anti-patterns because they appear frequently in systems that “worked at small scale” but fail under growth.

6.14 Hands-on lab: modeling an order lifecycle

The practical goal of this chapter is to connect abstract patterns to one realistic workflow.

6.14.1 Lab scenario

Assume a headless storefront, a payment service, an order service, Odoo, and a shipment service. Students are asked to model the lifecycle from checkout submission to shipment dispatch.

6.14.2 Lab tasks

1. Identify which interactions are synchronous and which are asynchronous.
2. Define an idempotency strategy for order creation and payment capture.
3. Design an `OrderCreated` event contract and one downstream consumer.
4. Show where the outbox pattern would be applied.
5. Define the operational metrics needed to detect broken flows.

6.14.3 Expected learning outcome

By the end of the lab, students should be able to explain integration design in terms of business guarantees rather than technology labels alone.

6.15 Managerial and architectural implications

Integration choices shape business agility. If every new channel or AI feature requires fragile point-to-point connections, innovation slows. If teams overuse asynchronous flows without good observability, the organization loses operational clarity.

Architecture therefore requires a balanced approach:

- synchronous calls where immediate business decisions are required,
- events where decoupling and fan-out matter,
- batch where latency is not critical,

- and governance across all three.

Well-designed integration is not invisible plumbing. It is a core business capability in enterprise commerce.

6.16 Multiple-choice questions (MCQs)

1. Which statement best characterizes synchronous vs asynchronous integration?
 - (a) Synchronous integration cannot fail; asynchronous always fails.
 - (b) Synchronous integration couples availability and latency; asynchronous trades immediacy for resilience and decoupling.
 - (c) Asynchronous integration is only for analytics; synchronous is only for ERP.
 - (d) They are equivalent as long as JSON is used.
2. Why is idempotency important in order processing APIs?
 - (a) It makes responses larger.
 - (b) It prevents duplicate effects when requests are retried.
 - (c) It eliminates the need for authentication.
 - (d) It guarantees exactly-once message delivery in all systems.
3. The *outbox pattern* is primarily used to:
 - (a) Store images for the catalog.
 - (b) Reliably publish events when database writes succeed, avoiding dual-write inconsistencies.
 - (c) Encrypt API traffic.
 - (d) Replace a message broker.
4. In event-driven systems, a *consumer* should typically be designed to be:
 - (a) State-less but non-retryable
 - (b) Idempotent and retryable
 - (c) Dependent on message ordering always being perfect
 - (d) Tightly coupled to producer database schemas
5. Which is a common reason to choose events over direct API calls for some flows?
 - (a) To avoid defining schemas/contracts
 - (b) To decouple systems and allow multiple downstream consumers without changing the producer
 - (c) To guarantee zero latency
 - (d) To avoid monitoring

Answer key

1. (b)
2. (b)
3. (b)
4. (b)
5. (b)

6.17 Exercises (short)

1. For each flow, choose sync API or async events and justify: (i) “place order”, (ii) “update inventory availability”, (iii) “invoice posted”, (iv) “shipment delivered”, (v) “customer profile updated”.
2. Design an idempotency strategy for the “create order” API. Specify the idempotency key and how long it is retained.
3. Write a short failure story: a dual-write bug between database and message broker. Explain how the outbox pattern prevents it.

6.18 Mini-case (Odoo-linked, vendor-neutral)

Scenario: A headless storefront integrates with Odoo for orders, inventory, invoices, and CRM updates.

Task: Propose an integration design that supports scale and reliability:

- Define the key domain events (at least 8) and which system publishes each.
- Identify where you need synchronous APIs (e.g., payment authorization) vs asynchronous events (e.g., shipment updates).
- Specify two data contracts and a versioning strategy to avoid breaking changes.
- List the minimum observability signals (logs/metrics/traces) required to operate the integration in production.

Chapter 7

ERP/CRM/SCM Integration in Commerce (Week 6)

7.1 Learning objectives

- Explain how commerce platforms integrate with ERP, CRM, and supply-chain systems across operational workflows.
- Identify systems of record for key entities such as customer, product, price list, order, invoice, and shipment.
- Analyze master data management, process ownership, and consistency risks in hybrid commerce architectures.
- Design integration flows for order-to-cash, returns, and B2B pricing using Odoo as a reference ERP environment.

7.2 Why ERP/CRM/SCM integration is central to commerce

Modern e-commerce experiences are often built in specialized commerce platforms, but commercial truth does not live only in the storefront. Financial commitments, stock movements, supplier processes, customer account structures, invoicing, and back-office controls typically depend on enterprise systems such as ERP, CRM, and warehouse or supply-chain platforms.

This means a commerce platform cannot be architected only as a customer-facing product. It must also participate in enterprise operating flows. If that integration is weak, the result is familiar:

- orders that appear confirmed online but fail in back-office processing,
- inventory figures that differ between storefront and warehouse,
- duplicate customer records across channels,
- B2B price lists that drift between sales and commerce systems,

- returns and refunds that create reconciliation problems in finance.

The core architectural task is not simply to “connect Odoo” or any ERP. It is to define ownership, process boundaries, and trustworthy data movement across the enterprise stack.

7.3 The enterprise systems perspective

In a hybrid commerce architecture, enterprise systems typically serve different roles.

7.3.1 ERP

ERP systems coordinate structured operational and financial processes. Using Odoo as a reference point, relevant modules often include:

- **Sales:** quotations, sales orders, customer terms, sales workflows,
- **Inventory:** stock moves, reservations, warehouse operations, replenishment,
- **Accounting:** invoices, credit notes, payments, reconciliation,
- **CRM:** leads, opportunities, account history, commercial context.

7.3.2 CRM

CRM systems focus on relationship history, account management, sales interactions, lead qualification, and customer-service context. In some organizations CRM and ERP overlap. What matters architecturally is not the product label but the business responsibility assigned to the system.

7.3.3 SCM and warehouse systems

Supply-chain and warehouse platforms manage sourcing, supplier coordination, inbound logistics, warehouse execution, and delivery orchestration. In smaller environments these functions may be embedded in ERP. In larger environments they may be split into specialized systems such as WMS, TMS, or procurement platforms.

7.4 Commerce versus enterprise ownership

A recurring source of failure in commerce architecture is treating ownership as a technical detail rather than a business decision. Teams must decide which domain owns each important entity and process state.

7.4.1 Common ownership patterns

- The **commerce layer** often owns web sessions, carts, checkout state, personalization context, and channel-specific merchandising behavior.
- The **ERP layer** often owns invoices, accounting entries, stock movements, contractual customer records, and internal fulfillment commitments.
- The **CRM layer** may own lead status, sales activity, account segmentation, and service interactions.
- The **data platform** may own analytical aggregates and experimentation outputs but should not silently become the master for operational entities.

7.4.2 Why ownership must be explicit

If pricing is editable in both storefront and ERP, or if both commerce and CRM create customer records independently, inconsistencies are not accidental. They are structural. Integration design cannot fix a domain-ownership problem after the fact; it must start by making ownership explicit.

7.5 Systems of record for key entities

One of the most important artifacts in enterprise commerce architecture is a system-of-record matrix. Representative entity choices include the following.

7.5.1 Customer

Customer identity is often split across systems:

- the storefront may own channel account credentials and preferences,
- ERP may own bill-to and ship-to parties for financial processing,
- CRM may own relationship context, lead lifecycle, and account history.

The correct design is usually not “one system owns everything.” Instead, different aspects of customer identity are mastered in different places and linked through governed identifiers.

7.5.2 Product

Product ownership is also frequently split:

- ERP may own procurement and internal stock-keeping information,
- a catalog or PIM layer may own rich digital merchandising content,

- analytics may derive behavioral product features.

The architecture must distinguish the commercial product definition from channel presentation and from analytical enrichment.

7.5.3 Price list

Price lists are especially sensitive in B2B environments. Negotiated terms, tax logic, discount hierarchies, and approval rules often belong closer to ERP or enterprise sales processes than to storefront presentation logic. If the storefront computes pricing independently while ERP maintains contractual prices, disputes are inevitable.

7.5.4 Order

The commerce platform often captures the initial order intent, but the enterprise question is whether ERP becomes the authoritative record for downstream operational and financial lifecycle. This must be decided explicitly. Some organizations treat commerce as the order system of record and export to ERP. Others treat the storefront order as a pre-order until ERP accepts and materializes it.

7.5.5 Invoice and payment reconciliation

Invoices, ledger effects, and payment reconciliation usually belong in ERP or finance-owned systems. Even if the storefront shows payment confirmation, the enterprise system remains authoritative for accounting truth.

7.5.6 Shipment

Shipment truth may belong in warehouse or logistics systems even if ERP or storefront mirrors the status. The architecture should distinguish customer-facing delivery status from the internal system that owns shipment execution and exception handling.

7.6 Master data management (MDM)

Master data management is the discipline of defining and governing shared core entities across systems. In commerce, the most common shared entities include customer, product, supplier, location, and price list.

7.6.1 Why MDM matters

Without MDM, organizations accumulate:

- duplicate customer records created by multiple channels,

- product records with inconsistent category or attribute definitions,
- seller or supplier entities that differ across procurement and commerce tools,
- incompatible codes used by analytics, ERP, and storefront systems.

7.6.2 Practical MDM principles

- Define authoritative identifiers and mapping rules.
- Separate creation rights from consumption rights.
- Document which attributes are mastered where.
- Maintain reconciliation workflows for conflicts and duplicates.
- Treat data stewardship as an operating responsibility, not an afterthought.

MDM is often perceived as bureaucratic. In reality, it is a scaling mechanism for preventing repeated operational confusion.

7.7 Order-to-cash in integrated commerce

Order-to-cash is one of the most important enterprise commerce value streams because it crosses customer experience, logistics, and finance.

7.7.1 Typical flow

A simplified integrated order-to-cash process may include:

1. customer places an order in the storefront,
2. payment is authorized or payment terms are validated,
3. order is persisted in commerce services,
4. order is exported or synchronized to ERP,
5. inventory is reserved or confirmed,
6. picking, packing, and shipment are triggered,
7. invoice is issued,
8. payment is captured or reconciled,
9. order and fulfillment statuses propagate back to customer-facing systems.

7.7.2 Architectural decisions inside the flow

Key design choices include:

- whether the storefront can confirm an order before ERP acceptance,
- whether stock is reserved in real time or reconciled later,
- whether ERP or commerce owns status transitions,
- how payment confirmation maps to invoicing and revenue recognition,
- which events trigger downstream systems such as warehouse or support tools.

7.7.3 Failure modes

Common failure modes include:

- the order is captured online but fails to post in ERP,
- payment succeeds but the order export fails,
- the storefront shows available stock that has already been consumed elsewhere,
- fulfillment completes but shipment status never propagates back to customer channels,
- invoice creation and refund processing become inconsistent during cancellations.

These are not edge cases. They are the expected complexity of distributed commercial workflows.

7.8 Inventory, fulfillment, and availability

Inventory integration is especially difficult because availability is both operationally critical and highly time-sensitive.

7.8.1 Inventory concepts that should not be conflated

- **On-hand inventory:** physical stock currently available in storage.
- **Reserved inventory:** stock allocated to pending commitments.
- **Available-to-promise:** stock that can still be sold considering current commitments and business rules.
- **Future availability:** stock expected from replenishment or transfer.

If the storefront treats all four as the same number, overselling and customer disappointment become likely.

7.8.2 Synchronization strategies

Common strategies include:

- near-real-time events for stock adjustments,
- short-lived availability caches,
- reservation services for high-demand items,
- periodic reconciliation against warehouse truth.

The right design depends on volume, latency needs, and oversell tolerance. Enterprise architecture should make these tradeoffs explicit rather than hiding them behind “inventory sync” language.

7.9 Returns and refunds

Post-purchase flows are often under-modeled in early architecture work, yet they are essential to both customer trust and accounting correctness.

7.9.1 Typical return flow

A return process may involve:

1. customer initiates a return,
2. policy eligibility is checked,
3. return merchandise authorization is created,
4. item is received or inspected,
5. stock is adjusted,
6. refund or credit note is issued,
7. customer and finance systems are updated.

7.9.2 Architectural risks

Returns create risks when:

- the storefront approves a return that finance policy rejects,
- inventory is restocked before inspection rules are applied,
- refunds are triggered twice through retries or manual intervention,

- credit notes in ERP do not align with customer-facing refund status.

These flows require the same rigor as order creation, including idempotency, state control, and clear ownership.

7.10 B2B pricing and account structures

B2B commerce introduces additional integration complexity because customer relationships are not purely transactional. They often involve negotiated pricing, approval chains, sales-assist workflows, credit limits, and account hierarchies.

7.10.1 Price lists and negotiated terms

B2B price lists may depend on:

- customer account,
- contract terms,
- region,
- product family,
- volume thresholds,
- approval status.

These rules are frequently managed in ERP or enterprise sales processes. The storefront may present the result, but it should not become a shadow source of truth for contractual pricing.

7.10.2 Account hierarchies

One B2B customer may represent:

- a legal entity,
- several branches,
- multiple buyers,
- different shipping destinations,
- separate billing relationships.

Architecture must therefore support organization-level identity and policy, not only individual user accounts.

7.11 Using Odoo as a reference ERP

Odoo is useful as a reference platform because it illustrates how enterprise workflows span modules rather than living in a single monolithic screen.

7.11.1 Illustrative Odoo-aligned mapping

- **Sales:** quotations, orders, commercial terms, B2B account workflows.
- **Inventory:** stock movements, reservations, warehouse updates.
- **Accounting:** invoices, payments, credit notes, reconciliation.
- **CRM:** opportunity context, customer engagement history, account management.

The point of using Odoo here is not vendor lock-in. It is to give students a concrete enterprise reference for how commerce-facing systems interact with structured back-office processes.

7.11.2 What should remain vendor-neutral

Even when Odoo is the teaching reference, the following concepts remain general:

- system-of-record analysis,
- event and API contracts,
- master-data governance,
- process ownership,
- enterprise reconciliation and auditability.

7.12 Integration risks and mitigation strategies

Representative risks in ERP/CRM/SCM integration include:

- **Duplicate sources of truth:** mitigated by explicit ownership and reconciliation.
- **Latency mismatch:** mitigated by choosing appropriate sync versus async boundaries.
- **Identifier mismatch:** mitigated by governed master identifiers and mapping tables.
- **Status divergence:** mitigated by explicit lifecycle models and state-mapping rules.
- **Operational blind spots:** mitigated by observability, error queues, and business-level monitoring.

Students should notice that these risks are architectural and organizational, not only technical.

7.13 Artifacts for this chapter

The expected architecture artifacts for ERP/CRM/SCM integration include:

- an integration context diagram showing commerce, ERP, CRM, warehouse, payments, and analytics boundaries,
- a system-of-record matrix for shared entities,
- a small set of data contracts for critical events or APIs,
- and a short narrative for order-to-cash and returns ownership.

7.14 Managerial implications

Enterprise integration decisions affect more than system design. They shape service reliability, customer trust, financial control, and the ability to scale channels. If commerce and ERP are tightly coupled in every user interaction, customer experience becomes hostage to back-office latency. If they are too loosely coupled without governance, the company loses commercial coherence.

The correct architectural aim is not maximal coupling or maximal decoupling. It is controlled coordination with explicit ownership and recoverable failure handling.

7.15 Multiple-choice questions (MCQs)

1. In an integrated commerce + ERP architecture, the *system of record* is best defined as:
 - (a) The system with the nicest UI
 - (b) The system that is most convenient for analytics
 - (c) The authoritative source of truth for a given entity (e.g., invoice) and its lifecycle
 - (d) The system that runs in the cloud
2. Which flow is most strongly associated with *order-to-cash*?
 - (a) Creating a product image embedding
 - (b) Quote/order → delivery → invoice → payment reconciliation
 - (c) Building a recommendation model
 - (d) Defining a Kafka topic
3. Master data management (MDM) is primarily concerned with:
 - (a) Training deep learning models for search
 - (b) Defining and governing shared entities (customer, product, supplier) and preventing duplicates/inconsistencies
 - (c) Choosing the best database index

- (d) Encrypting API requests
- 4. A common integration risk when connecting commerce with ERP is:
 - (a) Too many unit tests
 - (b) Duplicate sources of truth for prices, customers, or inventory
 - (c) Having an event schema
 - (d) Using idempotency keys
- 5. In practice, why do integrations often require both APIs and events?
 - (a) Because events are always cheaper than APIs
 - (b) Because some steps require immediate confirmation while others benefit from decoupling and retries
 - (c) Because ERP systems cannot expose APIs
 - (d) Because events guarantee zero duplicates without design effort

Answer key

- 1. (c)
- 2. (b)
- 3. (b)
- 4. (b)
- 5. (b)

7.16 Exercises (short)

- 1. For each entity, choose a system of record (commerce vs Odoo) and justify: customer, product, price list, inventory on-hand, order, invoice, payment, return.
- 2. Describe two failure modes of inventory sync and propose mitigations (e.g., eventual consistency, reservations, reconciliation jobs).
- 3. Draft a “data contract” for an `OrderCreated` event: required fields, optional fields, and versioning rules.

7.17 Mini-case (Odoo-linked, vendor-neutral)

Scenario: A hybrid company runs a headless storefront for B2C and a B2B portal for wholesale customers. Odoo is used for Sales, Inventory, Accounting, and CRM.

Task: Design the integration for three end-to-end processes:

- **Order-to-cash:** from checkout to invoice posting and payment reconciliation.
- **Returns/refunds:** from return request to stock adjustment and credit note.
- **B2B pricing:** negotiated price lists and approvals without duplicating sources of truth.

For each process, state: key events, required synchronous calls, idempotency strategy, and the minimum monitoring you would implement.

Chapter 8

Recommenders I: Classical Methods + Learning-to-Rank (Week 7)

8.1 Learning objectives

- Explain the business role of recommendation and ranking systems in e-commerce discovery.
- Distinguish memory-based collaborative filtering, matrix factorization, and learning-to-rank approaches.
- Understand the difference between explicit and implicit feedback, and why e-commerce is usually dominated by implicit signals.
- Evaluate recommendation systems using offline ranking metrics while recognizing the limits of offline evaluation.

8.2 Why recommenders matter in e-commerce

Recommendation systems are central to modern e-commerce because large catalogs create a matching problem. Customers rarely inspect all available items. The platform therefore needs mechanisms for deciding which products to surface, in what order, and in which context.

Recommendation is not only a technical personalization problem. It is also a business control layer for:

- product discovery,
- cross-sell and upsell,
- session continuation,
- re-engagement,
- merchandising efficiency,

- and marketplace supply exposure.

This chapter focuses on classical recommenders and ranking methods. These methods remain important because they are interpretable, computationally efficient, and often strong baselines even when deep-learning systems are eventually adopted.

8.3 Recommendation surfaces and tasks

Not all recommender problems are identical. A commerce platform usually contains several recommendation surfaces with different goals.

8.3.1 Common recommendation tasks

- **Homepage recommendations:** highlight likely-interest products for returning or anonymous users.
- **Product detail recommendations:** suggest similar, complementary, or substitute items.
- **Cart recommendations:** support bundle expansion and accessory attachment.
- **Email or push recommendations:** re-engage users outside the session.
- **Search ranking and re-ranking:** order retrieved items for a query.

8.3.2 Retrieval versus ranking

It is useful to distinguish two subproblems:

- **Candidate generation:** find a manageable subset of items from a very large catalog.
- **Ranking or re-ranking:** order the candidates for the current context.

This chapter emphasizes ranking-oriented thinking and classical candidate selection methods. Chapter 8 will extend this to deep retrieval and embedding-heavy systems.

8.4 Feedback in commerce recommendation

Recommendation systems learn from user behavior. The nature of that behavior strongly affects modeling choices.

8.4.1 Explicit feedback

Explicit feedback is directly provided by users, such as:

- star ratings,
- likes or dislikes,
- structured review scores,
- preference selections.

Explicit feedback is valuable because it is interpretable, but in many commerce settings it is sparse. Most customers do not rate enough items to support the full recommendation problem.

8.4.2 Implicit feedback

Implicit feedback is inferred from behavior:

- clicks,
- product views,
- add-to-cart events,
- purchases,
- dwell time,
- repeat visits,
- returns or cancellations as negative indicators.

E-commerce is largely an implicit-feedback domain. This makes modeling harder because a missing interaction does not necessarily mean dislike. It may simply mean the user never saw the item.

8.4.3 Signal strength

Not all implicit signals are equally informative. For example:

- a purchase usually signals stronger intent than a click,
- an add-to-cart event may indicate interest but not final preference,
- a return may weaken confidence in the original purchase signal,
- repeated category browsing may indicate exploration rather than commitment.

Feature and label design should therefore reflect business meaning, not just event counts.

8.5 Memory-based collaborative filtering

Classical collaborative filtering starts from the idea that similar users may like similar items, and similar items may be preferred by similar users.

8.5.1 User-based collaborative filtering

In user-based collaborative filtering, the system identifies users whose historical behavior resembles the target user and recommends items those similar users engaged with. Conceptually:

- define a user-item interaction matrix $R = [r_{ui}]$,
- compute user-user similarity,
- aggregate neighbor preferences for unseen items.

This approach is intuitive but often scales poorly for large consumer systems with many users and rapidly changing behavior.

8.5.2 Item-based collaborative filtering

Item-based collaborative filtering compares items rather than users. If two items are frequently co-viewed, co-purchased, or similarly interacted with, one can recommend one when the other is observed.

Item-based methods are often more stable in commerce because items change more slowly than user sessions. They work well for:

- “similar items”,
- “customers also bought”,
- and complementary recommendation widgets.

8.5.3 Similarity measures

Common similarity measures include cosine similarity, Pearson correlation, and co-occurrence-based heuristics. The choice depends on data type and normalization assumptions. Architecturally, what matters is that similarity calculations must be based on well-defined interaction windows and item identities.

8.6 Matrix factorization

Memory-based methods are useful, but they struggle with sparsity and scale. Matrix factorization addresses this by learning low-dimensional latent representations for users and items.

8.6.1 Basic idea

Given an interaction matrix R , matrix factorization approximates it as:

$$\hat{r}_{ui} = \mathbf{p}_u^\top \mathbf{q}_i$$

where $\mathbf{p}_u$ is a latent vector for user u and $\mathbf{q}_i$ is a latent vector for item i .

These latent factors capture hidden structure such as taste patterns, style clusters, or price sensitivity without requiring manually defined categories.

8.6.2 Why factorization works well

Matrix factorization often performs strongly because:

- it compresses sparse interactions into learnable structure,
- it generalizes beyond direct co-occurrence,
- it supports efficient ranking once vectors are learned,
- it forms a bridge toward later embedding-based methods.

8.6.3 Limitations

Despite its strengths, factorization has limits:

- it handles cold-start users and items poorly without additional features,
- it may ignore rich session context,
- it can struggle when preferences shift quickly,
- it usually needs careful treatment of implicit feedback and negative sampling.

8.7 Implicit-feedback recommendation

Most commerce recommenders do not predict star ratings. Instead, they infer preference from implicit interactions. This changes both objective design and evaluation.

8.7.1 Positive-only observations

In implicit-feedback settings, observed interactions are often treated as positive evidence, while unobserved interactions are ambiguous. Training usually requires assumptions such as:

- sampling unobserved items as negatives,
- weighting confidence by interaction type,
- or optimizing pairwise ranking losses instead of rating prediction.

8.7.2 Confidence weighting

Different interactions can be assigned different strengths. For example:

- purchase > add-to-cart > click > impression,
- repeat purchase may strengthen confidence,
- return or refund events may weaken confidence.

This is not just a modeling choice. It is a business-rule decision and should align with how the organization interprets engagement quality.

8.8 Cold start and sparsity

Cold start is one of the defining challenges of recommendation systems. It occurs when the system has insufficient data about a new user, a new item, or a new marketplace participant.

8.8.1 User cold start

New users may have little or no history. Common mitigation strategies include:

- popularity-based defaults,
- category-entry onboarding questions,
- contextual recommendations from landing page or query behavior,
- session-based short-term signals.

8.8.2 Item cold start

New items may not yet have interaction data. Mitigation strategies include:

- using metadata such as category, brand, and price range,
- manual merchandising boosts,
- seller and catalog-quality priors,
- controlled exploration to gather initial interactions.

8.8.3 Marketplace and fairness considerations

In marketplaces, recommendation systems also influence seller exposure. If the system only amplifies already popular items, new or smaller sellers may never gather enough interaction data to compete. This creates a business and governance concern, not just a model-performance issue.

8.9 Learning-to-rank

Recommendation is often framed as a ranking problem rather than a rating-prediction problem. Learning-to-rank (LTR) methods optimize the order of items shown to users for a given context.

8.9.1 Pointwise, pairwise, and listwise views

LTR methods are often grouped into three families:

- **Pointwise:** predict a score or label for each item independently.
- **Pairwise:** learn which of two items should be ranked higher.
- **Listwise:** optimize over the ordering of whole ranked lists.

8.9.2 Why LTR is important in commerce

E-commerce surfaces almost always have limited slots. The platform is not merely predicting whether an item is relevant; it is deciding which few items to prioritize. LTR is therefore especially appropriate for:

- search result ordering,
- recommendation carousel ranking,
- category page re-ranking,
- email content selection.

8.9.3 Feature design for LTR

Classical ranking models often use hand-engineered or structured features such as:

- user affinity to category or brand,
- historical click-through rate,
- purchase propensity,
- price distance from user norms,
- popularity and freshness,
- inventory availability,
- promotion eligibility,
- query-item lexical match for search contexts.

This makes LTR highly practical because structured business features are often already available in commerce systems.

8.10 Offline evaluation

Recommendation systems need offline evaluation for rapid iteration, model comparison, and failure analysis before live deployment.

8.10.1 Common ranking metrics

Frequently used metrics include:

- $\text{precision@}k$,
- $\text{recall@}k$,
- mean average precision ($\text{MAP@}k$),
- normalized discounted cumulative gain ($\text{NDCG@}k$),
- mean reciprocal rank (MRR) in some search-style tasks.

8.10.2 What these metrics measure

- **Precision@ k :** how many of the top- k items are relevant.
- **Recall@ k :** how much of the relevant set is recovered in the top- k .
- **NDCG@ k :** rewards placing more relevant items higher in the ranked list.

These metrics are useful because they reflect ranking quality rather than raw score accuracy.

8.10.3 Limits of offline evaluation

Offline metrics are not enough on their own. They cannot fully capture:

- causal business impact,
- position bias and feedback loops,
- changing user behavior under new experiences,
- interference with other surfaces such as promotions or search.

This is why offline evaluation should be seen as a filtering step, not the final decision criterion.

8.11 Online evaluation and business impact

Ultimately, recommender systems exist to improve user and business outcomes in production. Typical online measures include:

- click-through rate,
- add-to-cart rate,
- conversion rate,
- average order value,
- revenue per session,
- repeat engagement,
- return rate or dissatisfaction signals.

However, online KPI interpretation must be careful. A recommender can increase clicks while harming profitability or reducing diversity. Guardrail metrics are therefore important.

8.12 Biases and practical risks

Recommendation systems are affected by several structural biases.

8.12.1 Popularity bias

Popular items accumulate more exposure and therefore more interactions, reinforcing their future dominance. This can suppress niche discovery and reduce long-tail coverage.

8.12.2 Selection bias

Models only observe interactions for items that were shown. Unseen items do not generate evidence, so naive training pipelines may overestimate the quality of historically exposed items.

8.12.3 Position bias

Higher-ranked items receive more clicks partly because they are higher ranked, not only because they are more relevant. This complicates interpretation of click data.

8.12.4 Business-constraint tension

The “best” recommendation by predicted engagement may conflict with:

- margin objectives,
- inventory constraints,
- seller-fairness rules,
- compliance or policy requirements,
- user-experience diversity needs.

Recommendation is therefore not a pure prediction task. It is a constrained decision problem embedded in business policy.

8.13 A practical recommender architecture

Even for classical methods, recommender systems are not just models. They require an end-to-end architecture that includes:

- event collection and identity linking,
- feature generation,
- offline training,
- candidate generation,
- ranking logic,
- serving interfaces,
- monitoring and experimentation.

This chapter remains focused on the modeling layer, but students should keep the surrounding system in mind. The model only creates value when the architecture can serve it reliably and measure its effect.

8.14 Hands-on lab: a baseline recommender

The hands-on goal for this chapter is to build and evaluate a classical recommendation baseline.

8.14.1 Lab scenario

Students are given a synthetic or public interaction dataset containing user IDs, item IDs, timestamps, and event types such as view, add-to-cart, and purchase.

8.14.2 Lab tasks

1. Construct a user–item interaction matrix using implicit feedback.
2. Implement a simple popularity baseline and one collaborative filtering or matrix-factorization baseline.
3. Evaluate the models using $\text{MAP}@K$ and $\text{NDCG}@K$.
4. Analyze cold-start and sparsity cases explicitly.
5. Propose one feature-based re-ranking rule that incorporates a business constraint such as inventory or diversity.

8.14.3 Expected learning outcome

By the end of the lab, students should be able to explain not only which model performed better offline, but why the evaluation setup and business constraints matter.

8.15 Managerial implications

Classical recommendation methods remain valuable in practice because they provide:

- strong baselines,
- fast iteration cycles,
- reasonable interpretability,
- lower infrastructure cost than many deep systems.

Organizations often make a mistake by jumping directly to deep-learning architectures without first establishing strong baselines, high-quality data, and trustworthy evaluation pipelines. For many commerce surfaces, the real bottleneck is not model sophistication but poor signal design or unclear business objectives.

8.16 Multiple-choice questions (MCQs)

1. In e-commerce recommendation, which statement best distinguishes *implicit* from *explicit* feedback?
 - (a) Implicit feedback is always more accurate than explicit feedback.
 - (b) Explicit feedback comes directly from user judgments such as ratings, whereas implicit feedback is inferred from behaviors such as clicks or purchases.
 - (c) Explicit feedback can only be used in search ranking.
 - (d) Implicit feedback means the platform has no user identifiers.

2. Item-based collaborative filtering is often attractive in commerce because:
 - (a) items are usually more stable than users and item-item similarities can support scalable recommendation widgets.
 - (b) it eliminates cold start entirely.
 - (c) it requires no interaction data.
 - (d) it guarantees causal lift.
3. In matrix factorization, the expression $\hat{r}_{ui} = \mathbf{p}_u^\top \mathbf{q}_i$ is intended to represent:
 - (a) a network timeout policy.
 - (b) the predicted affinity between user u and item i using latent factors.
 - (c) the exact accounting value of an order.
 - (d) a lossless compression ratio for logs.
4. Which metric is most appropriate for evaluating the quality of a top- k ranked recommendation list offline?
 - (a) NDCG@ k
 - (b) Mean absolute error on shipping cost
 - (c) CPU utilization
 - (d) Database replication lag
5. Why is learning-to-rank especially relevant in e-commerce?
 - (a) Because commerce systems never need candidate generation.
 - (b) Because the platform must decide which limited set of items to place at the top of a ranked surface.
 - (c) Because it replaces online experimentation.
 - (d) Because it removes the need for features.

Answer key

1. (b)
2. (a)
3. (b)
4. (a)
5. (b)

8.17 Exercises (short)

1. Compare user-based, item-based, and matrix-factorization approaches for a medium-sized commerce catalog. State one strength and one weakness of each.
2. Define an implicit-feedback weighting scheme for view, add-to-cart, purchase, and return events. Explain the business reasoning behind your weights.
3. Design a simple learning-to-rank feature set for a “recommended for you” widget. Include at least one user feature, one item feature, one context feature, and one business-rule feature.

8.18 Mini-case (Odoo-linked, vendor-neutral)

Scenario: A company uses a headless storefront, Odoo for ERP, and a central data platform. The business wants recommendation widgets on the homepage, product pages, and cart page.

Task: Propose a classical recommender design:

- Define which interaction signals from storefront and Odoo-derived outcomes should be used for training.
- Choose a baseline method (item-based CF, matrix factorization, or feature-based ranking) for each surface and justify the choice.
- Identify two business constraints the ranking layer should respect, such as inventory, margin, seller exposure, or return risk.

Chapter 9

Recommenders II: Deep Learning, Embeddings, and Retrieval (Week 8)

9.1 Scope (placeholder)

- Two-tower retrieval, embeddings, sequence models, and candidate generation vs re-ranking.
- Practical concerns: cold-start, diversity, bias, and latency.

9.2 Hands-on lab (placeholder)

- Train embeddings; build a retrieval + re-ranking pipeline (high level).

Chapter 10

Pricing and Promotions with ML + Optimization (Week 9)

10.1 Scope (placeholder)

- Demand forecasting, elasticity, uplift, and constraints (inventory/competition).
- Optimization approaches and experiment design for pricing.

10.2 Hands-on lab (placeholder)

- Forecast baseline + simple constrained optimization scenario.

Chapter 11

Fraud, Risk, and Trust & Safety (Week 10)

11.1 Scope (placeholder)

- Fraud types: payment, account takeover, promotion abuse, refund abuse.
- Methods: supervised learning, anomaly detection, graph ML, and rule+ML hybrids.

11.2 Hands-on lab (placeholder)

- Build a risk scoring baseline + thresholding and cost-sensitive evaluation.

Chapter 12

Customer Service Automation with LLMs (Week 11)

12.1 Scope (placeholder)

- RAG, tool use, workflow orchestration, evaluation, and guardrails.
- When to use LLMs vs classical NLP vs rules.

12.2 Hands-on lab (placeholder)

- Prototype: FAQ assistant with retrieval + citations (local documents only).

Chapter 13

MLOps + DataOps for E-Commerce (Week 12)

13.1 Scope (placeholder)

- Platform-agnostic MLOps/DataOps: CI/CD for ML, pipelines, and reproducibility.
- Feature stores (concept), training-serving skew, monitoring, drift, and SLOs.
- Experiment tracking, model governance, and evaluation at scale.

13.2 Artifacts (placeholder)

- Deliverable: ML system design doc + monitoring plan.

Chapter 14

Security, Privacy, Compliance, and Responsible AI (Week 13)

14.1 Scope (placeholder)

- Payments/security basics, privacy-by-design, and data minimization.
- Threat modeling for e-commerce and AI features (LLM-specific risks included).
- Responsible AI: fairness, explainability, auditability, and human-in-the-loop controls.

14.2 Discussion (placeholder)

- Risk assessment template for an AI feature.

Chapter 15

Capstone: Enterprise-Grade AI E-Commerce Architecture (Week 14)

Capstone scenario (agreed)

- Scenario: end-to-end “AI commerce + Odoo” reference architecture for a single company.
- Goal: demonstrate EA alignment, ERP integration, and measurable AI value with operational readiness.
- Business model: hybrid (B2C + B2B).
- Commerce front-end: both (compare Odoo Website/eCommerce vs a separate headless storefront integrating to Odoo).
- AI stack: both (ML for recs/fraud/pricing + LLMs for support/ops).

15.1 Scope (placeholder)

- End-to-end architecture: channels, services, data platform, ML/LLM layer, ERP integration.
- Evaluation: business metrics, technical metrics, and operational readiness.

15.2 Capstone deliverables (placeholder)

- Architecture pack (diagrams), backlog, data contracts, and evaluation plan.

Appendix A

Datasets and Synthetic Data for Teaching (placeholder)

Appendix B

Template: System Design Document (placeholder)

Appendix C

Template: Data Contract (placeholder)

Appendix D

Template: Model Card and Risk Assessment (placeholder)